

Kamer Kaya
.. Sabancı ..
Üniversitesi

Yüksek Başarımlı (Paralel) Hesaplama

Veri Bilim - Yapay Öğrenme Yaz Okulu, 2017
"Matematiksel Temeller ve Vaka Çalışmaları"

Yüksek başarım?

- 1. (isim) Elde edilen bir başarı
- 2. Herhangi bir olayı veya durumu başarma isteği ve gücü
- 3. Kişinin yapabileceği **en iyi** derece, performans
- 4. Herhangi bir eseri, oyunu, işi vb.ni ortaya koyarken gösterilen başarı, performans

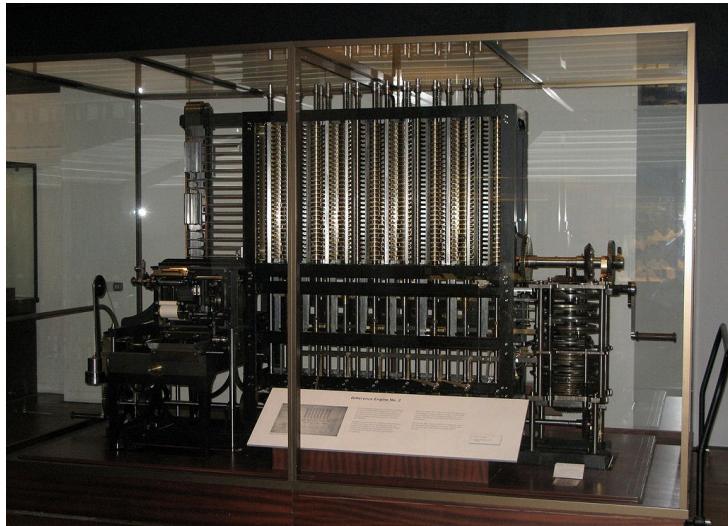
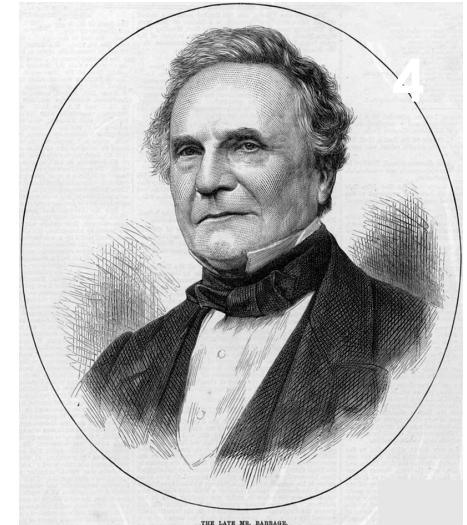
1800ler



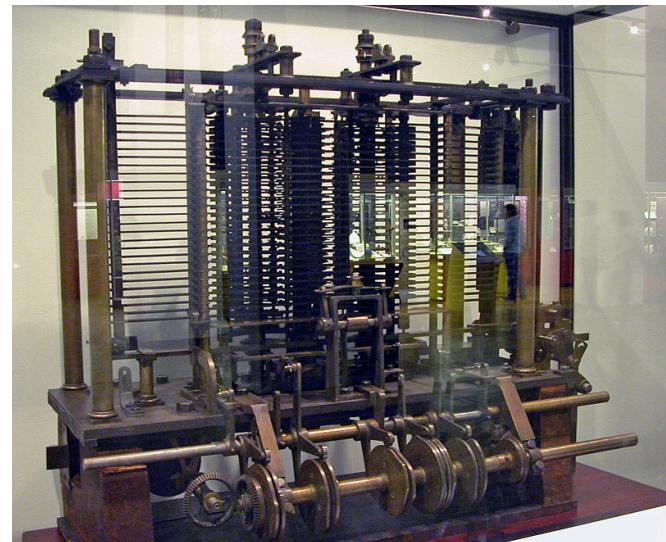
"Computer" kelimesinin ilk defa New York Times gazetesinde gözüktüğü tarih: May 2, 1892; US Civil Service Commission tarafından verilen bir ilan:

"A Computer Wanted. [...] The examination will include the subjects of algebra, geometry, trigonometry, and astronomy."

1800'lər Charles Babbage



The Science Museum Difference Engine No. 2, Babbage'ın birebir tasarımlı



Analytical Engine prototipi, The Science Museum

1930lar

Alan Turing



ON COMPUTABLE NUMBERS, WITH AN APPLICATION TO
THE ENTScheidungsproblem

By A. M. TURING.

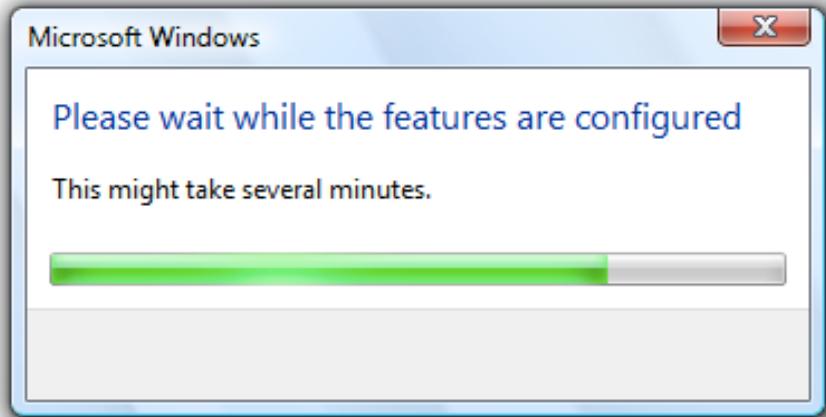
[Received 28 May, 1936.—Read 12 November, 1936.]

1936

The “computable” numbers may be described briefly as the real numbers whose expressions as a decimal are calculable by finite means. Although the subject of this paper is ostensibly the computable *numbers*, it is almost equally easy to define and investigate computable functions of an integral variable or a real or computable variable, computable predicates, and so forth. The fundamental problems involved are, however, the same in each case, and I have chosen the computable numbers for explicit treatment as involving the least cumbrous technique. I hope shortly to give an account of the relations of the computable numbers, functions, and so forth to one another. This will include a development of the theory of functions of a real variable expressed in terms of computable numbers. According to my definition, a number is computable if its decimal can be written down by a machine.

1930lar

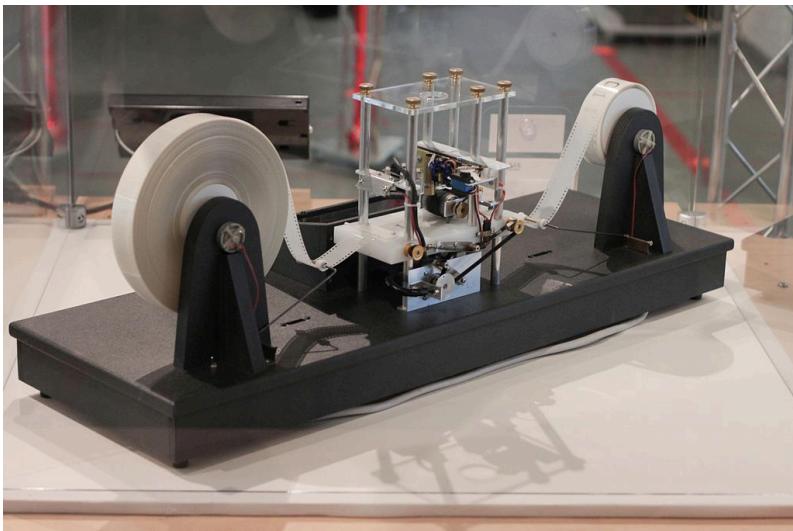
Alan Turing



Hesaplanabilirlik teorisinde, Durma problemi, keyfi bir bilgisayar programının ve girdinin açıklamasından, programın koşmayı bitirip bitirmeyeceğine ya da sonsuza dek koşmaya devam edip etmeyeceğine karar verme problemidir.

Büyük olasılıkla durmayacak bir program

1930lar Alan Turing



Bir Turing makinesi

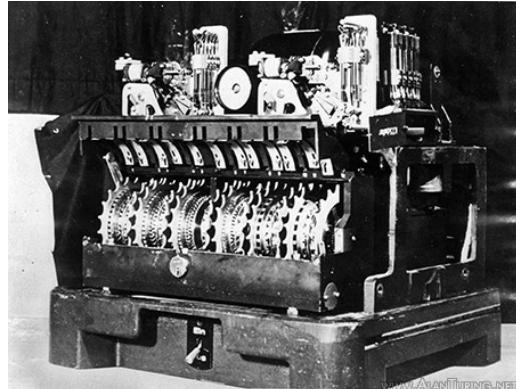
Hesaplanabilirlik teorisinde, bir veri-manipülasyon kuralları sistemi (bir bilgisayarın komut kümesi, bir programlama dili veya hücresel bir otomat gibi) tek bir bantlı Turing makinesini simüle etmek için kullanılabiliyorsa, Turing tam veya hesaplamalı evrensel olduğu söylenir .

1939-1945 II. Dünya Savaşı

İYİ ADAMLAR

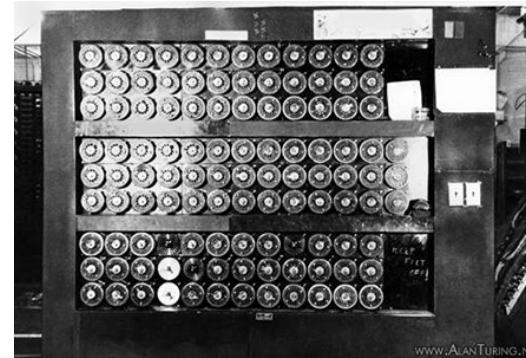


Enigma



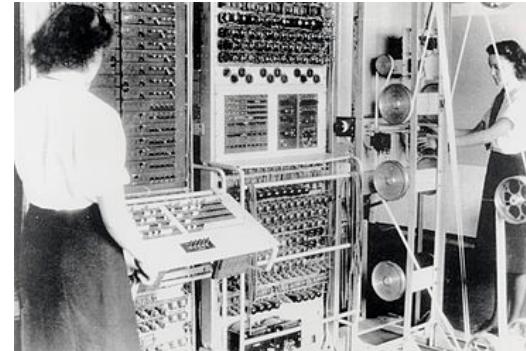
A Lorenz Şifreleme Makinası

vs.

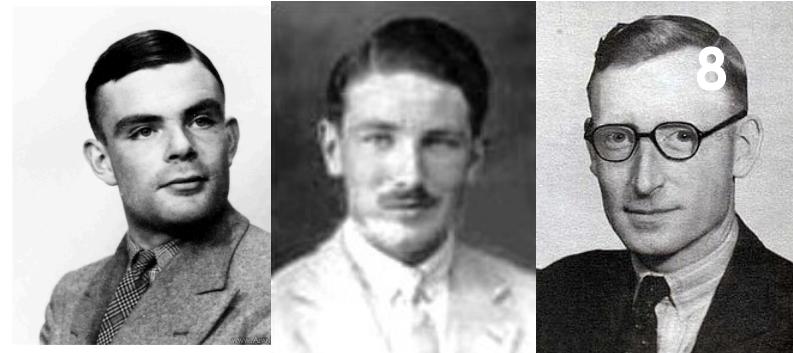


Bombe

vs.



Colossus



Turing Welchman Flowers

KÖTÜ ADAMLAR

1936-38

Z1

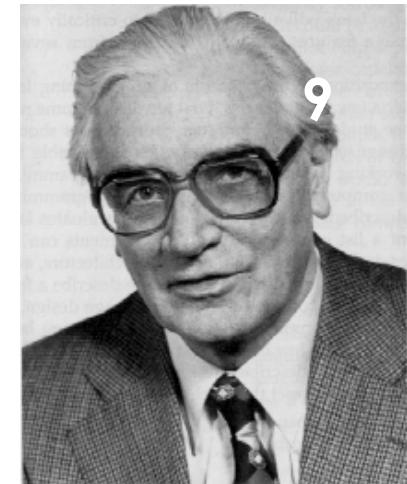
1939-40

Z2

1941...

Z3...

(TURING TAM)



Konrad Zuse



Zuse Z3 replikası
Deutsches Museum,
Munich

Ortalama hesaplama hızı: ek - 0.8 saniye, çarpma - 3 saniye

Aritmetik birim: İkili kayan noktalı, 22 bit, ekle, çıkart, çarp, böl, karekök

Veri belleği: 22 bit uzunluğa sahip 64 sözcük

Program hafızası: Delinmiş selüloit bant

Girdi: Ondalık kayan nokta sayıları

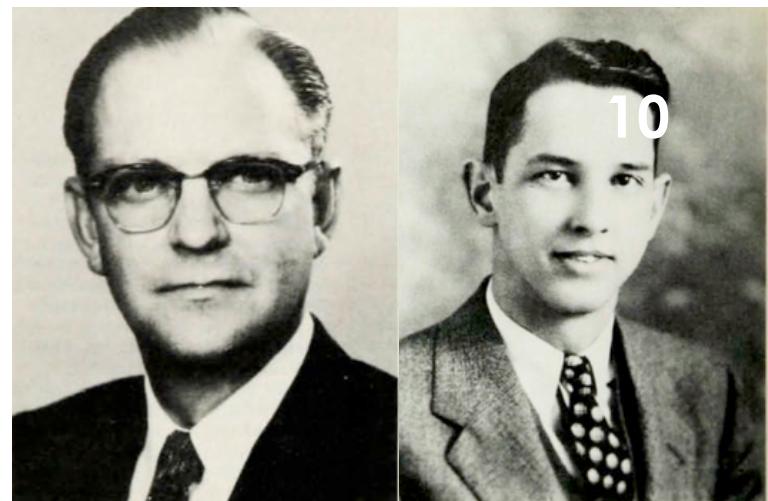
Frekans: 5.3 Hertz

Güç tüketimi: Yaklaşık 4.000 watt

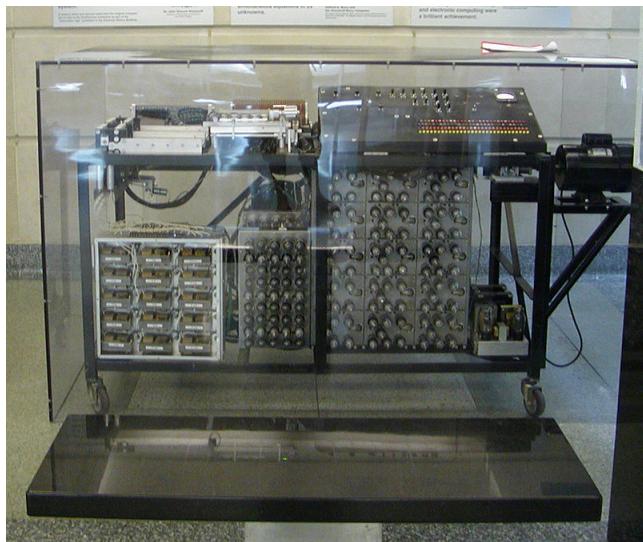
Ağırlık: Yaklaşık 1 ton (2,200 lb)

1939-42

Atanasoff-Berry Bilgisayarı(ABC)



10



Atanasoff-Berry bilgisayarı
Durham Center, Iowa State
University

Veri Bilim - Yapay Öğrenme Yaz Okulu, 2017 "Matematiksel Temeller ve Vaka Çalışmaları"

Atanasoff

Berry

- **Üç önemli fikir:**
- Tüm sayıları ve verileri temsil etmek için ikili basamakları kullanma
- Tüm hesaplamaları, tekerlekler, mandallar veya mekanik anahtarlar yerine elektronik kullanarak gerçekleştirmek
- Hesaplama ve hafızanın ayrıldığı bir sistemi kullanmak.

1945-1960s



ENIAC (1945)



IBM 5Mb HD (1956)

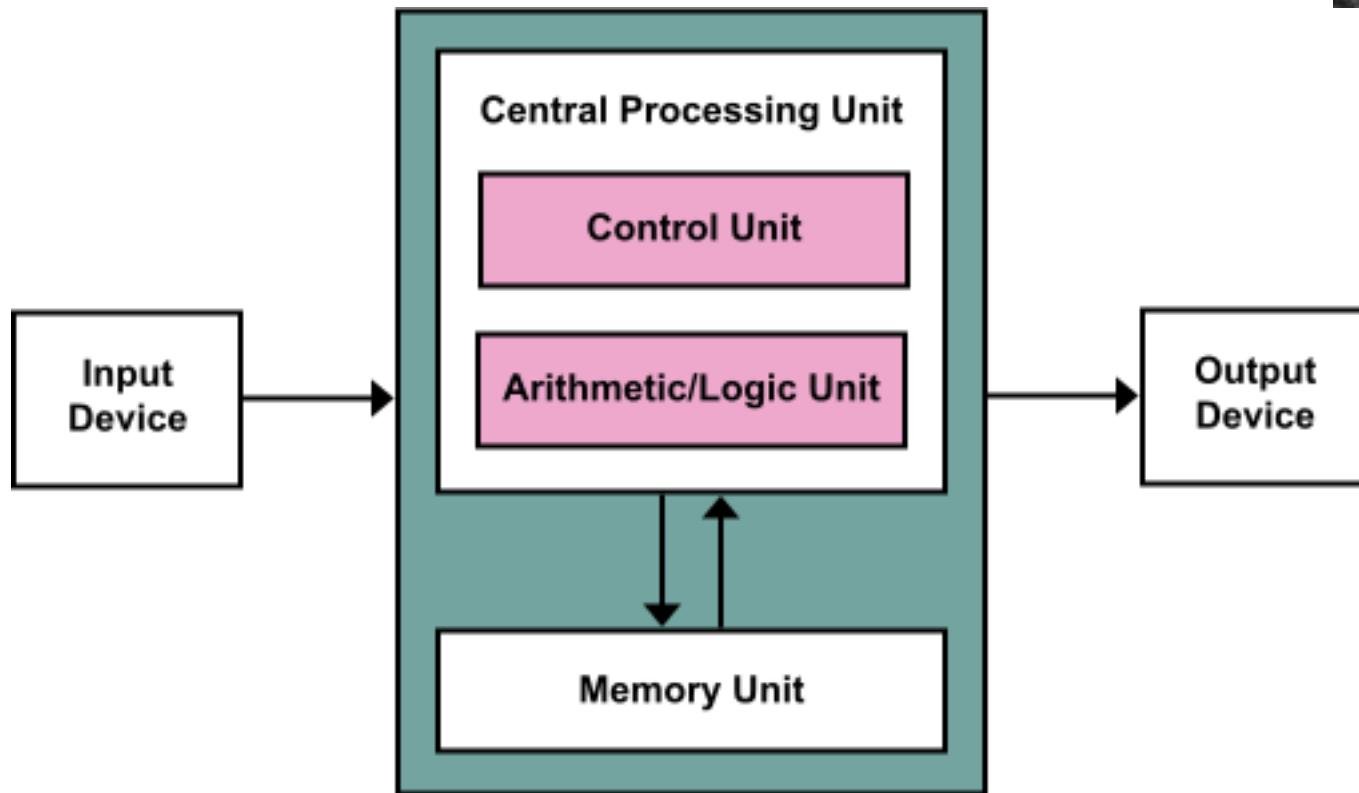


1945

Von Neumann Mimarisi



John von
Neumann



o zamanlar geleceği tahmin etmek
oldukça güçtü

“Sanırım bütün dünyada sadece 5
bilgisayarlık bir pazar olacak”

IBM başkanı, 1943



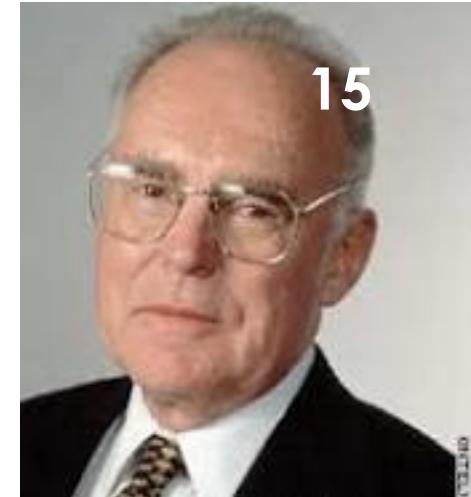
ama bazıları dğour tahmin
edebilmişti

“Gelecekte bilgisayarların ağırlığı en fazla

1.5 ton
olacak.”

Popular Mechanics, 1949

1965 Moore's Law



Cramming More Components onto Integrated Circuits

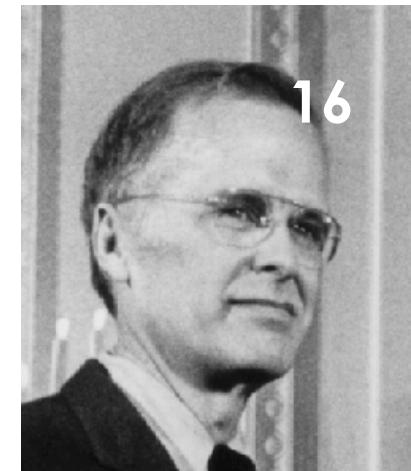
Gordon Moore

GORDON E. MOORE, LIFE FELLOW, IEEE

With unit cost falling as the number of components per circuit rises, by 1975 economics may dictate squeezing as many as 65 000 components on a single silicon chip.

1975: Yarı iletken karmaşıklığı, yaklaşık 1980 yılına kadar yıllık iki katına çıkmaya devam edecek ve bundan yaklaşık iki yılda bir iki katına çıkacaktı

1977 von Neumann darboğazı



John Backus

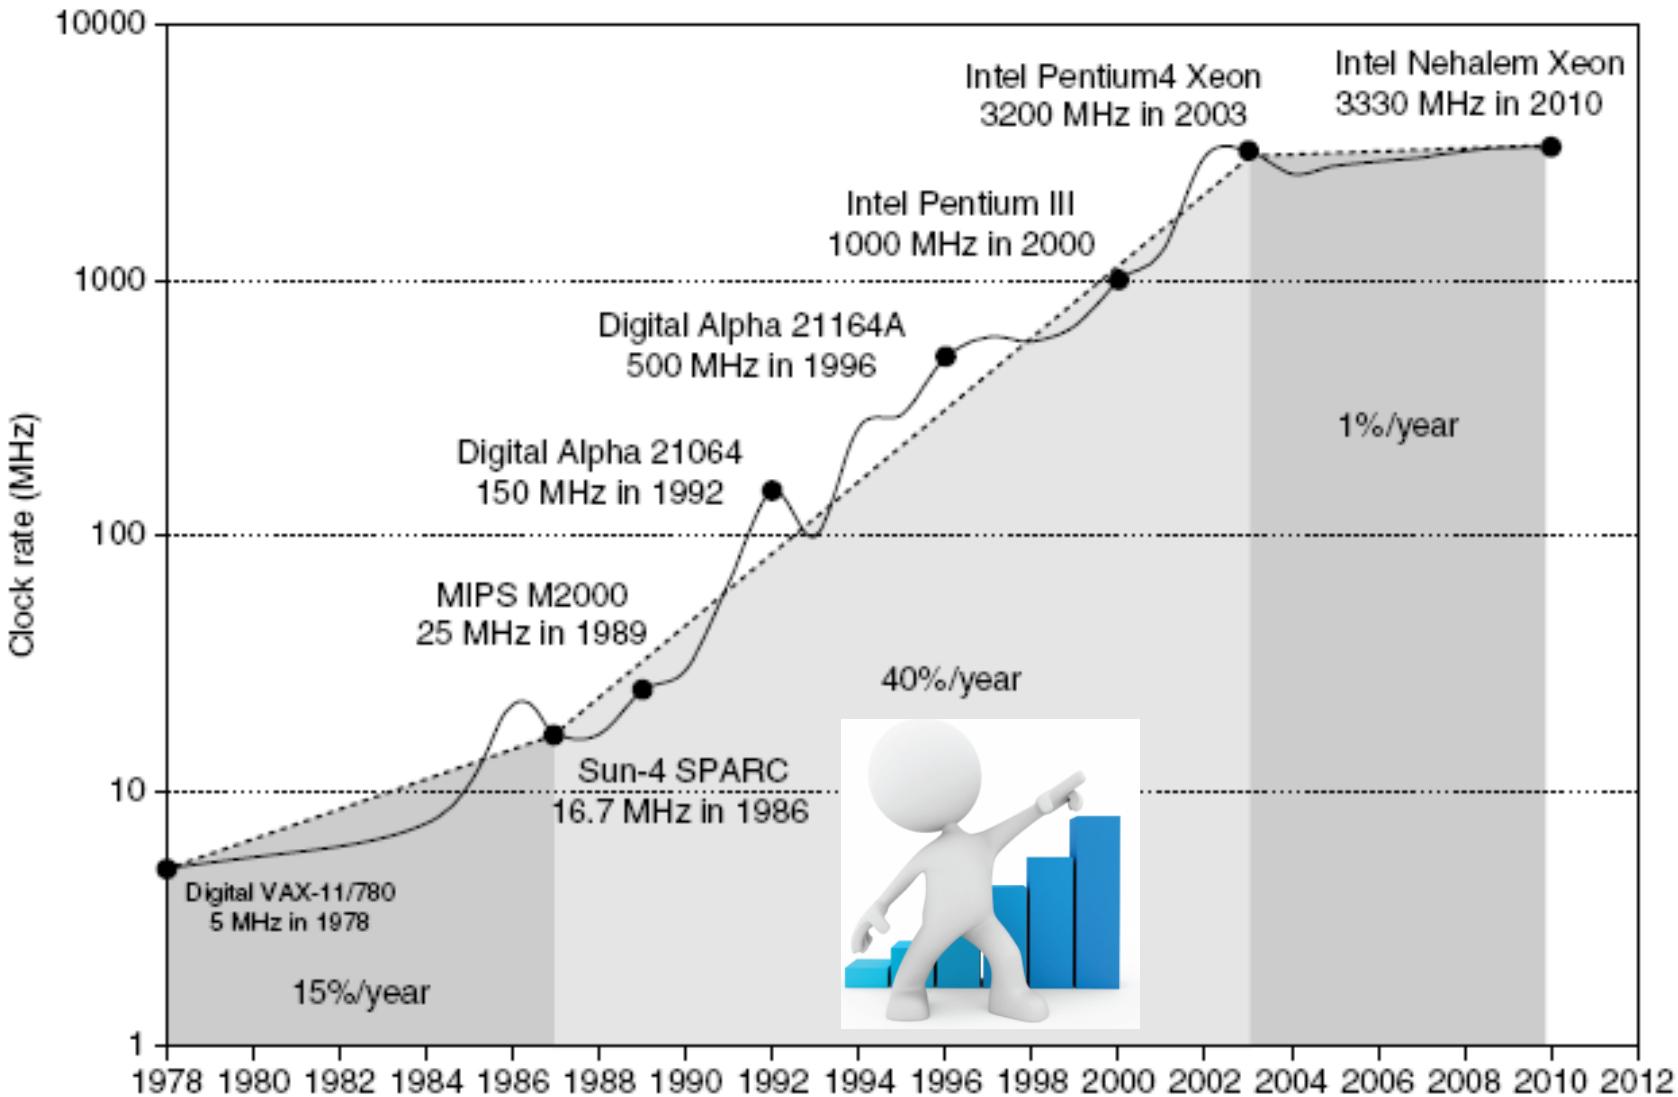
- Bir sistem darboğazı: Merkezi İşlemci Birimleri (CPU'lar) ve Rasgele Erişim Belleği arasındaki bant genişliği tipik bir CPU'nun dahili olarak veri işleyebileceği hızdan çok daha düşüktür.
- Bir 'entelektüel darboğaz': programcılar CPU ve RAM arasında gidip gelen "bir sürü gereklı/gereksiz kelimeyi" azaltmak için gereklı kod optimizasyonlarını düşünmek için çok zaman harcıyorlar.

"... programming is basically planning and detailing the enormous traffic of words through the von Neumann bottleneck, and much of that traffic concerns not significant data itself, but where to find it." [1977 Turing Award Lecture]

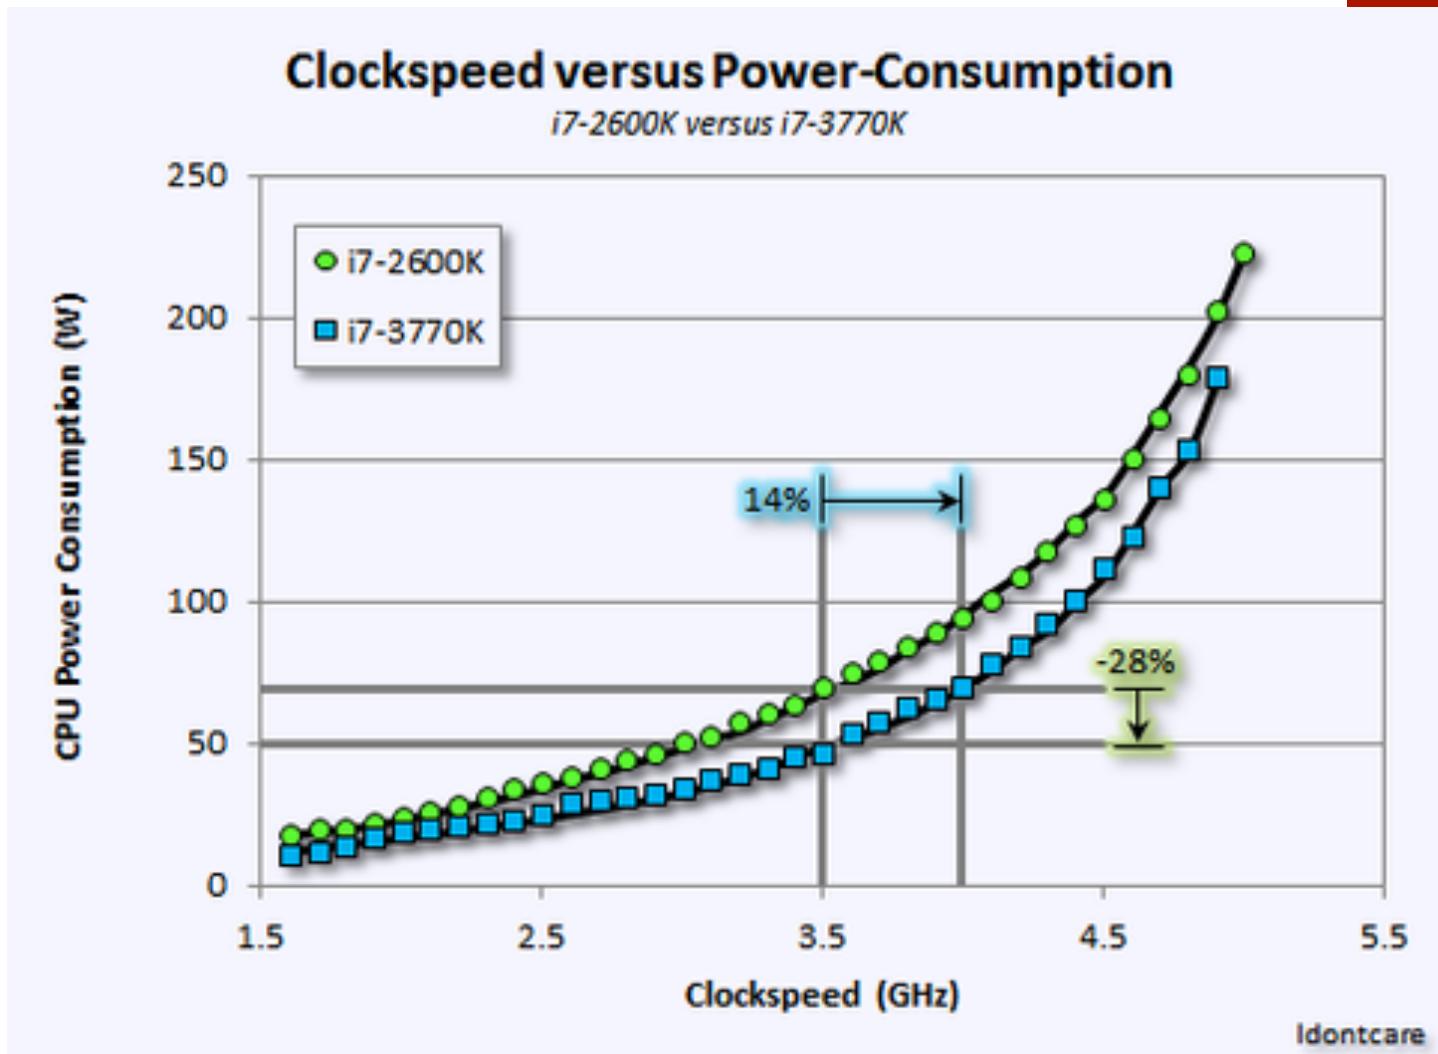
Bir işlemci



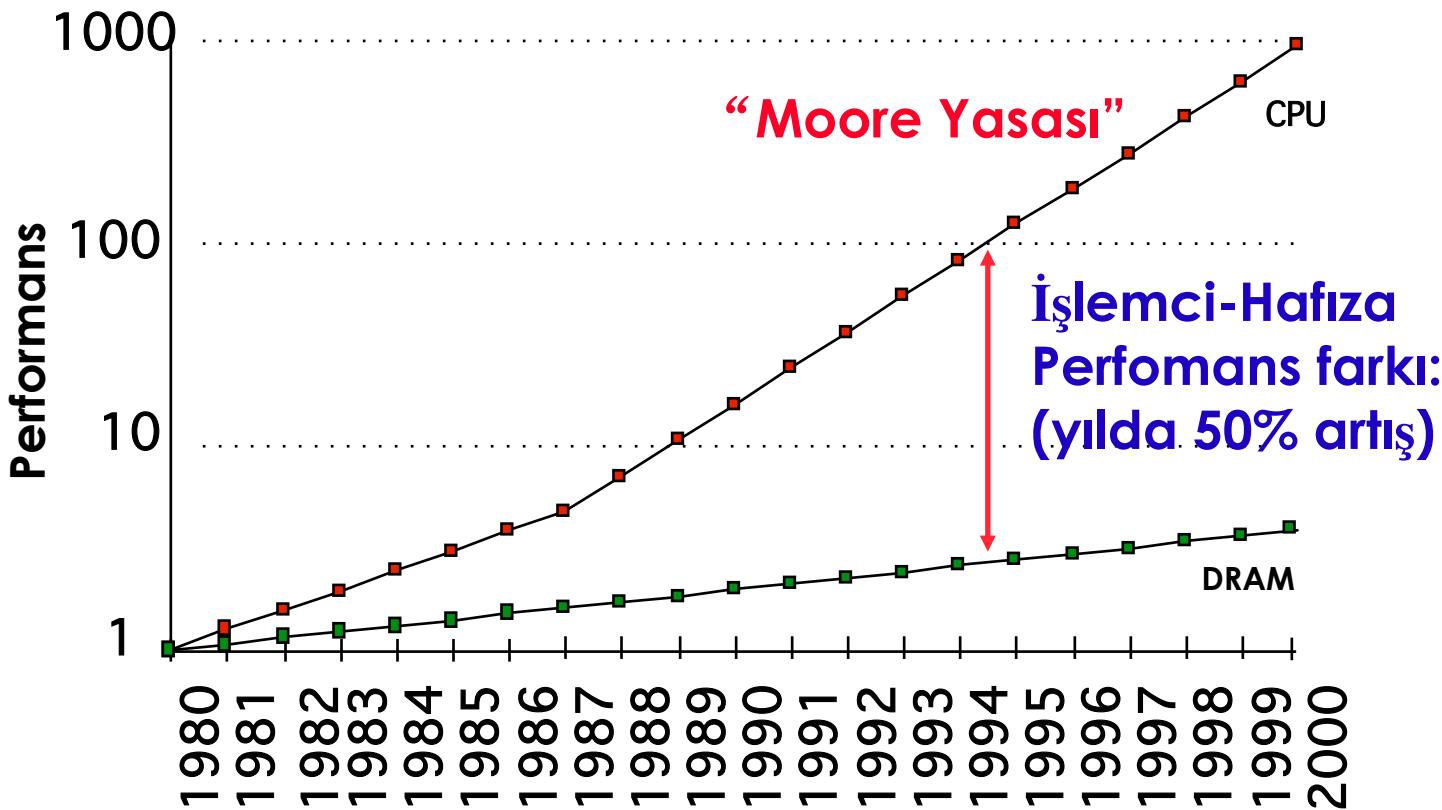
Bir zamanlar...



GÜÇ duvarı



Hafıza duvarı



Hafıza duvarını aşmak

CPU yazmaçları

100s Byte

300 - 500 ps (0.3-0.5 ns)

L1 ve L2 önbelleği

10ns-100ns

\$1000s/ GByte

G Bytes

80ns- 200ns

~\$100/ GByte

T Byte, 10 ms

(10,000,000 ns)
~\$1 / GByte

Sonsuz

saniye-dakika

~\$1 / GByte

Yazmaç

Komut işlenenleri

derleyici
1-8 byte

L1 Önbellek

Blok

L2 Önbellek

Blok

Hafıza

Sayfa

Disk

Dosya

Kaset

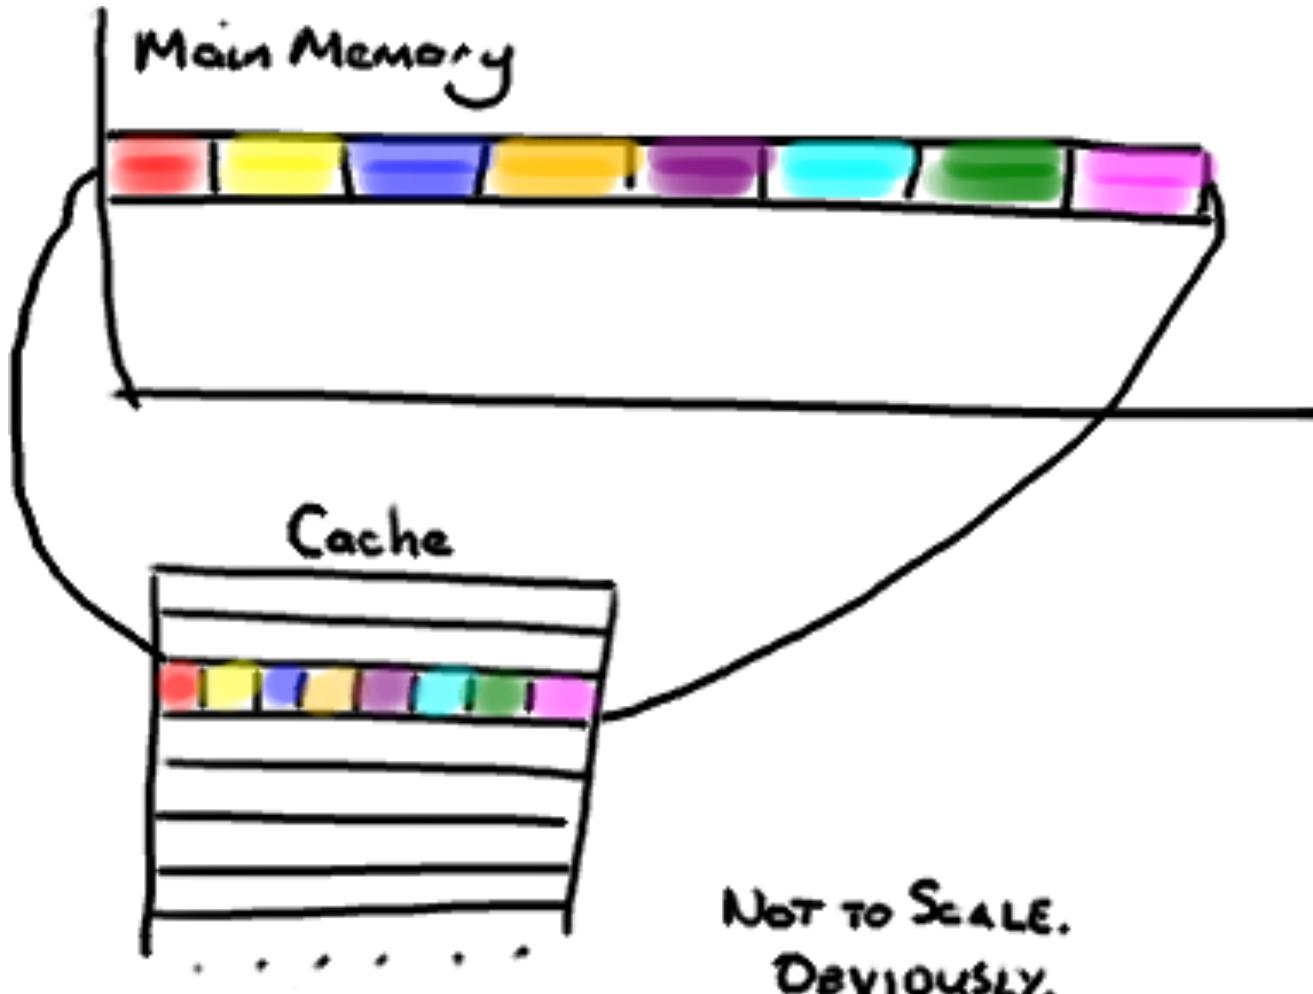
Önbellek yönetici
32-64 byte

Önbellek yönetici
64-128 byte

İşletim sistemi
4K-8K byte

kullanıcı
Mbyte

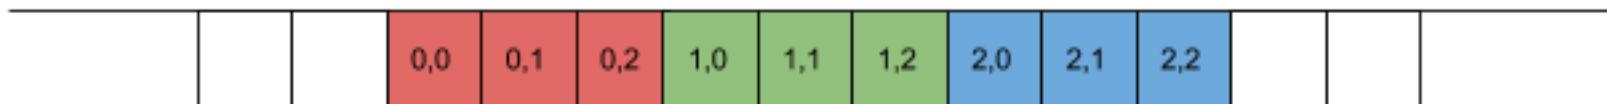
Önbellek: en iyi dostunuz



Matrislerin sütunlarının toplAMI

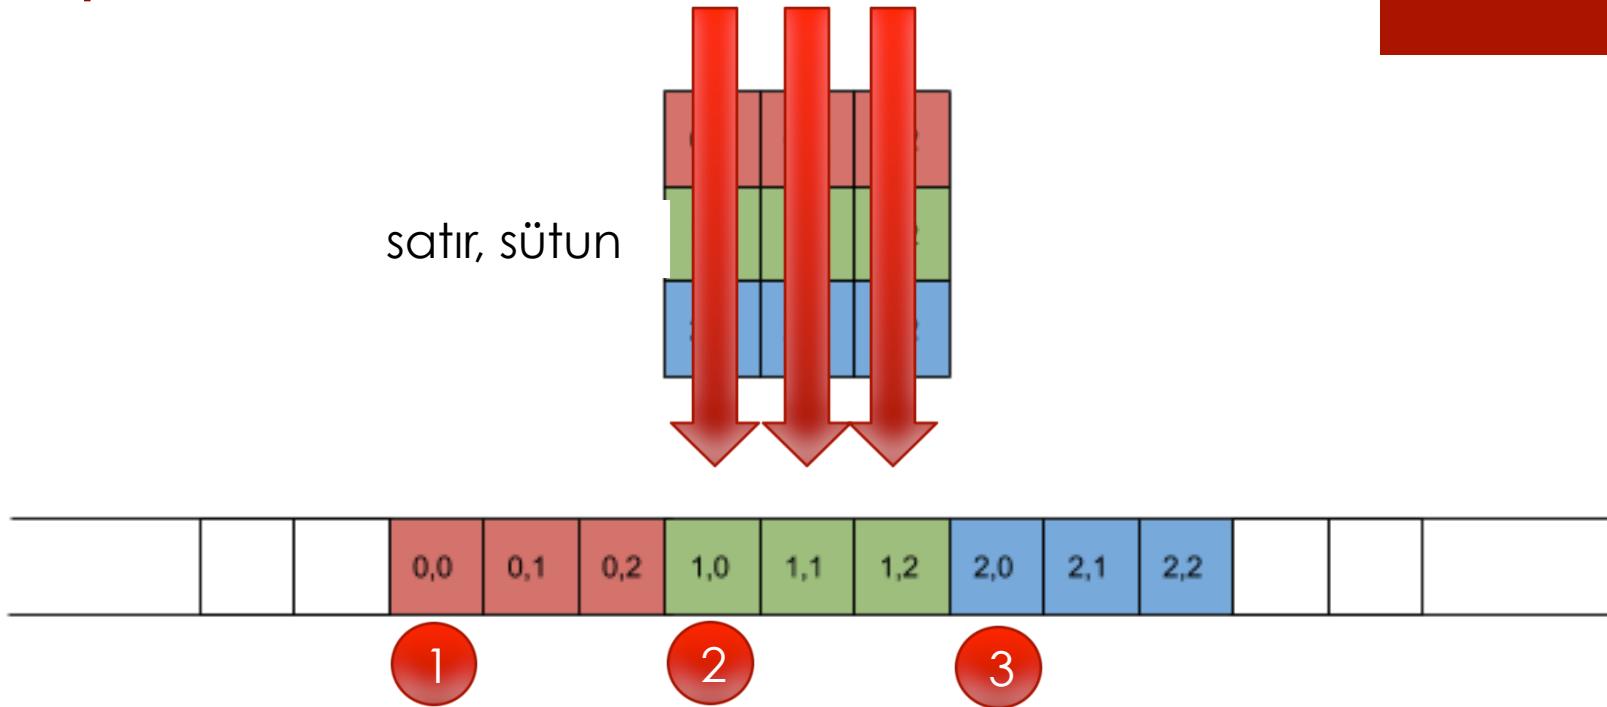
satır, sütun

0,0	0,1	0,2
1,0	1,1	1,2
2,0	2,1	2,2

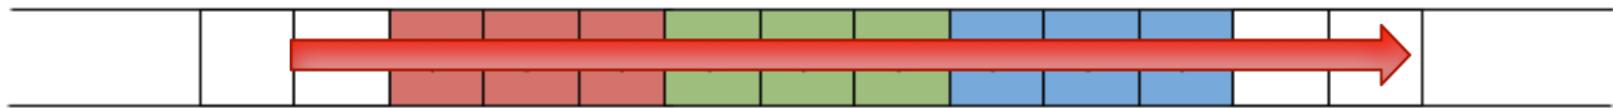
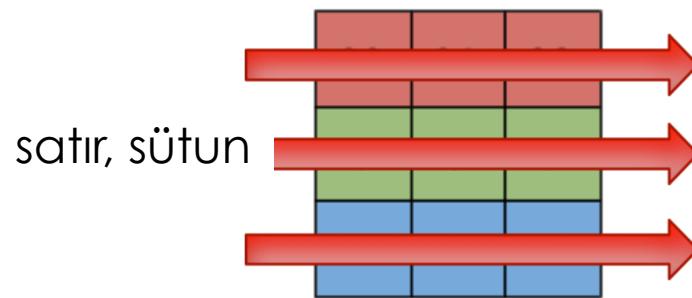


2d verilerin hafıza yapısı

Matrislerin sütunlarının toplamı



Matrislerin sütunlarının toplamı



Matrislerin sütunlarının toplamı

```
for (i = 0; i < N; i++) {  
    for (j = 0; j < N; j++) {  
        column_sum[i] += b[j][i];  
    }  
}
```

2 saniye

```
for (j = 0; j < N; j++) {  
    for (i = 0; i < N; i++) {  
        column_sum[i] += b[j][i];  
    }  
}
```

0.1 saniye

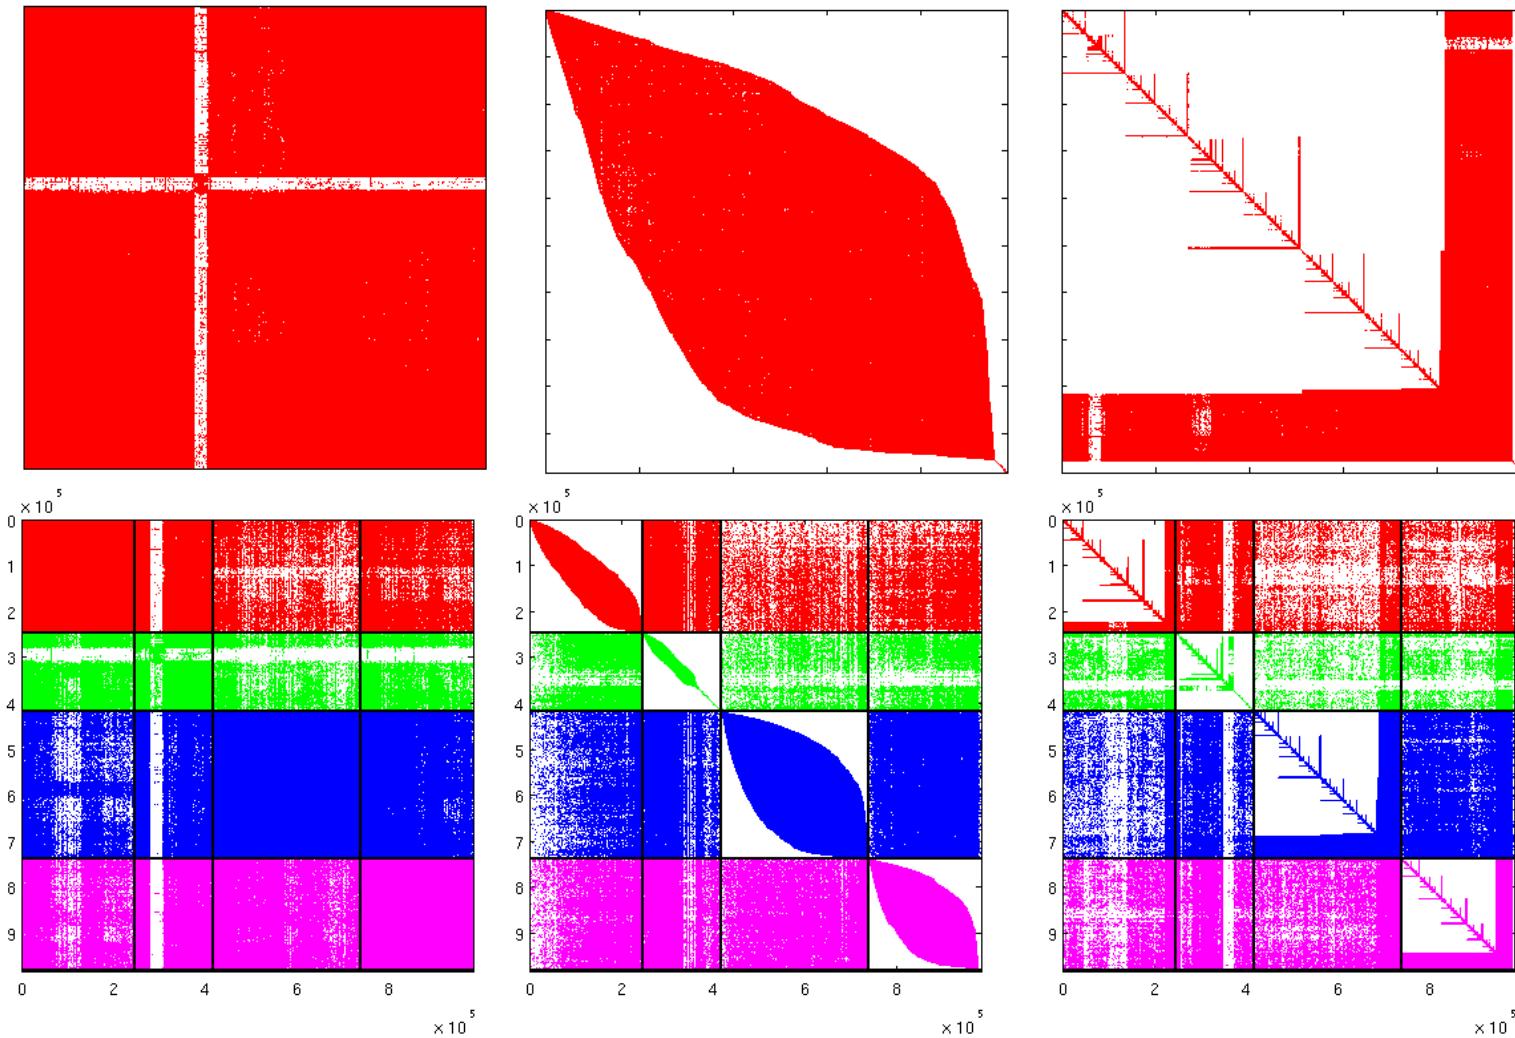
20x daha hızlı

Matris Çarpımı

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 2 & 0 \\ 1 & 2 \end{bmatrix} = \begin{bmatrix} 4 & 4 \\ 10 & 8 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 0 \\ 1 & 2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 2 & 4 \\ 7 & 10 \end{bmatrix}$$

Matris sıralama



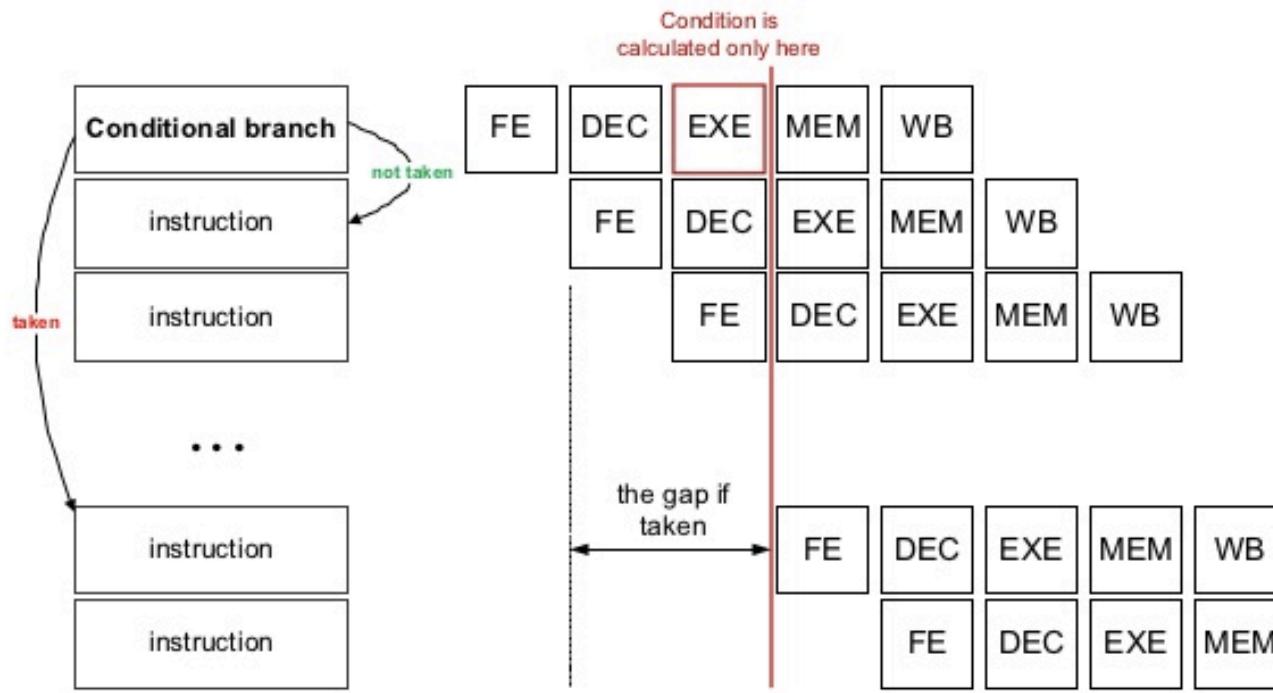
Komut seviyesinde koşutluk

Program to evaluate $X = (A + B) * (C + D)$:

ADD	R1, A, B	$\quad /* \quad R1 \leftarrow M[A] + M[B]$	*/
ADD	R2, C, D	$\quad /* \quad R2 \leftarrow M[C] + M[D]$	*/
MUL	X, R1, R2	$\quad /* \quad M[X] \leftarrow R1 * R2$	*/

Komutlar: Düzenleme, yeniden sıralama vb.

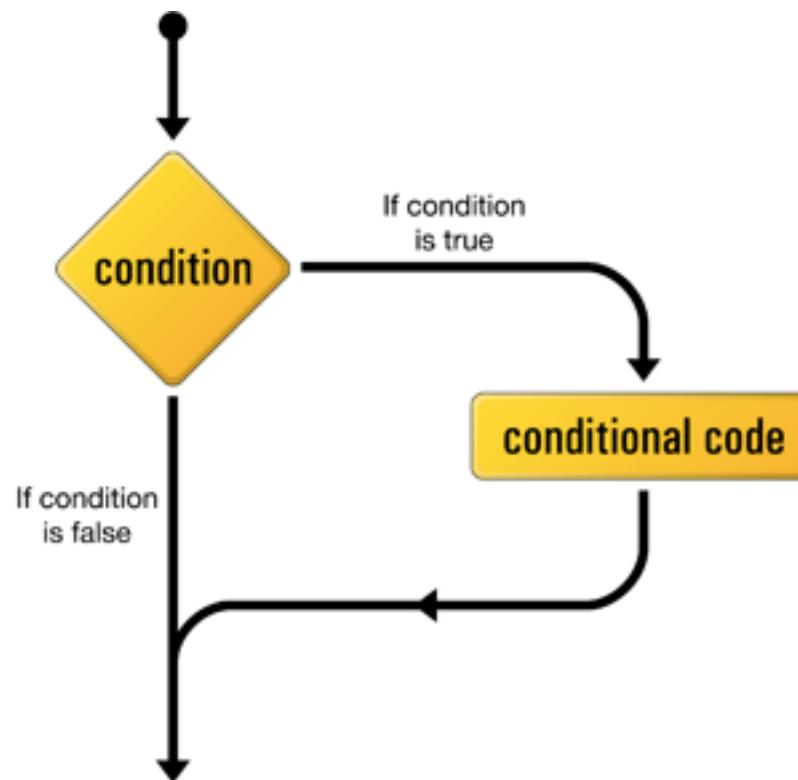
- Her zaman kodladığınız gibi çalışmıyor



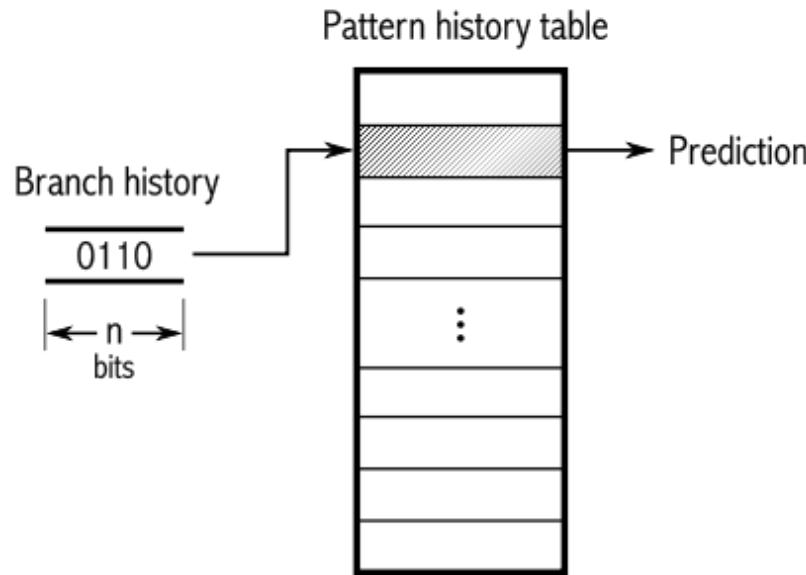
Pipelining

Komutlar: Düzenleme, yeniden sıralama vb.

- Ne çalıştırırmak istediğimizi bilmiyorsak



Komutlar: Düzenleme, yeniden sıralama vb.



Dal tahmincisi

Komutlar: Düzenleme, yeniden sıralama vb.

```
for (unsigned c = 0; c < arraySize; ++c) {
    data[c] = std::rand() % 256;
}

long long sum = 0;
for (unsigned i = 0; i < 100000; ++i) {
    for (unsigned c = 0; c < arraySize; ++c) {
        if (data[c] >= 128)
            sum += data[c];
    }
}
```

13.2 saniye, 1.5 B yanlış tahmin

Komutlar: Düzenleme, yeniden sıralama vb.

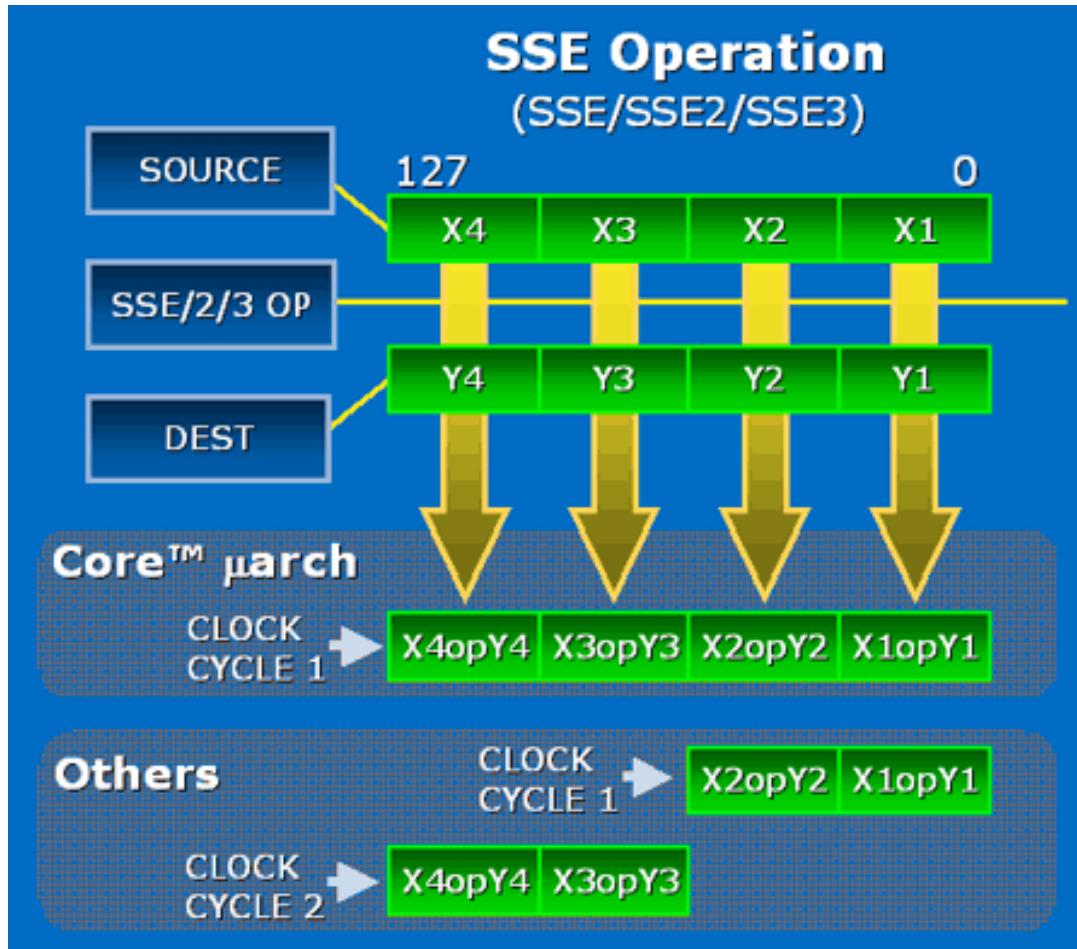
```
for (unsigned c = 0; c < arraySize; ++c) {
    data[c] = std::rand() % 256;
}

std::sort(data.begin(), data.end());

long long sum = 0;
for (unsigned i = 0; i < 100000; ++i) {
    for (unsigned c = 0; c < arraySize; ++c) {
        if (data[c] >= 128)
            sum += data[c];
    }
}
```

2.1 saniye, 364 K yanlış tahmin

Komut seviyesinde koşutluk



Komut seviyesinde koşutluk: matris-vektör çarpımı örneği

Matris-vektör çarpımı

Hafıza duvarını aşmak

- Bu değil **for** (i = 0; i < N; i++) {
 sum += a[i] * b[i];
 }
- Belki de bu **for** (i = 0; i < N; i++) {
 fetch (&a[i+1]);
 fetch (&b[i+1]);
 sum += a[i] * b[i];
 }

Hafıza duvarını aşmak

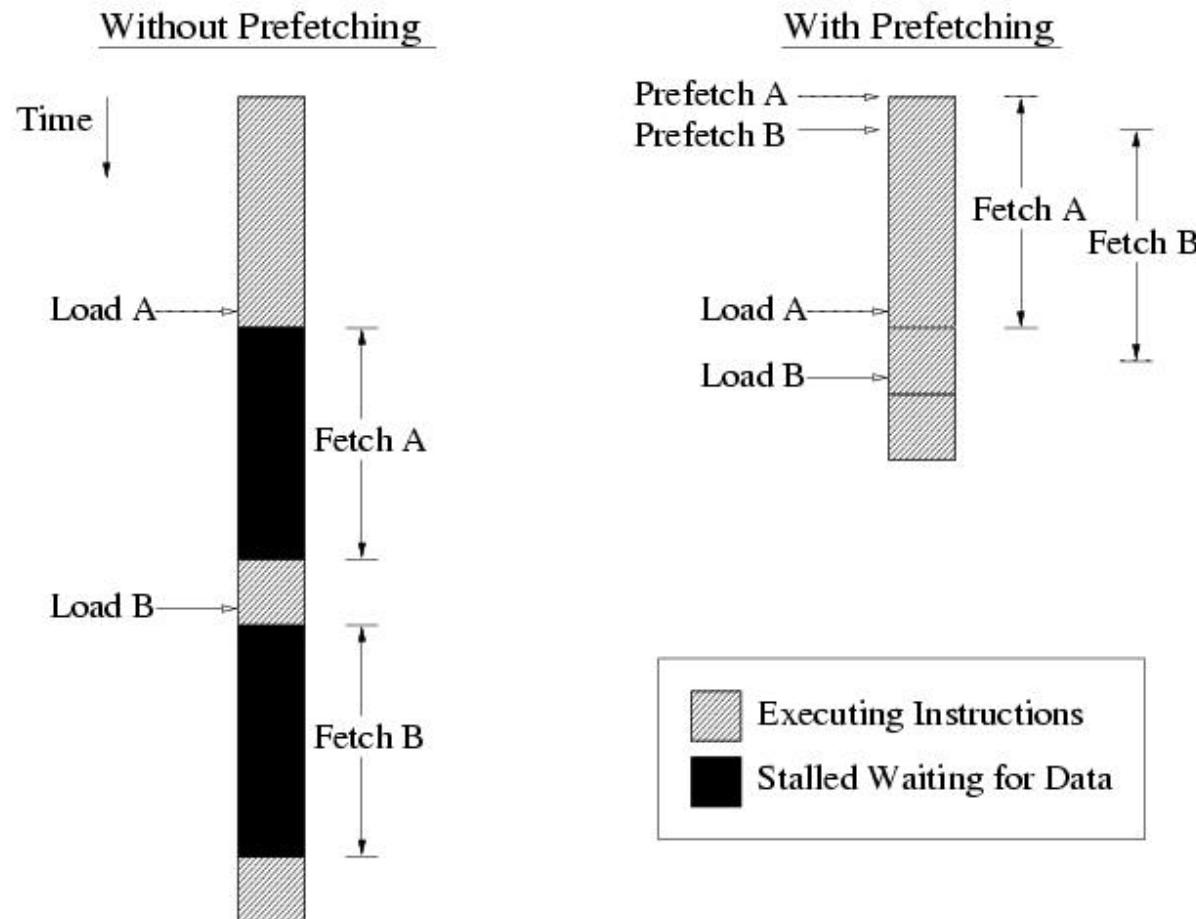
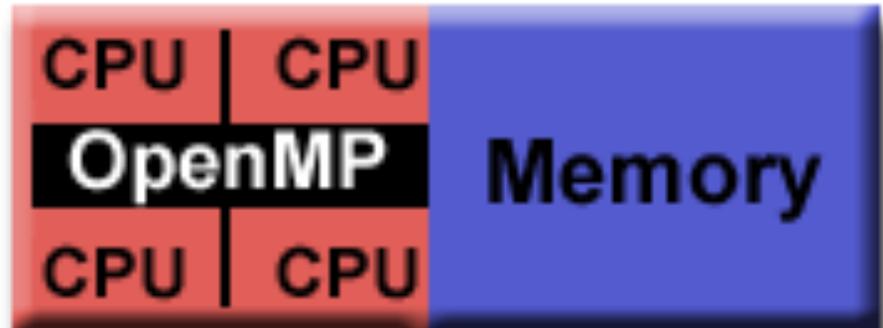


Figure 1.4: Illustration of how prefetching tolerates memory latency.

Bir işlemciden daha fazlası



Bugünün işlemcileri

- Intel E7-8890v3
 - ❑ Frekans: 2.2 GHz
 - ❑ 14nm
 - ❑ Çekirdek: 24
 - ❑ Hafıza: 3.07TB / soket
 - ❑ Hafıza veriyolu: 85 GBs
 - ❑ 60 MB önbellek

Paralel (Koşut) İşlemler

- Aynı anda birden fazla işlemin işlemci tarafından gerçekleştirilmesi.

Problem

Pr
ob
I
em

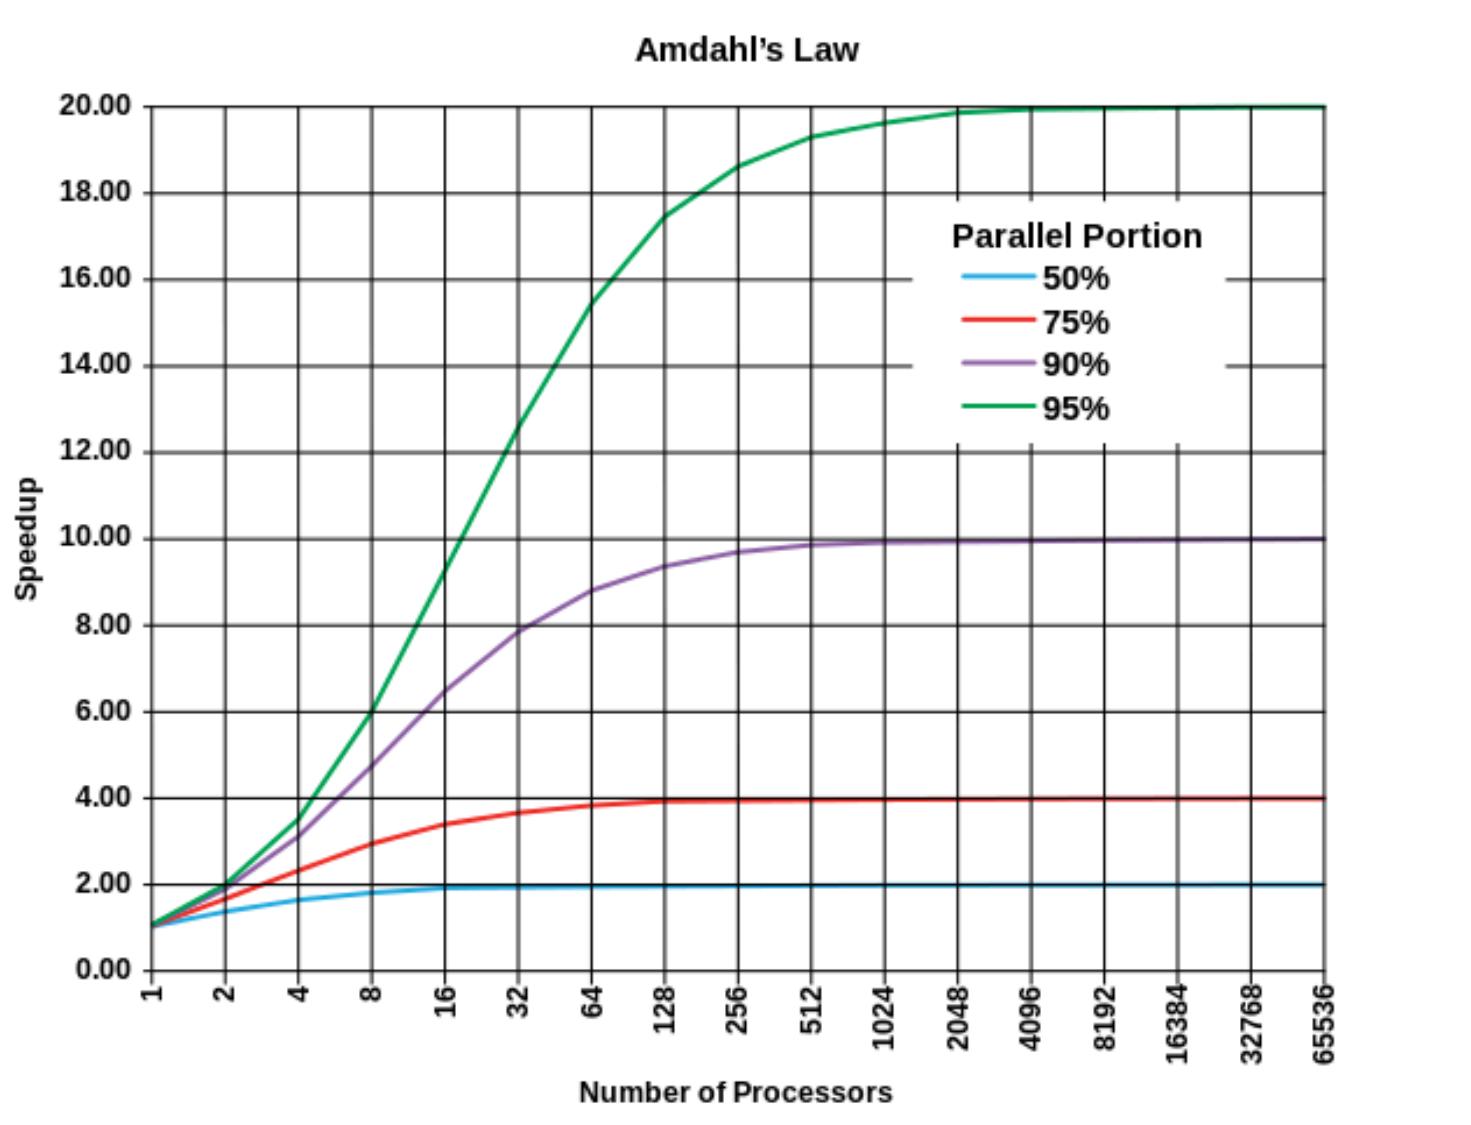
Birden fazla işlemci ve hızlanma

Bir problemi tek işlemci ile $T(1)$ zamanda çözelim.

Aynı problem p işlemci ile $T(p)$ zamanında çözülsün.

Hızlanma oranı $T(1) / T(p)$ 'dır.

Amdahl Kuralı



Aykırı bir problem

1	2	3	4
5	6	7	8
9	10	7	11
13	14	15	12

(a)

1	2	3	4
5	6	7	8
9	10	11	11
13	14	15	12

(b)

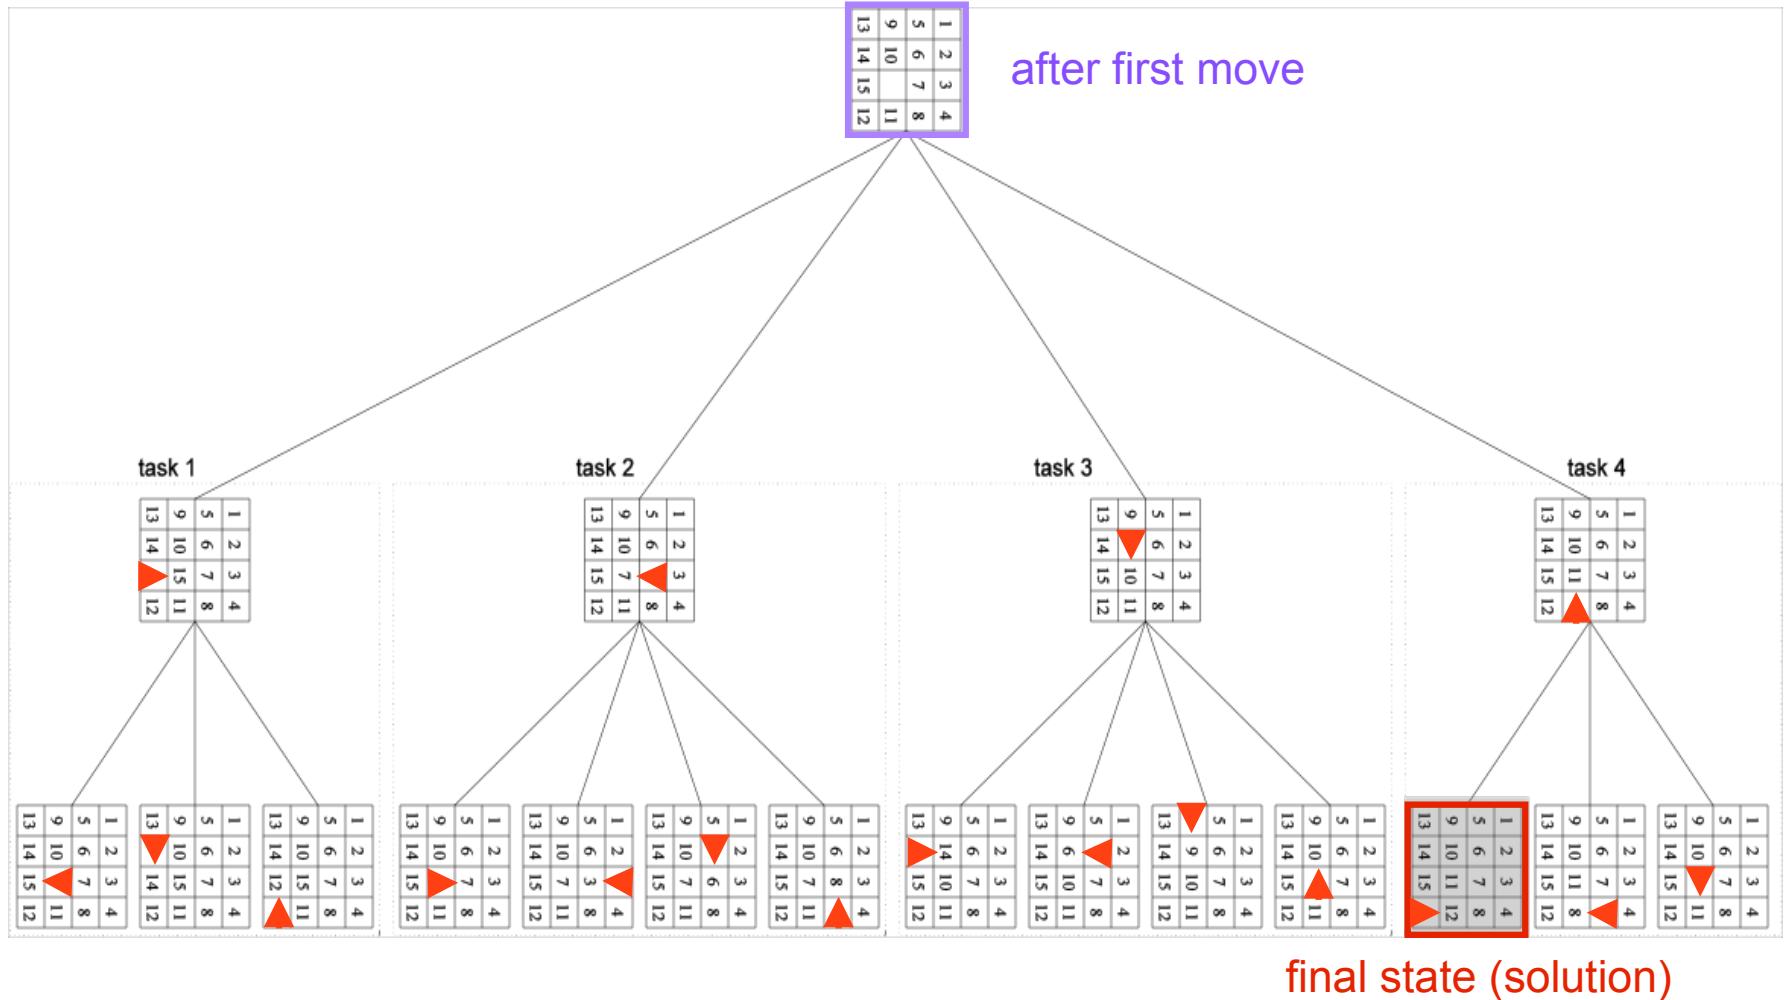
1	2	3	4
5	6	7	8
9	10	11	11
13	14	15	12

(c)

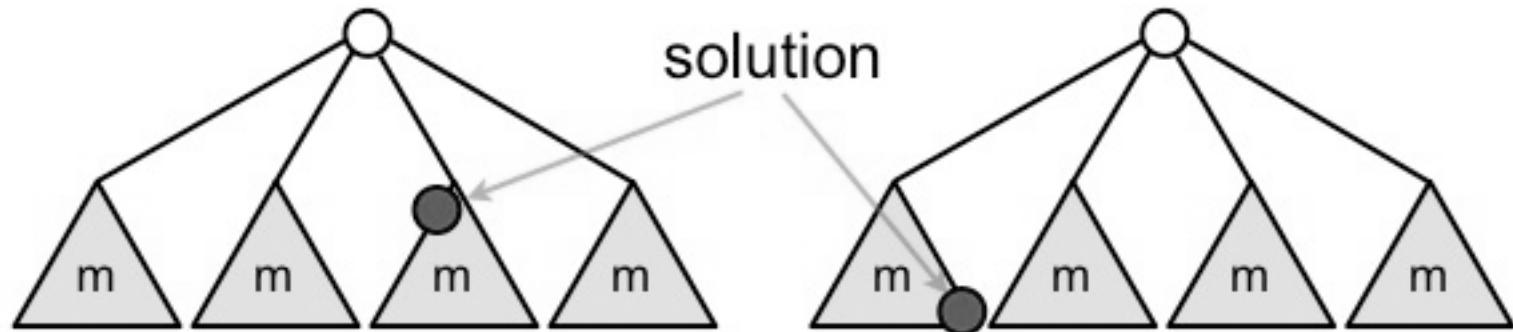
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

Aykırı bir problem



Aykırı bir problem



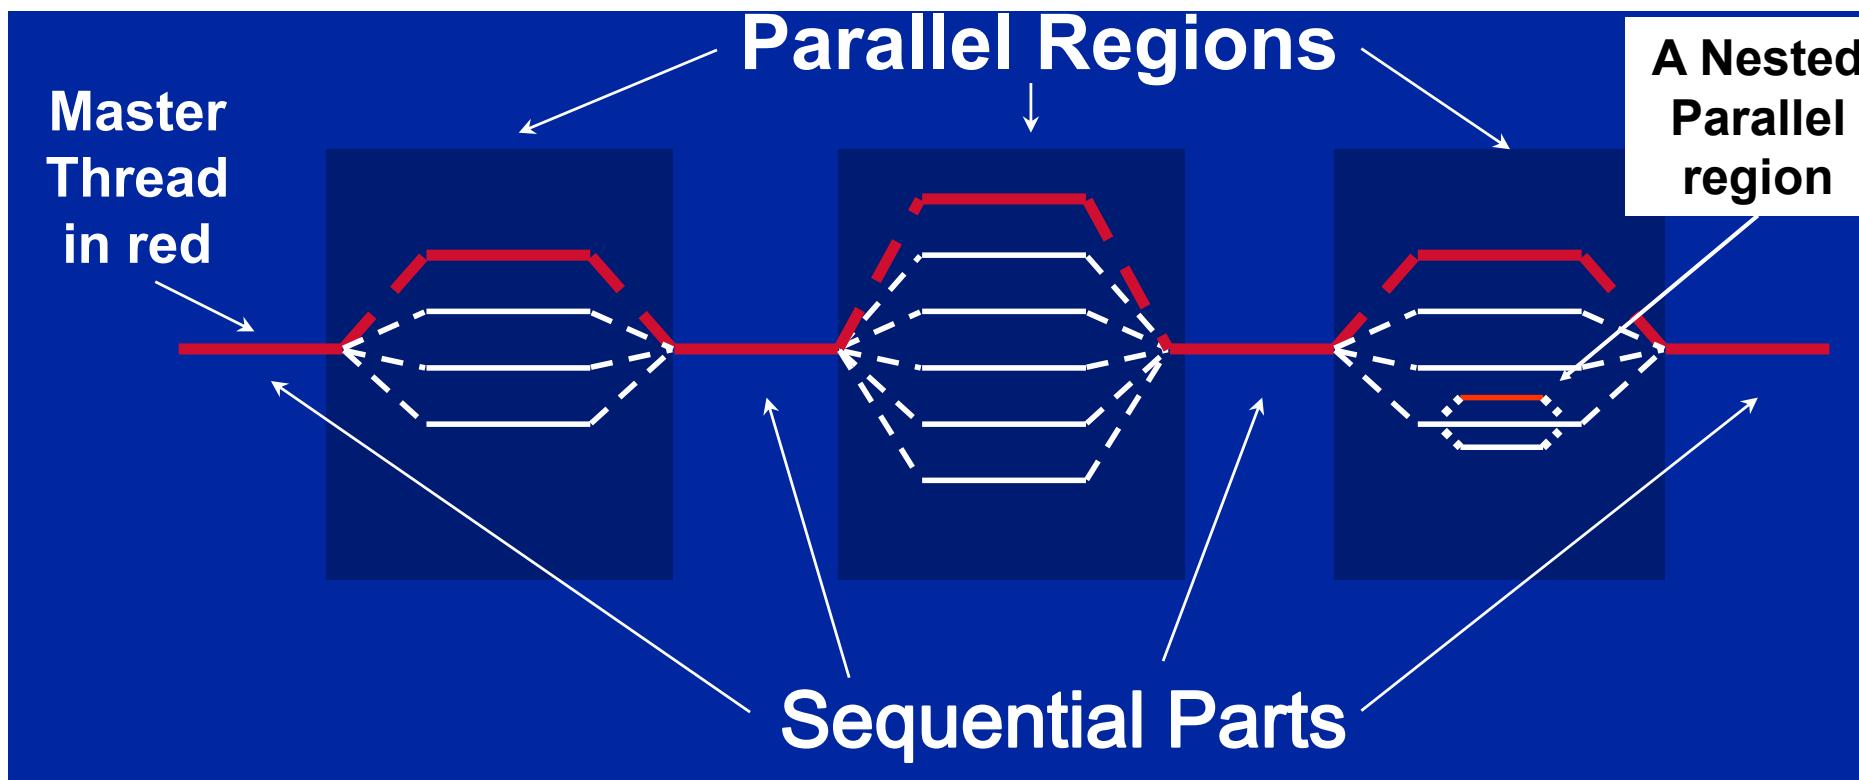
$$\text{total serial work} = 2m + 1$$

$$\text{total parallel work} = 4$$

$$\text{total serial work} = m$$

$$\text{total parallel work} = 4m$$

Çok çekirdekli işlemcileri programlama: OpenMP



Çok çekirdekli işlemcileri programlama: OpenMP

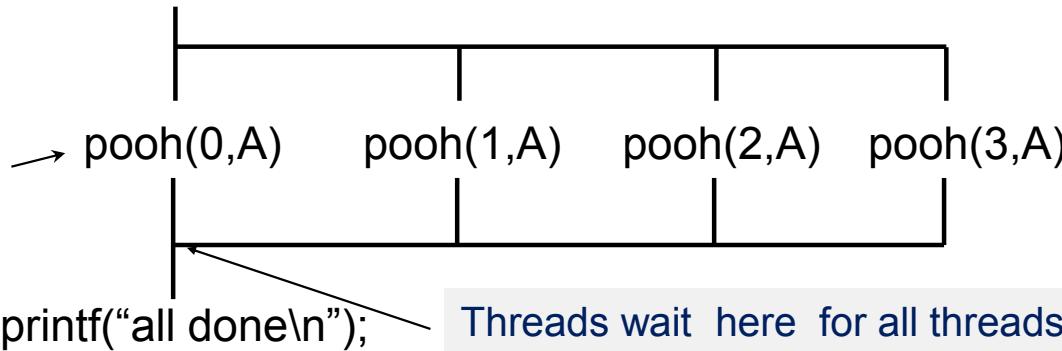
- Each thread executes the same code redundantly.

```
double A[1000];
```

```
omp_set_num_threads(4)
```

A single copy of A is shared between all threads.

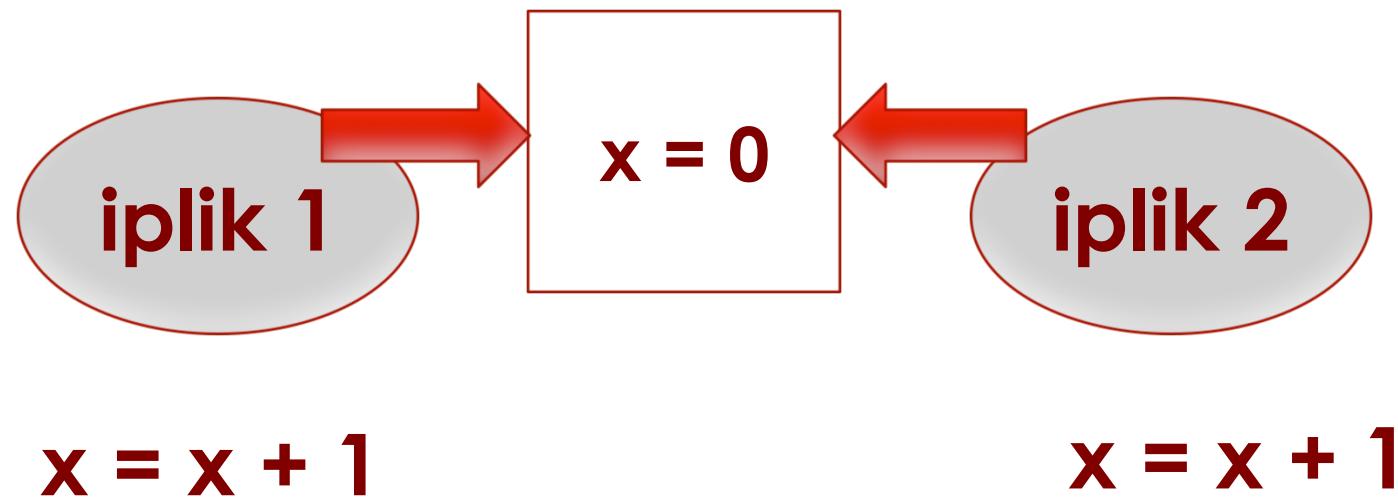
```
double A[1000];
#pragma omp parallel num_threads(4)
{
    int ID = omp_get_thread_num();
    pooh(ID, A);
}
printf("all done\n");
```



Threads wait here for all threads to finish before proceeding (i.e. a *barrier*)

Çok çekirdekli işlemcileri programlama: OpenMP

PARALEL PROGRAMLAMA 101

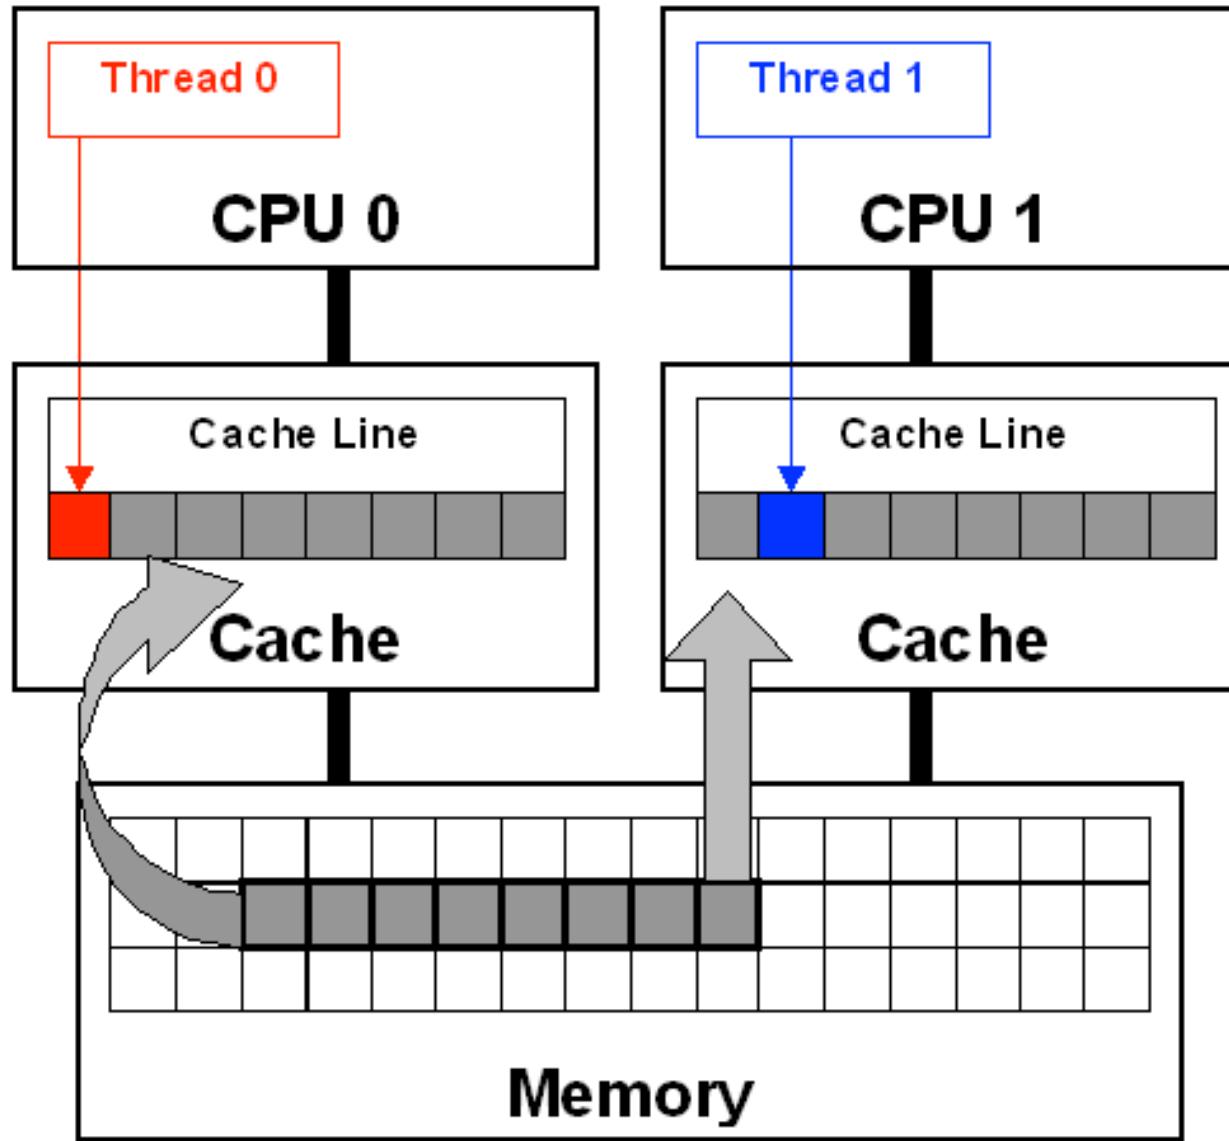


OpenMP: vektör çarpımı

```
double sum = 0.0;
double sum_local[NUM_THREADS];
#pragma omp parallel num_threads(NUM_THREADS)
{
    int me = omp_get_thread_num();
    sum_local[me] = 0.0;
    #pragma omp for
    for (i = 0; i < N; i++)
        sum_local[me] += x[i] * y[i];
    #pragma omp atomic
    sum += sum_local[me];
}
```

OpenMP

51



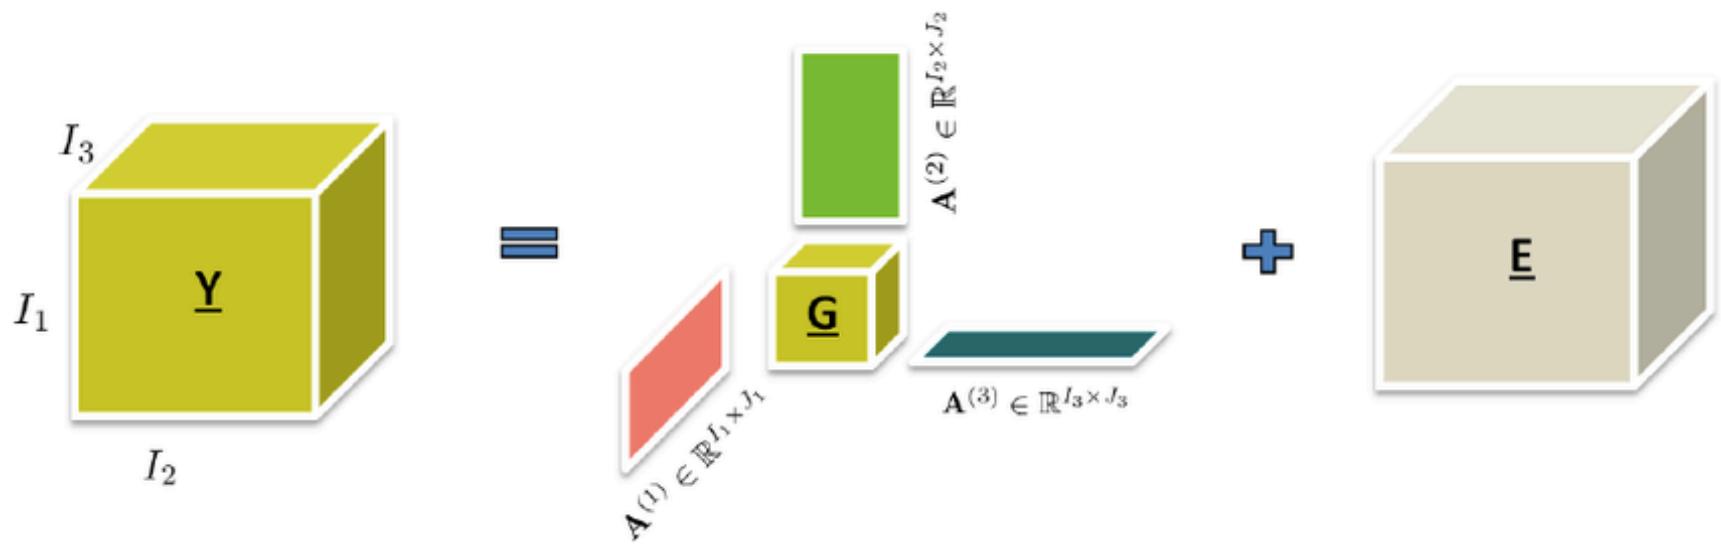
OpenMP: vektör çarpımı

```
double sum = 0.0;
double sum_local[NUM_THREADS][16];
#pragma omp parallel num_threads(NUM_THREADS)
{
    int me = omp_get_thread_num();
    sum_local[me][0] = 0.0;
    #pragma omp for
    for (i = 0; i < N; i++)
        sum_local[me] += x[i] * y[i];
    #pragma omp atomic
    sum += sum_local[me][0];
}
```

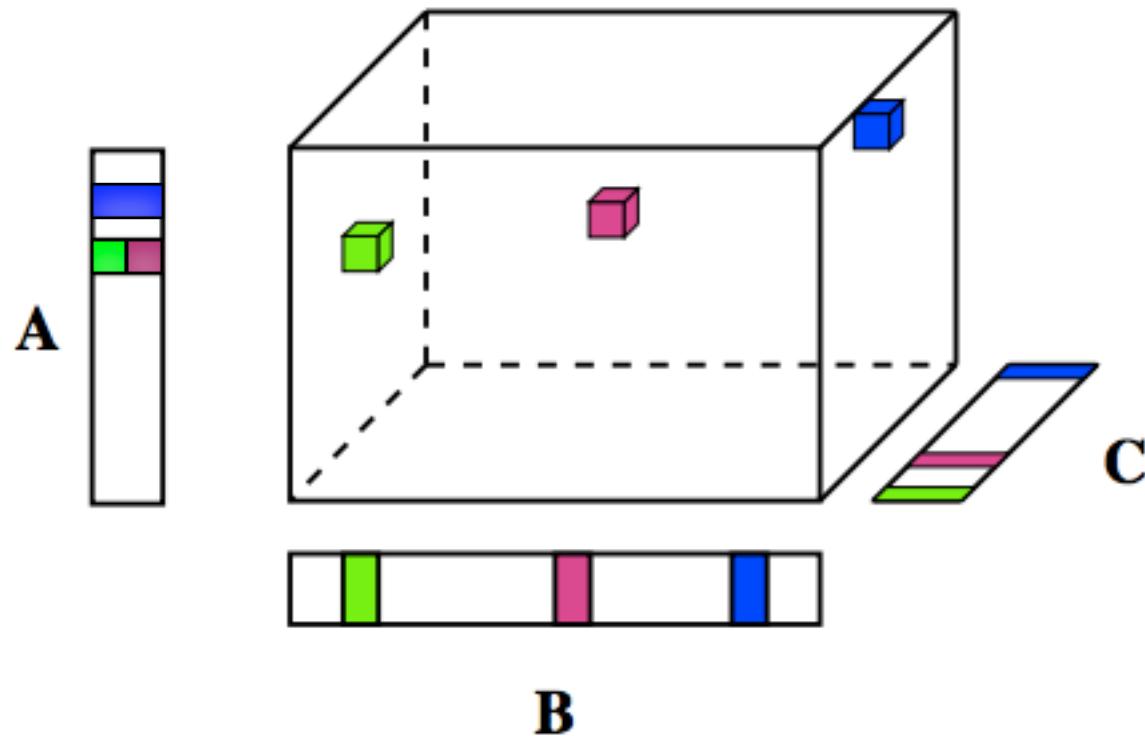
**Threads wait
their turn –
only one at a
time calls
consume()**

```
float res;  
#pragma omp parallel  
{   float B;  int i, id, nthrds;  
    id = omp_get_thread_num();  
    nthrds = omp_get_num_threads();  
    for(i=id;i<niters;i+=nthrds){  
        B = big_job(i);  
        #pragma omp critical  
        res += consume (B);  
    }  
}
```

Örnek problem: Tensör ayrıştırma



Örnek problem: Koşut tensör ayrıştırma



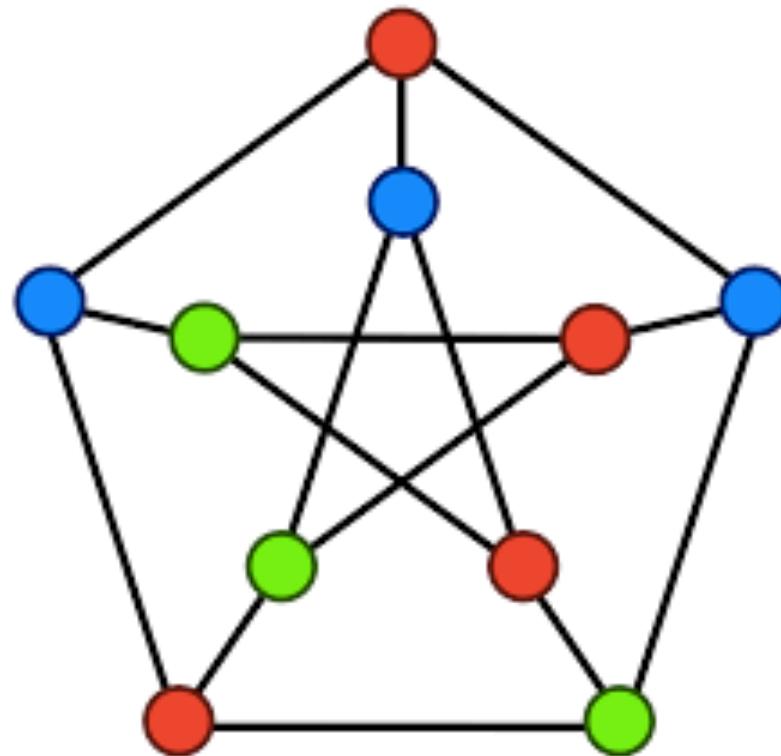
Örnek problem: Koşut tensör ayrıştırma

$$\begin{array}{|c|c|c|c|c|} \hline & a & & b & \\ \hline c & & d & & e \\ \hline & f & & & \\ \hline g & & & h & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \times \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|c|} \hline & a & & b & \\ \hline c & & d & & e \\ \hline & f & & & \\ \hline g & & & h & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \times \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array}$$

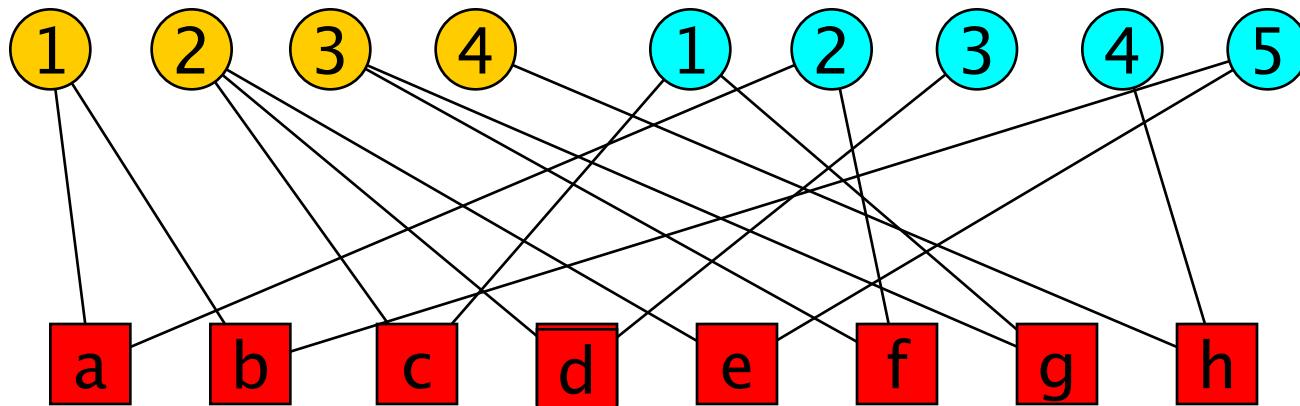
$$\begin{array}{|c|c|c|c|c|} \hline & a & & b & \\ \hline c & & d & & e \\ \hline & f & & & \\ \hline g & & & h & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \times \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array}$$

Çizge boyama problemi



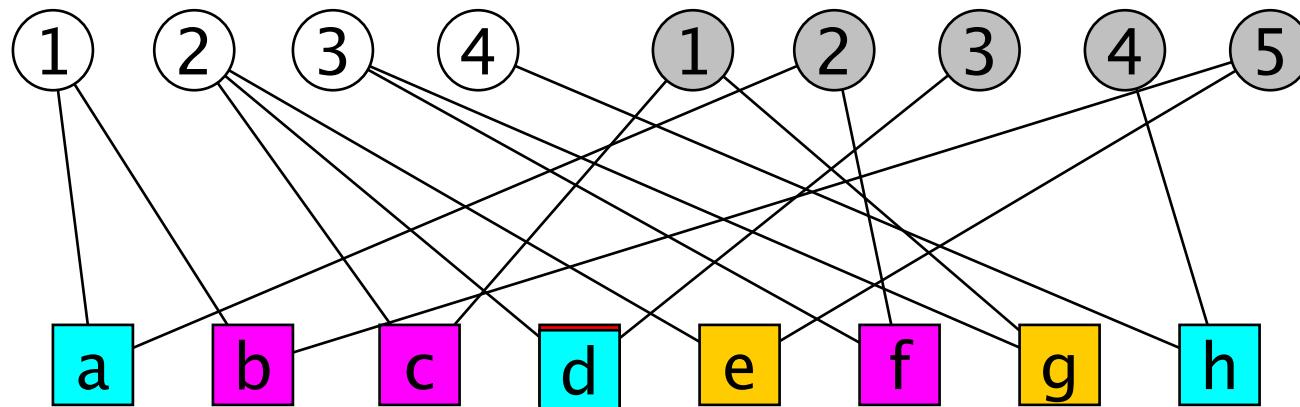
Hiperçizge ile modelleme

$$\begin{array}{|c|c|c|c|} \hline & a & & b \\ \hline c & & d & e \\ \hline & f & & \\ \hline g & & h & \\ \hline \end{array} = \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array} \times \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array}$$



Hiperçizge boyama

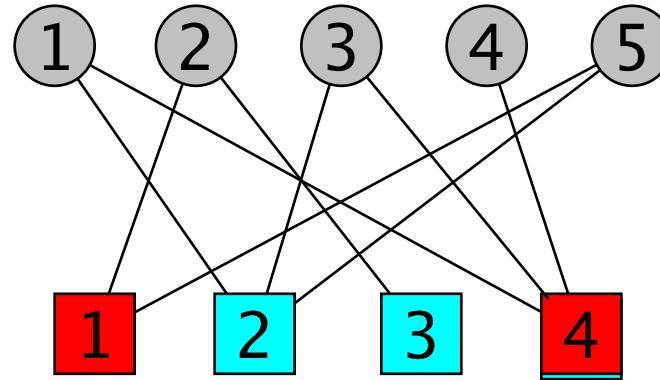
$$\begin{array}{|c|c|c|c|c|} \hline & a & & b & \\ \hline c & & d & & e \\ \hline & f & & & \\ \hline g & & & h & \\ \hline \end{array} = \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array} \times \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array}$$



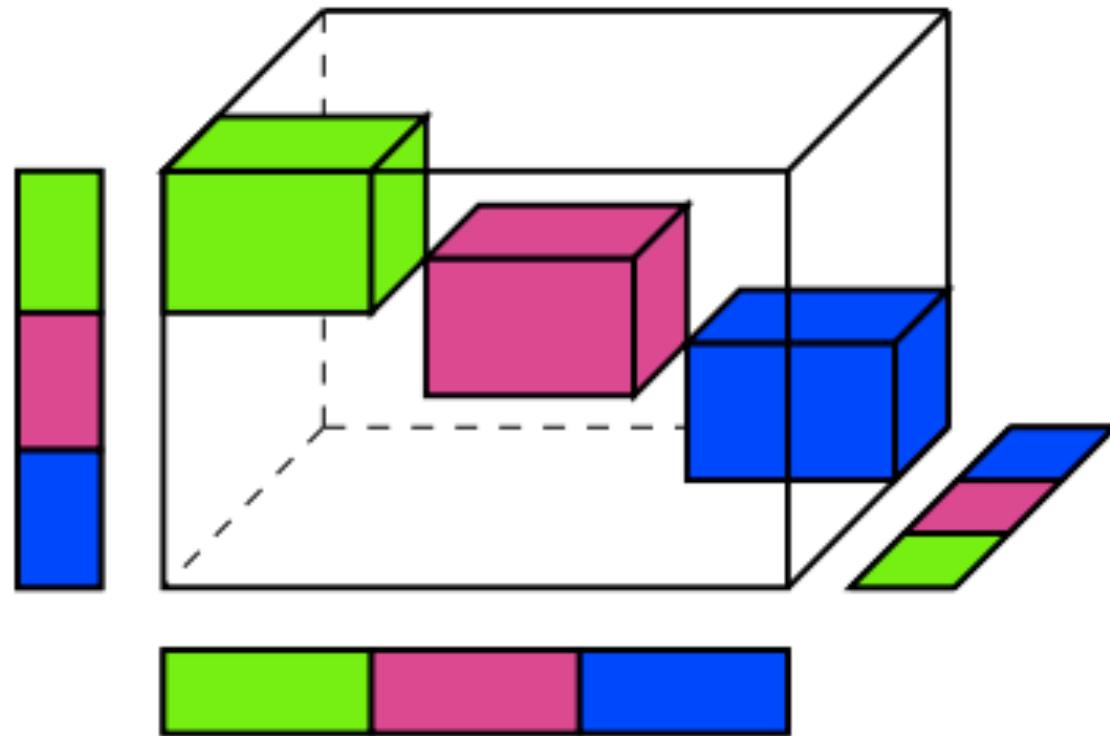
Hiperçizge modelleme/ boyama

60

$$\begin{matrix} & a & & b \\ c & & d & & e \\ & f & & & \\ g & & & h & \end{matrix} = \begin{matrix} & & & & \\ & & & & \\ & & & & \\ & & & & \\ & & & & \end{matrix} \times \begin{matrix} & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \end{matrix}$$



Koşut tensör ayrıştırma

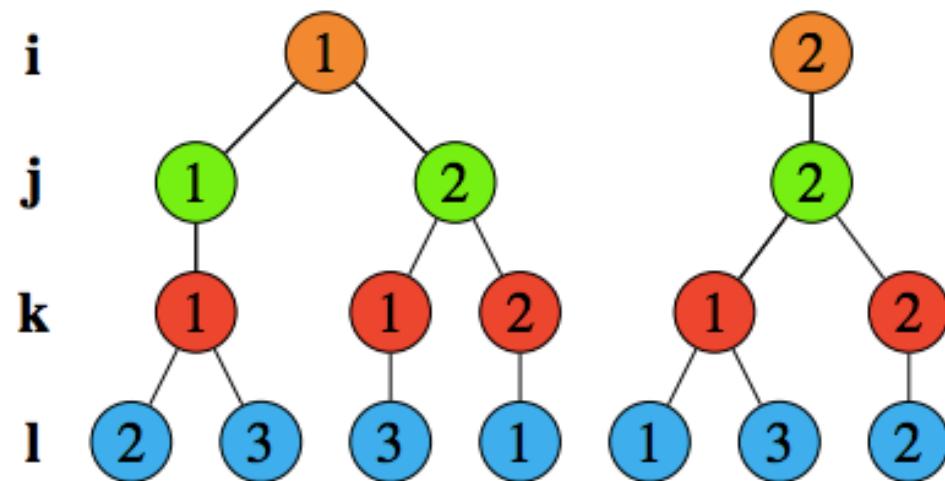


Koşut tensör ayrıştırma

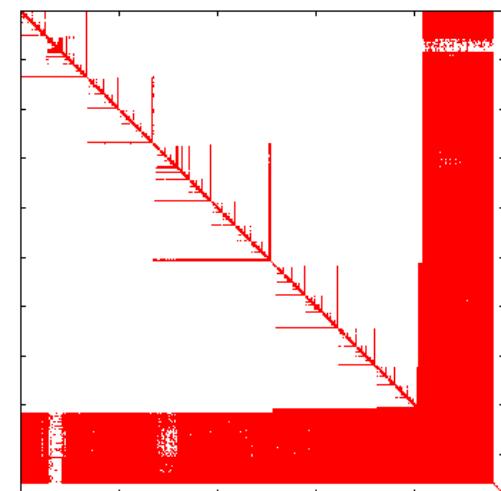
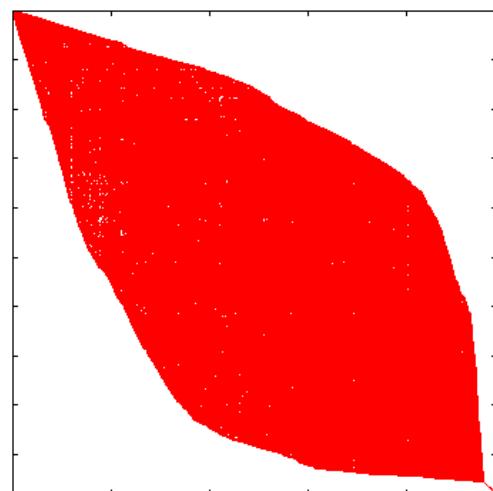
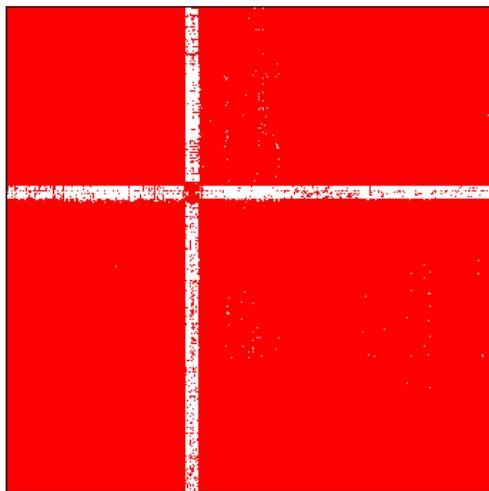
155758	39848
113027	37751
75902	38389
69064	37083
47450	41190
39880	39587
20634	41029
22481	39657
15741	37098
14061	39542
16174	37720
8107	39300
6846	40733
12103	39121
11866	40461
3680	38434
INT 15 0.518887 5.99912	INT 15 0.200038 5.99988
INT 31 1.05262 5.74374	INT 31 0.415625 5.74477
INT 47 1.58465 5.57151	INT 47 0.629195 5.57245
INT 63 2.11765 5.43935	INT 63 0.843154 5.44008

Tensör depolama

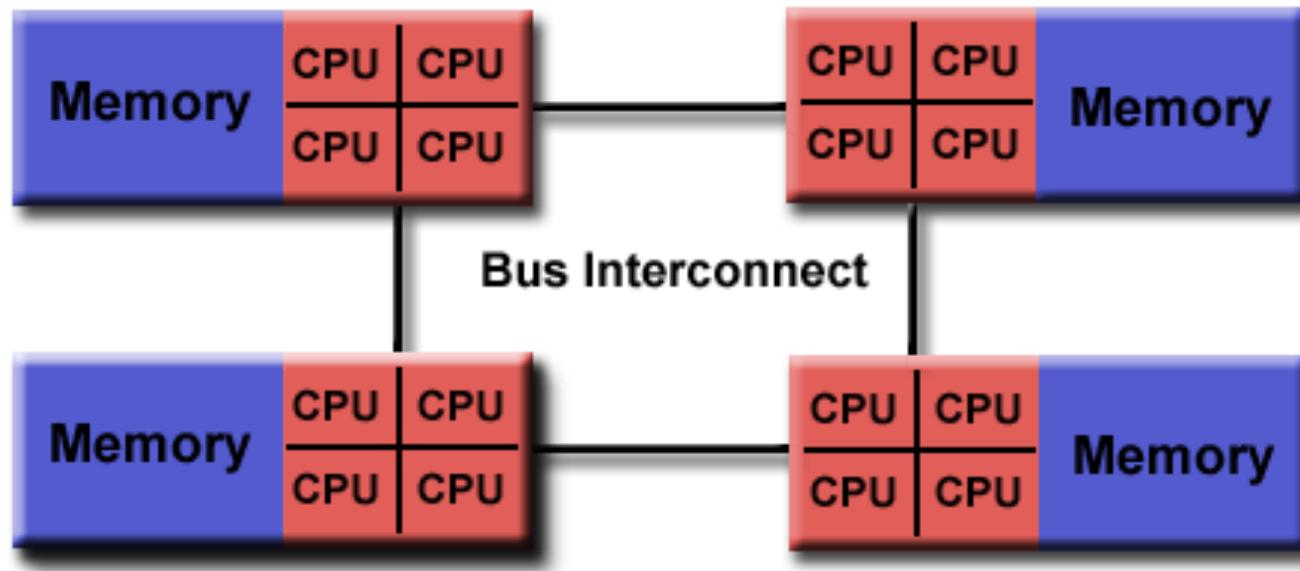
i	j	k	l
1	1	1	2
1	1	1	3
1	2	1	3
1	2	2	1
2	2	1	1
2	2	1	3
2	2	2	2



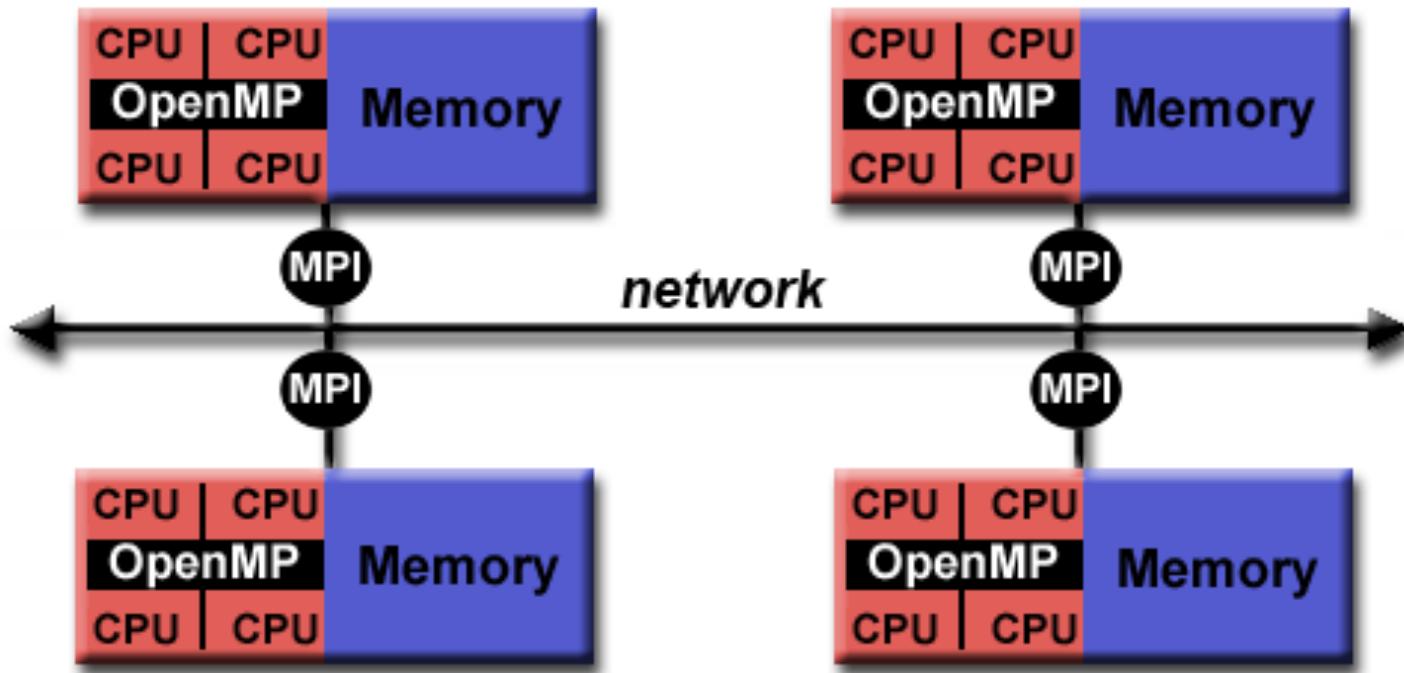
Tensör sıralama?



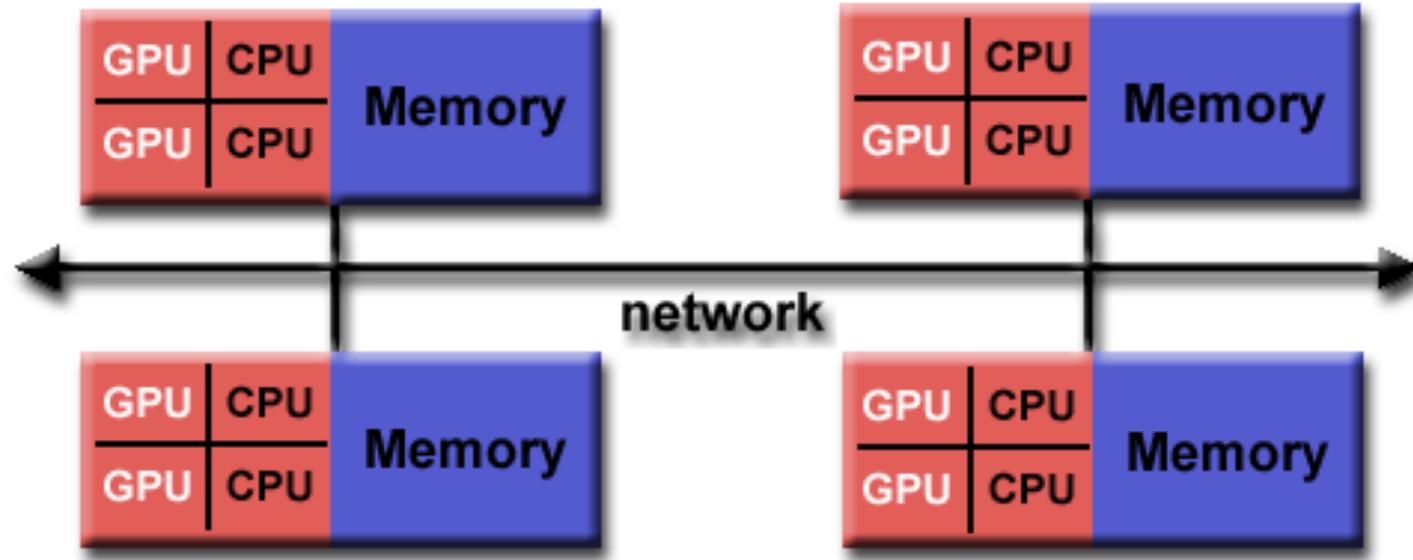
Bir bilgisayar



Bir bilgisayardan daha fazlası



Bir bilgisayardan daha fazlası



Bir bilgisayardan çok daha fazlası



1,024 işlemci
59.7 gigaflop/s

M-5: LOS ALAMOS NATIONAL LAB
No 1. süperbilgisayar, Haziran 1993

Bir bilgisayardan çooooook
daha fazlası



478.2 Teraflop/s

BLUEGENE/L: LAWRENCE LIVERMORE NATIONAL
LABORATORY, No 1. süperbilgisayar, 2004-2007

Neredeyse en süper bilgisayar



33.86 Petaflop/s (san. kuadrilyon işlem)
3,120,000 çekirdek
Intel Xeon Phi

TIANHE-2 (MILKYWAY-2) : NATIONAL UNIVERSITY OF DEFENSE TECHNOLOGY

En süper bilgisayar

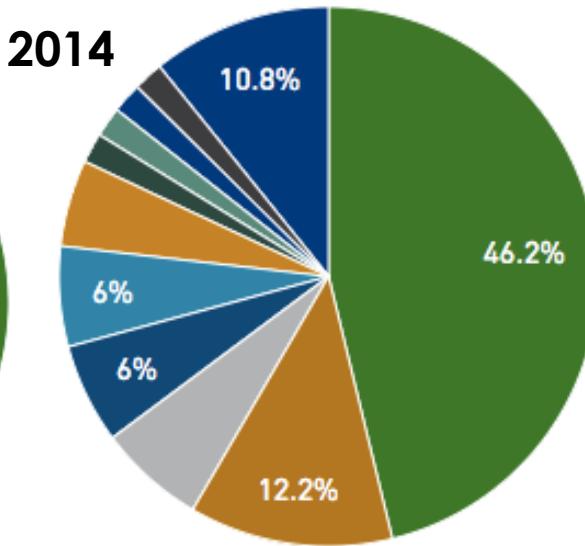
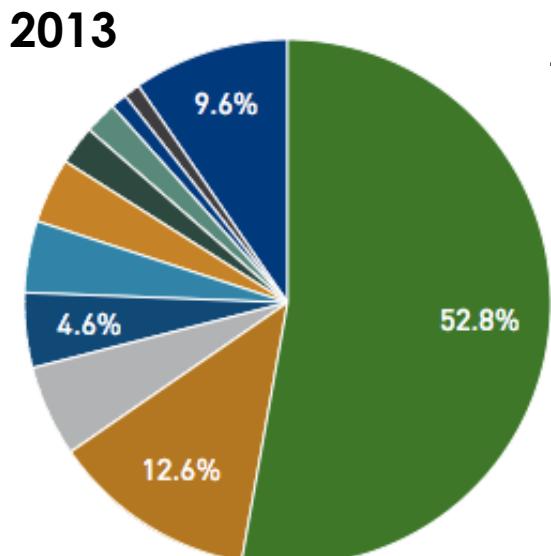


- **SUNWAY TAIHULIGHT – 125 Petaflop/s**
 - NATIONAL SUPERCOMPUTING CENTER IN WUXI
- Araştırma geliştirme amaçlı
 - Meteoroloji,
 - Derin öğrenme,
 - İmalat ve malzeme bilimleri,
 - Veril analizi

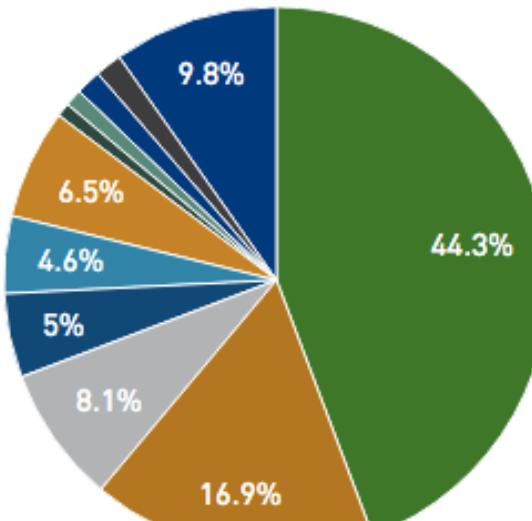
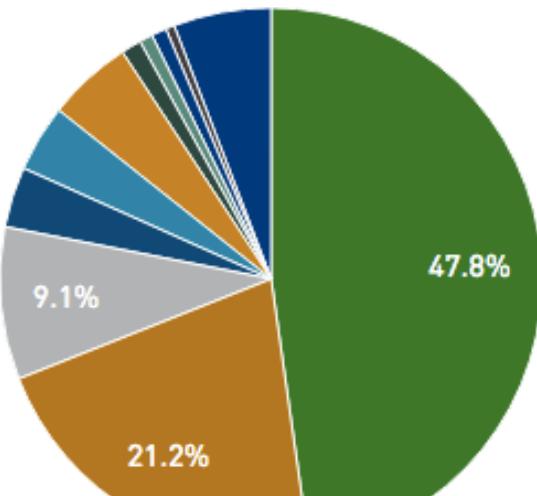
En süper 500 bilgisayar: Ülkeler

72

Number



Performance

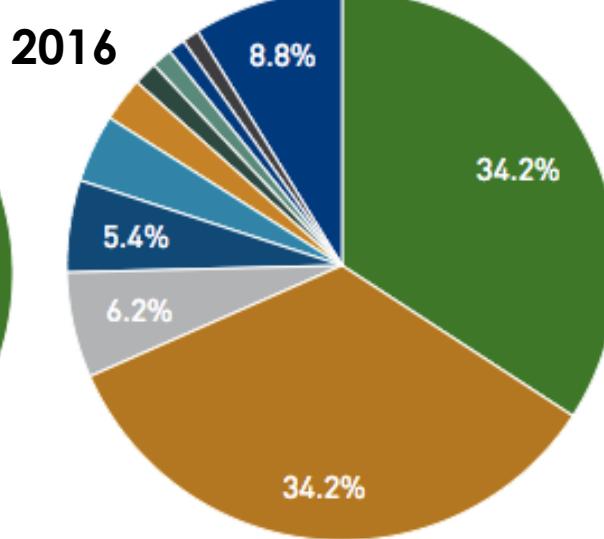
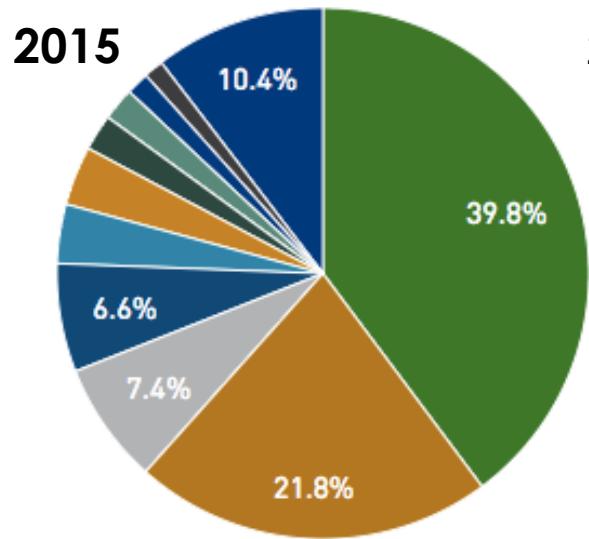


- United States
- China
- Japan
- United Kingdom
- France
- Germany
- India
- Canada
- Russia
- Sweden
- Others

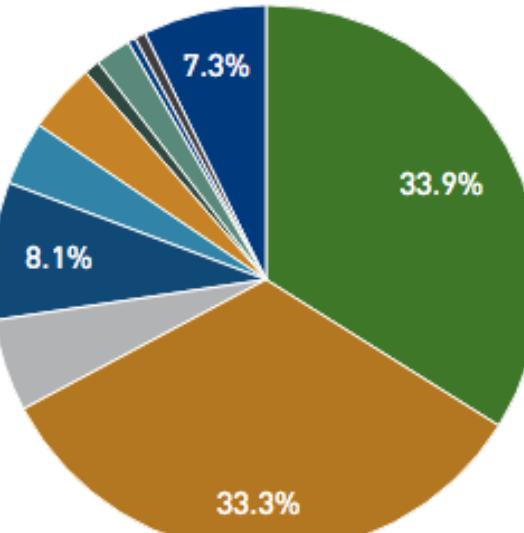
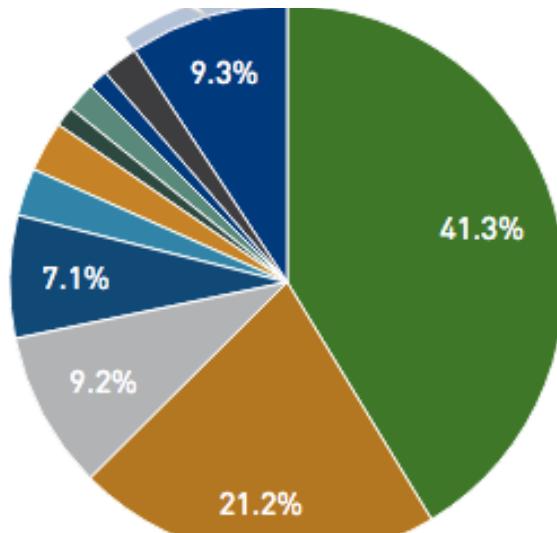
En süper 500 bilgisayar: Ülkeler

73

Number



Performance

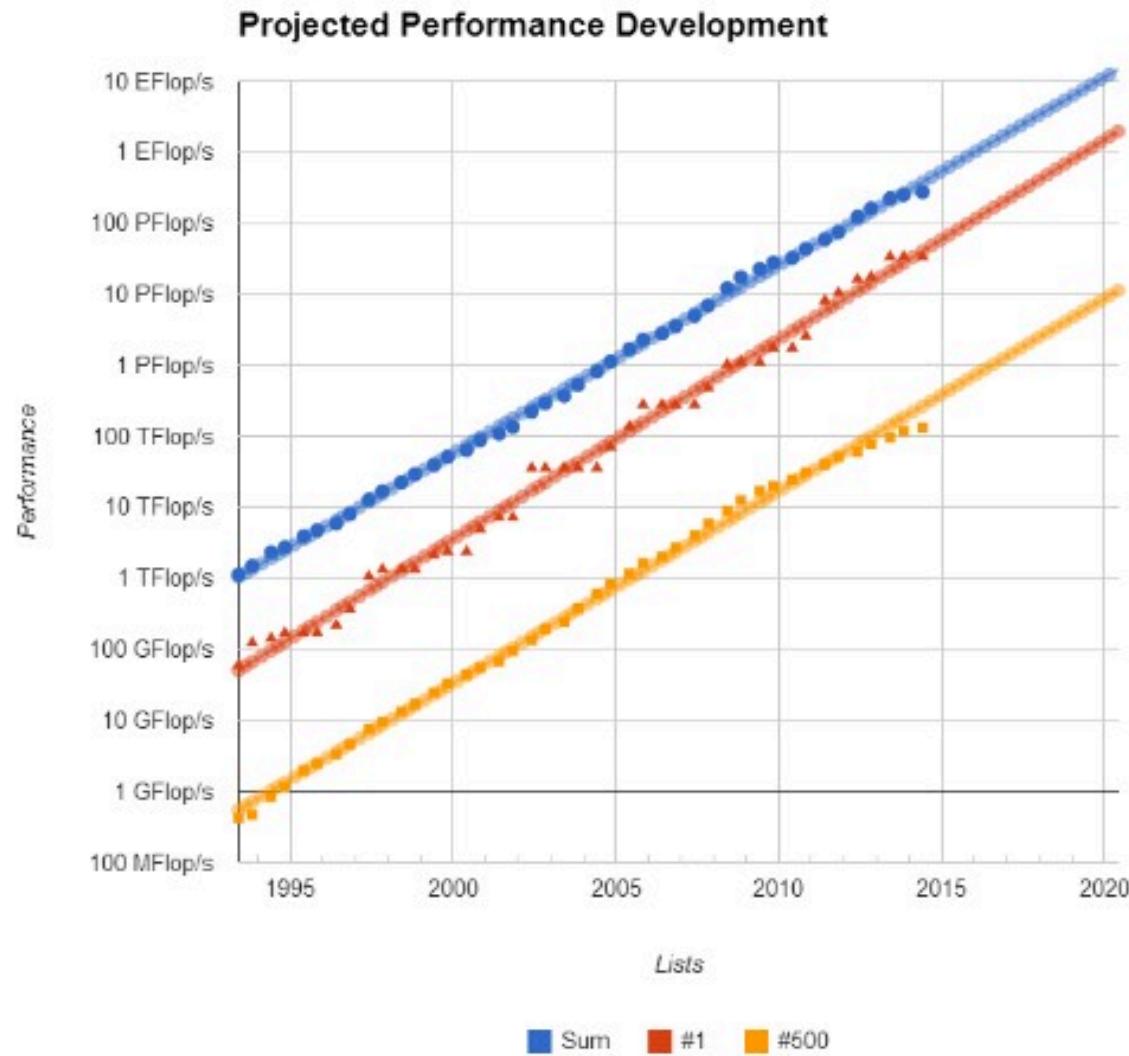


- United States
- China
- Japan
- United Kingdom
- France
- Germany
- India
- Canada
- Russia
- Sweden
- Others

HPC Yatırımları

Country	HPC Strategy/Program and Description	Investment Level
United States	National Strategic Computing Initiative (NSCI)	@\$320 million/year
China	13th Five-Year Development Plan (Develop Multiple Exascale Systems)	\$200 million/year (for next five years)
Japan	Flagship2020 Program	@\$200 million/year (for next five years)
European Union	ExaNeSt; PRACE; ETP4HPC	@\$1.1 in billion total allocated through 2020 (annual allocations N/A)
India	National Supercomputing Mission	\$140 million/year (for five years from 2016-2020)
South Korea	National Supercomputing Act	\$20 million/year (for five years from 2016-2020)
Russia	HPC Focus of Medvedev Modernisation Programme	N/A

Süperbilgisayar trendleri



Exascale: Kilo, Mega, Giga, Tera, Peta, Exa.

- Saniye başına işlem sayısı (FLOPS, TEPS, ...)

- | Terascale: 10^{12} , 1,000,000,000,000

- | Petascale: 10^{15} , 1,000,000,000,000,000

- | Exascale: 10^{18} , 1,000,000,000,000,000,000

Exascale: Kilo, Mega, Giga, Tera, Peta, Exa.

77



Exascale: Kilo, Mega, Giga, Tera, Peta, Exa.

- Exascale hesaplama Yüksek Başarılı Hesaplama alanının şu andaki en önemli hedeflerinden birisidir
- Başlangıç: 2008/USA, 2011/Europe
- Hedef: America, 2021-2023, China önumüzdeki yıl (gerçek?)

HPC'de geçen ay

AI: Deeper Learning with Intel® Omni-Path Architecture

Deep learning is a powerful tool that identifies patterns, extracts meaning from large, diverse datasets, and solves complex problems. However, integrating neural networks into existing compute ...

Sponsored Content by Intel

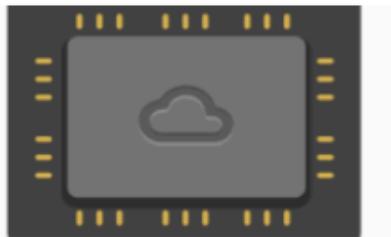
HPC'de geçen ay



September 22, 2017

The 66th HPC User Forum held September 5-7, in Milwaukee, Wisconsin, at the elegant and historic Pfister Hotel, highlighting the 1893 Victorian décor and art of "The Grand Hotel Of The West," contrasted nicely with presentations on the latest trends in modern computing – deep learning, machine learning and AI.

HPC'de geçen ay



Google Cloud Makes Good on Promise to Add Nvidia P100 GPUs

September 21, 2017

Google has taken down the notice on its cloud platform website that says Nvidia Tesla P100s are “coming soon.” That’s because the search giant has announced the beta launch of the high-end P100 Nvidia Tesla GPUs on t [Read more...](#)

By George Leopold

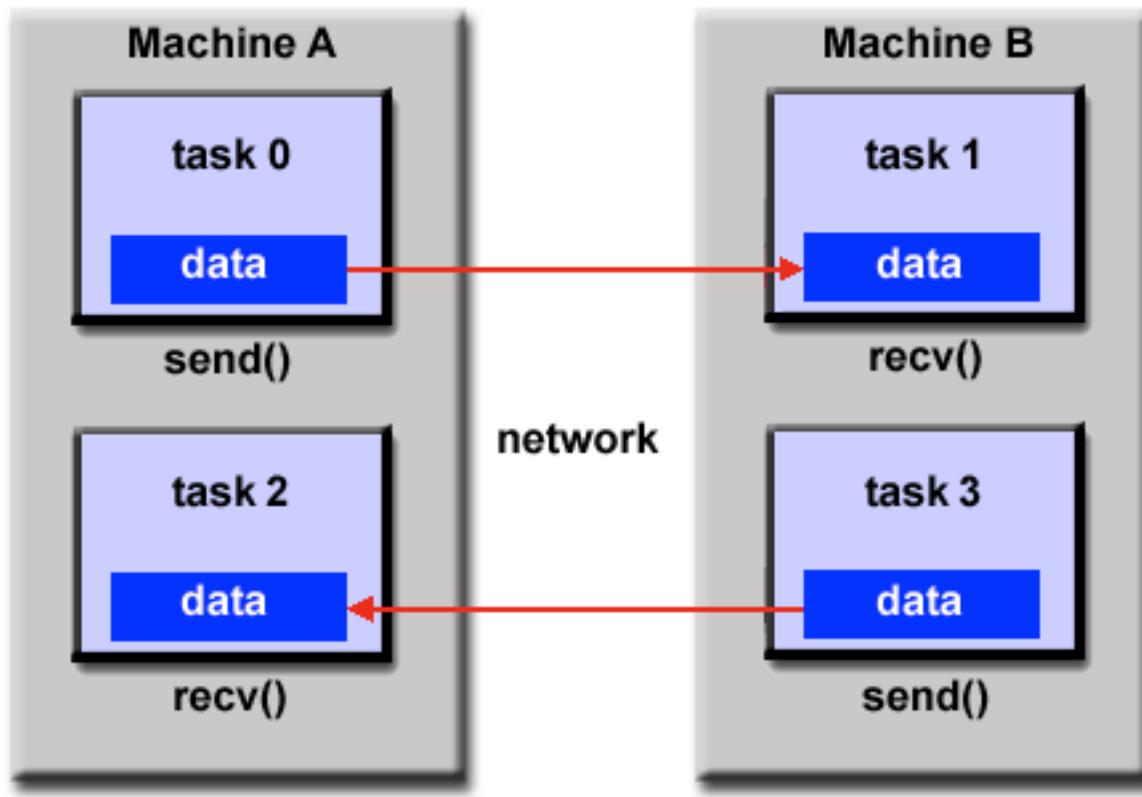


HPC'de geçen ay

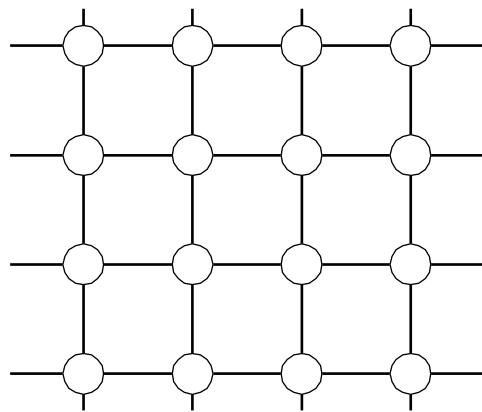
**DeepL Anchors Neural Machine Translator at Verne
Global's HPC-Optimised Data Center**

September 20, 2017

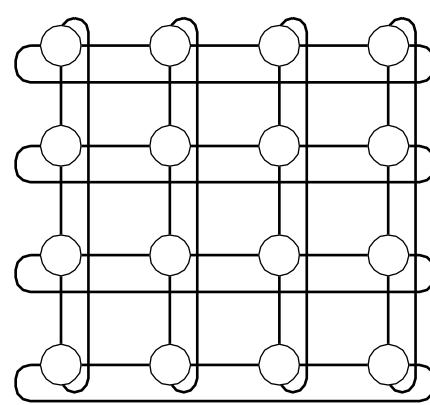
Veri alışverişesi



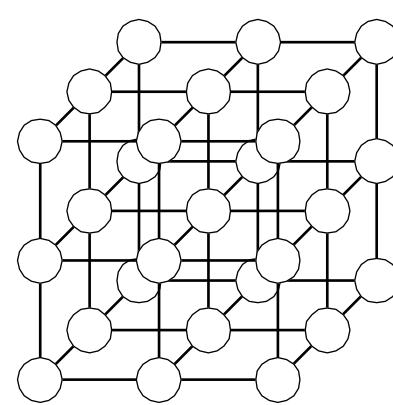
Ağ topolojileri



(a)

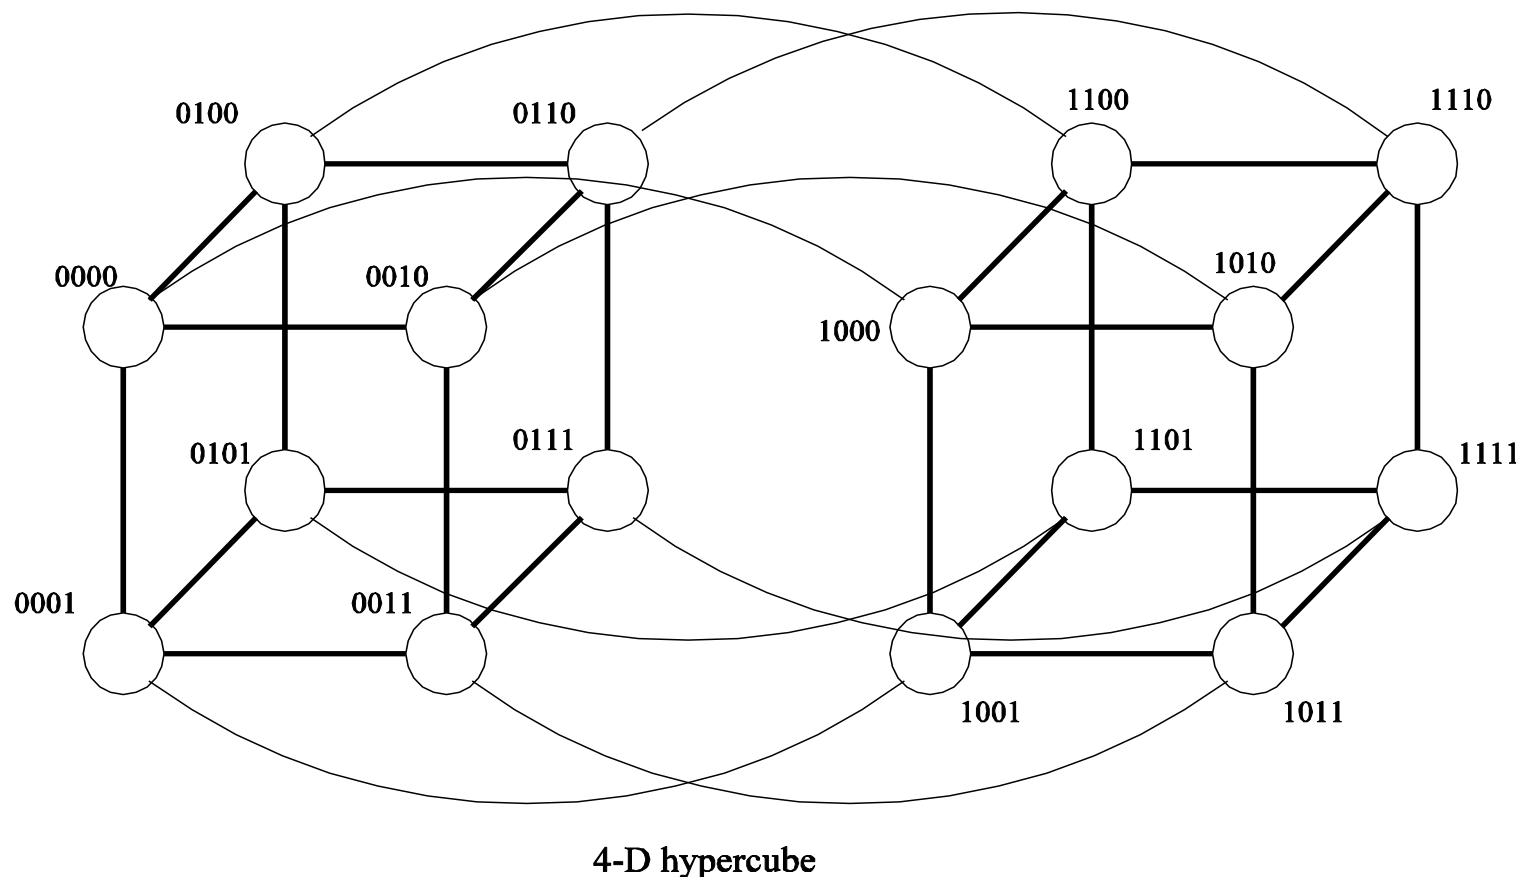


(b)

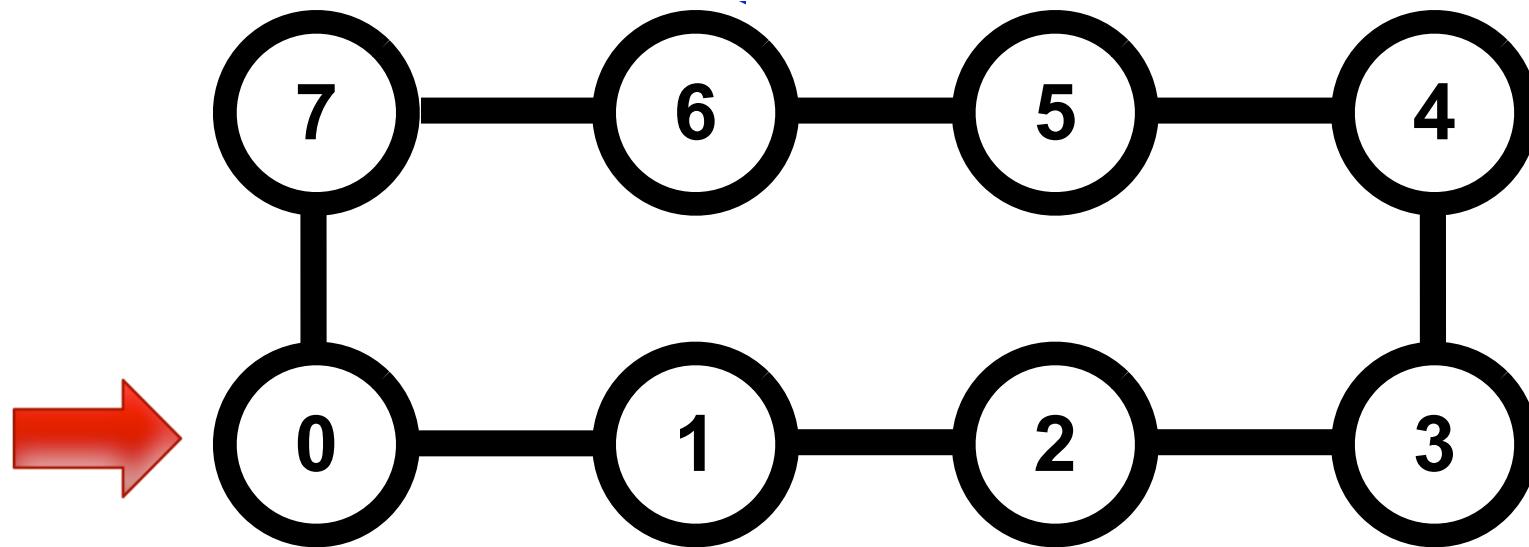


(c)

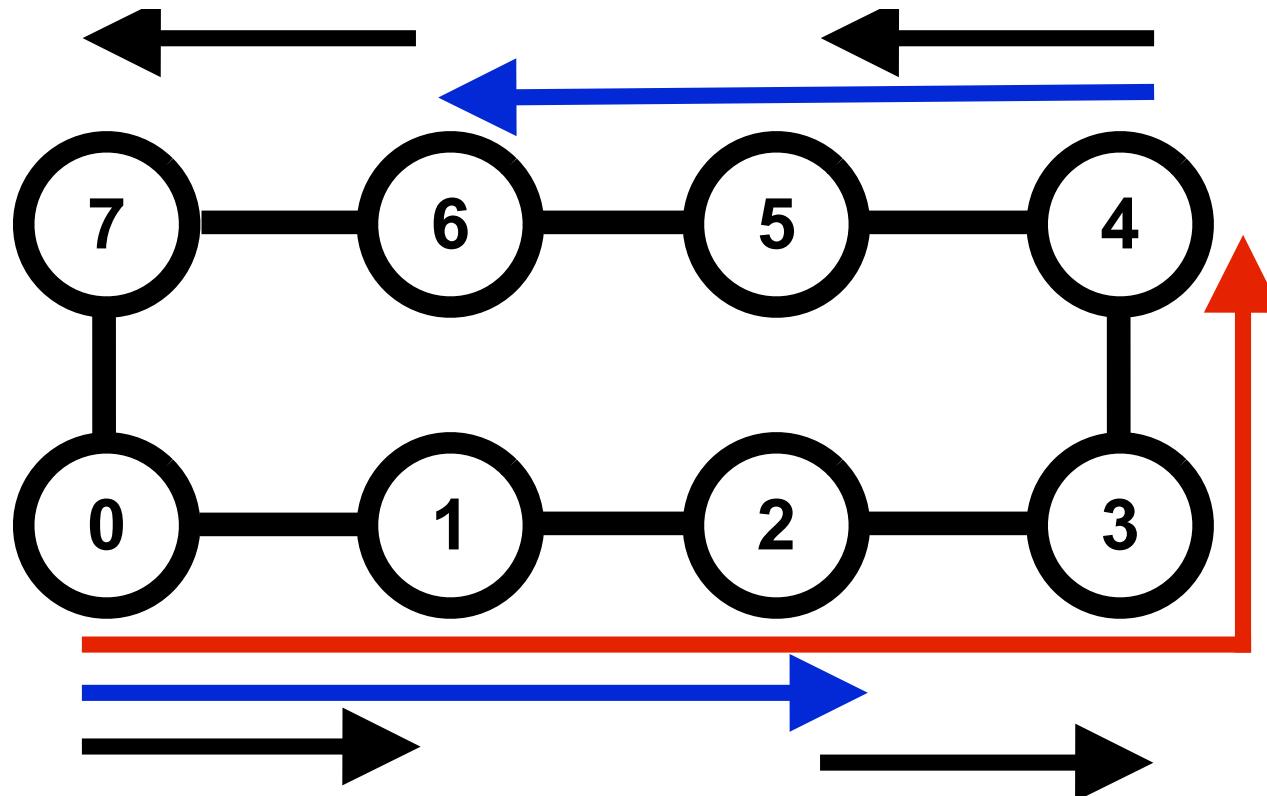
Ağ topolojileri



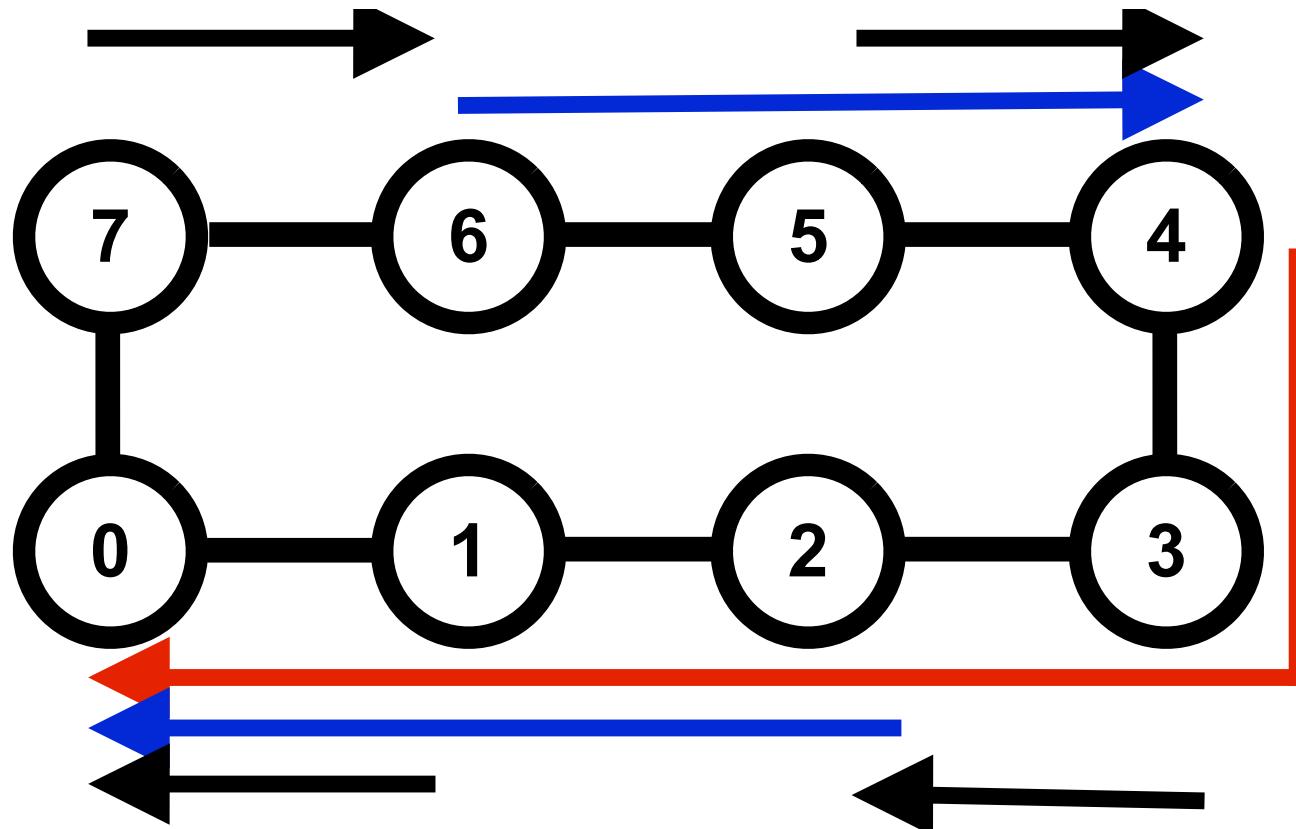
Örnek problem: bir mesajı
her yere gönderme



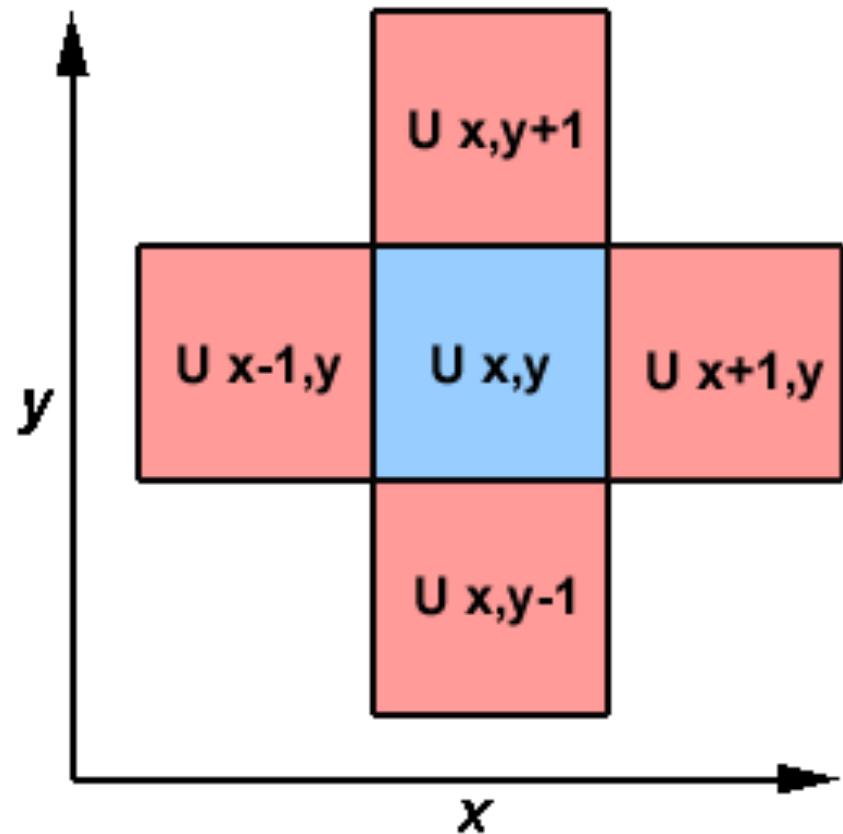
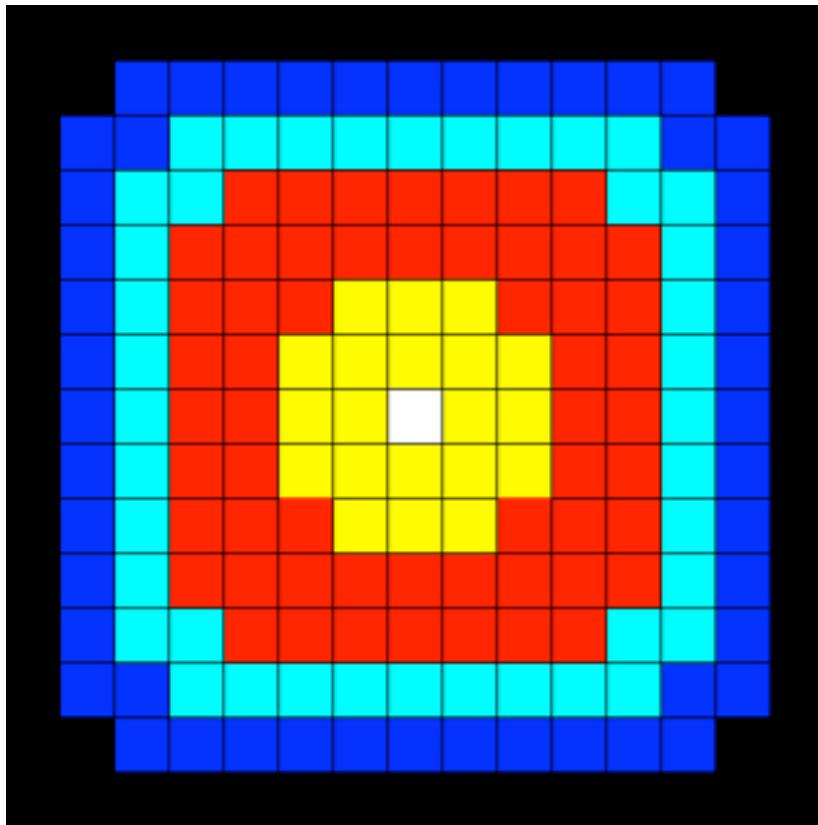
Örnek problem: bir mesajı
her yere gönderme



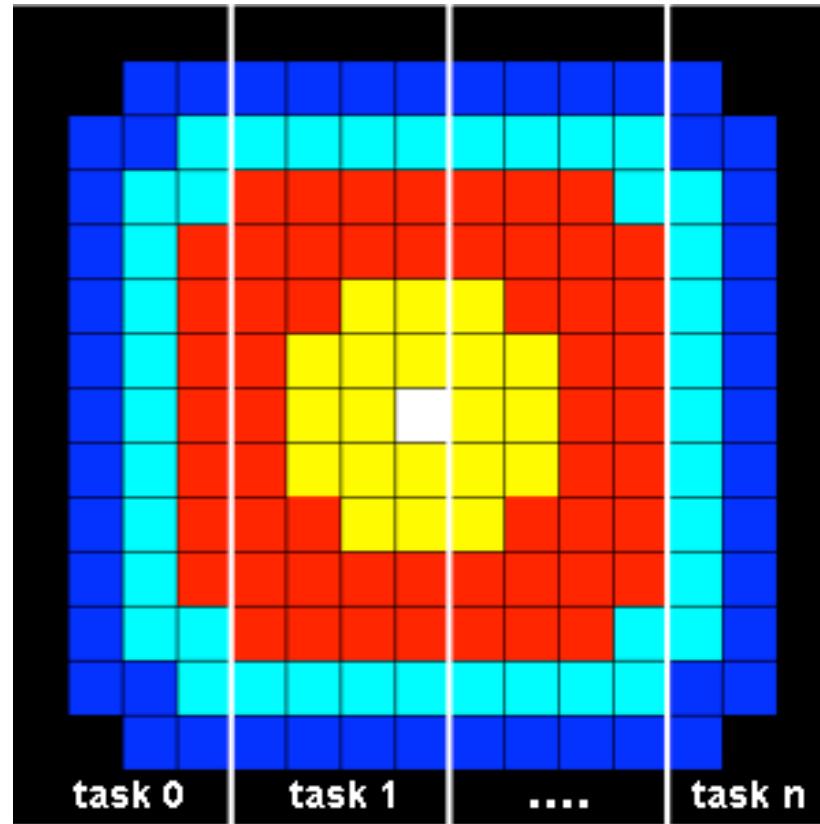
Örnek problem: mesajları bir yerde toplama



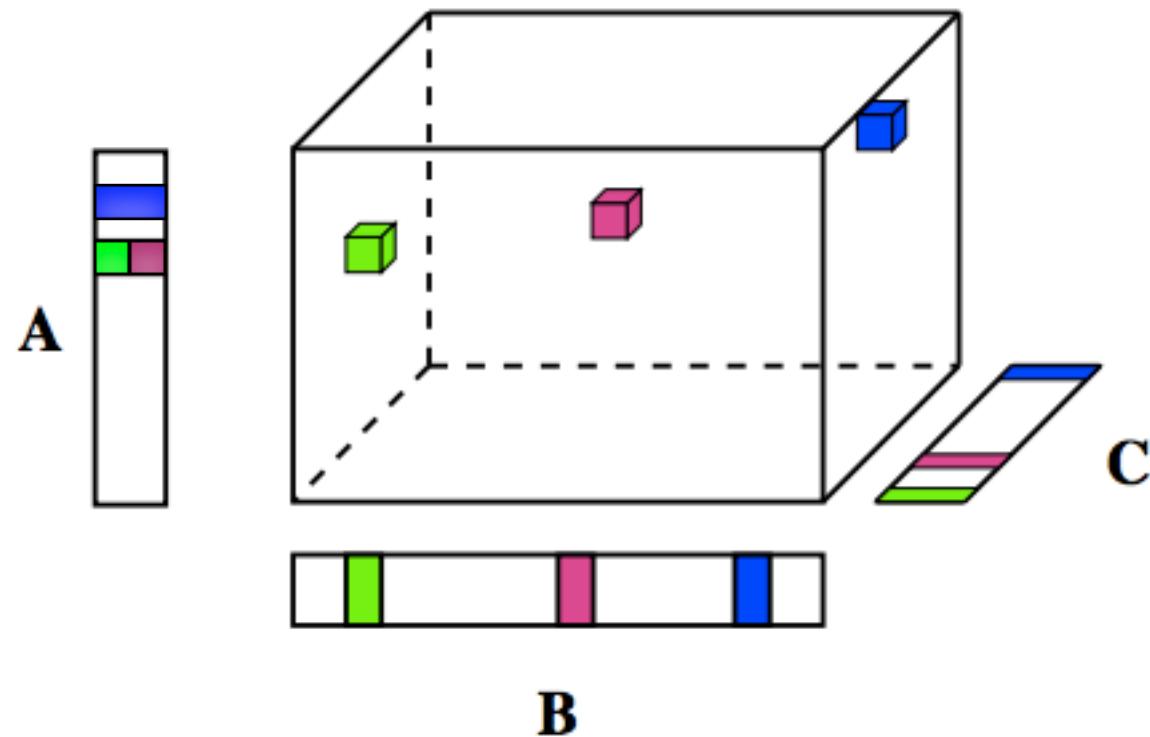
Neden veri gönderiyoruz?



Neden veri gönderiyoruz?



Neden veri gönderiyoruz?



MPI

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    int npes, myrank;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &npes);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    printf("From process %d out of %d, Hello World!\n",
           myrank, npes);
    MPI_Finalize();
    return 0;
}
```

MPI

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
            int dest_pe, int tag, MPI_Comm comm)  
  
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
            int source_pe, int tag, MPI_Comm comm,  
            MPI_Status *status)
```

MPI: herkes bir sonraki işlemciye mesaj gönderirse

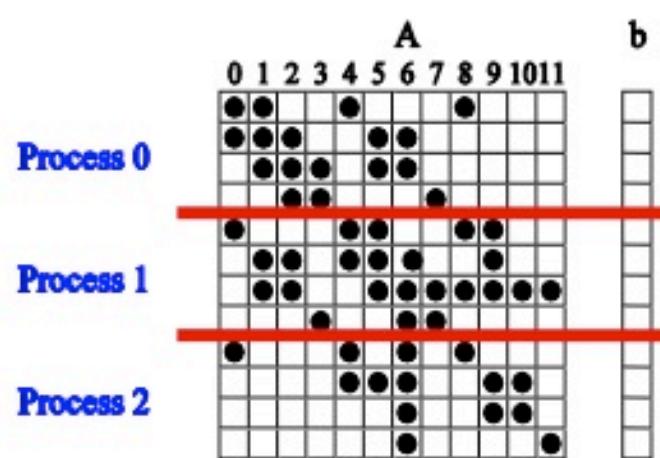
```
int a[10], b[10], npes, myrank;  
MPI_Status status;  
...  
MPI_Comm_size(MPI_COMM_WORLD, &npes);  
  
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);  
  
MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,  
         MPI_COMM_WORLD);  
  
MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,  
         MPI_COMM_WORLD, &status);  
...
```

MPI: herkes bir sonraki işlemciye mesaj gönderirse

```
if (myrank%2 == 1) { // odd processes send first, receive second
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
              MPI_COMM_WORLD);
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
              MPI_COMM_WORLD, &status);
}
else { // even processes receive first, send second
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
              MPI_COMM_WORLD, &status);
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
              MPI_COMM_WORLD);
}
```

Matrisler ve (Hiper)çizgeler

96



17 items to communicate

$$C_0 = (4, 5, 6, 7, 8)$$

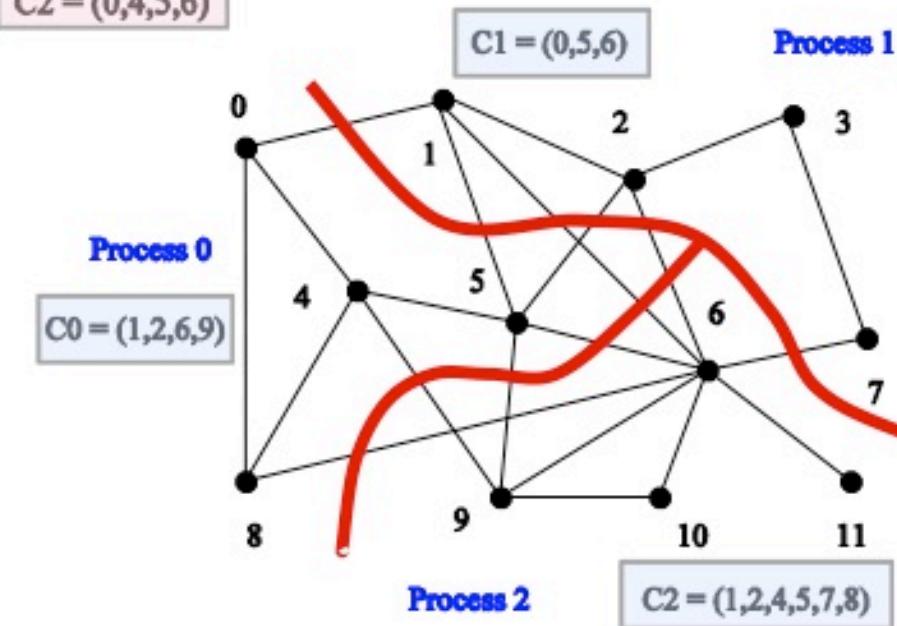
$$C_1 = (0, 1, 2, 3, 8, 9, 10, 11)$$

$$C_2 = (0, 4, 5, 6)$$

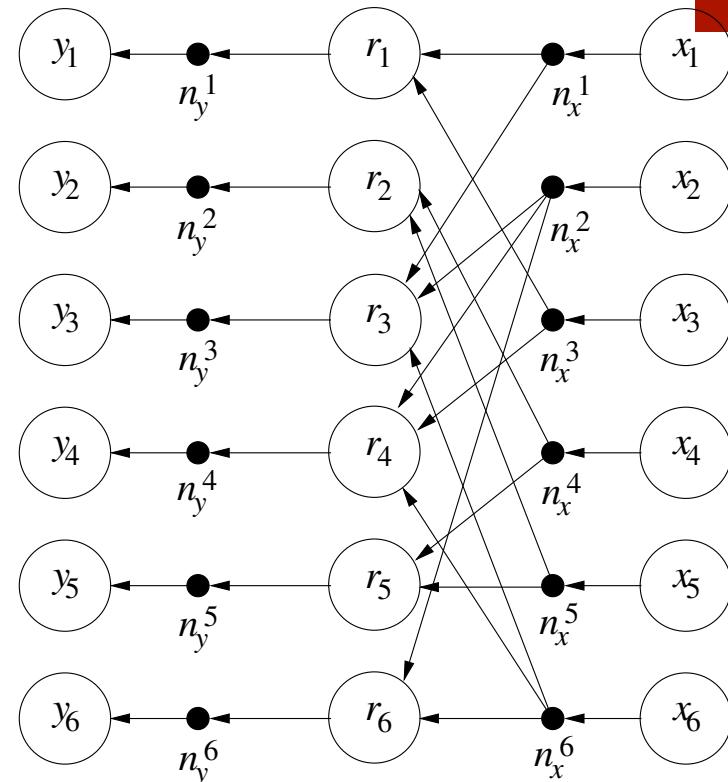
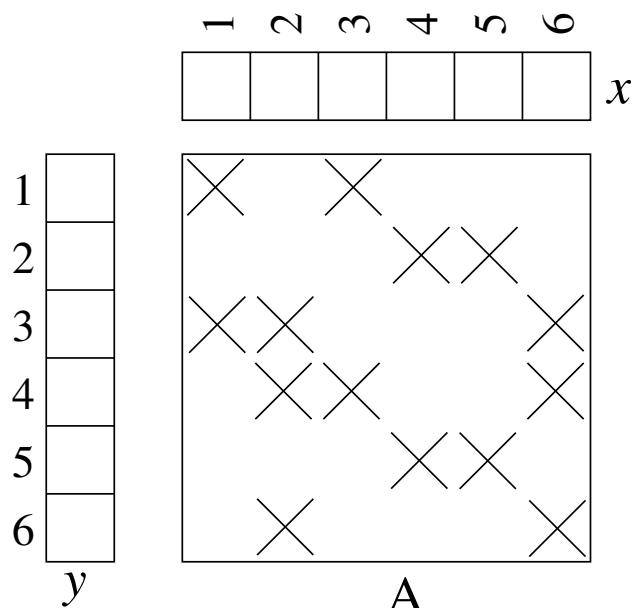
13 items to communicate

sparse matrix structure

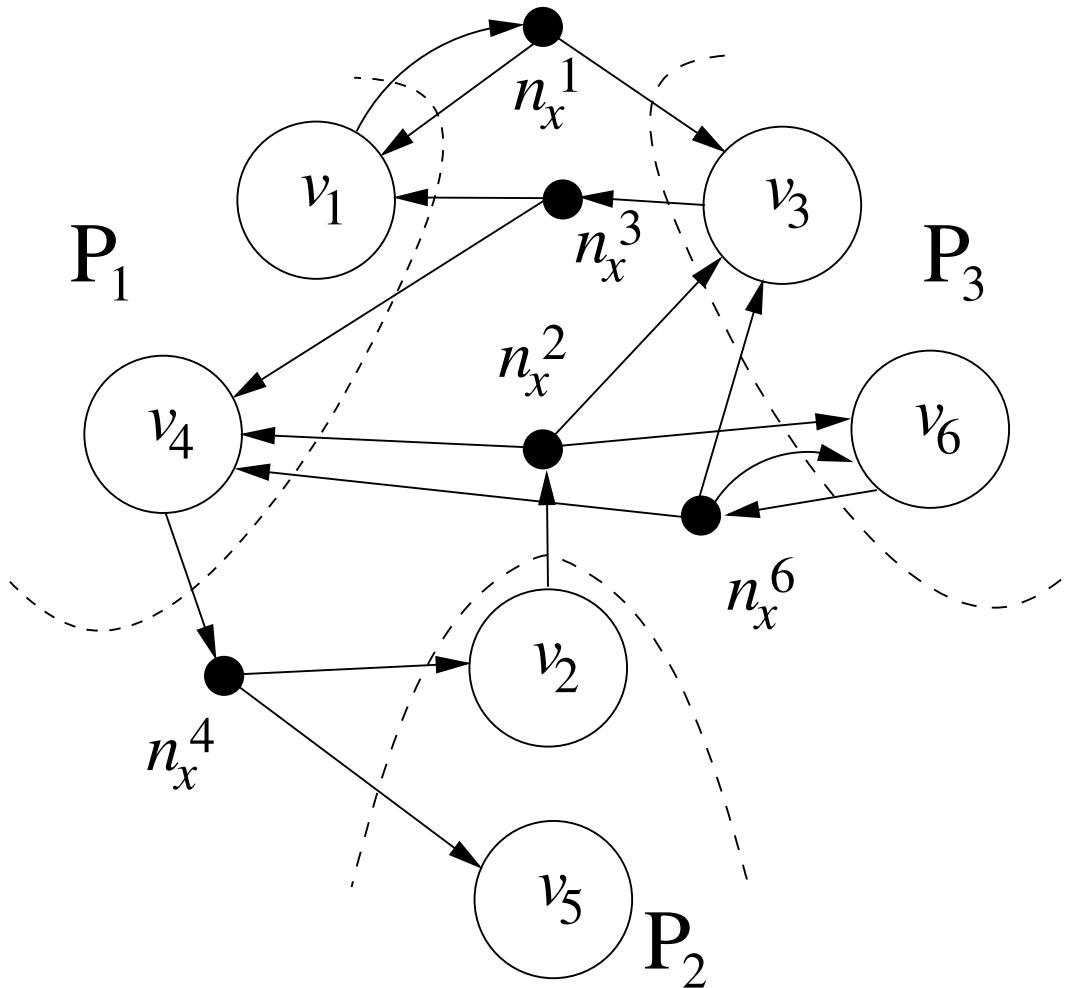
mapping
partitioning



Örnek problem: SpMV



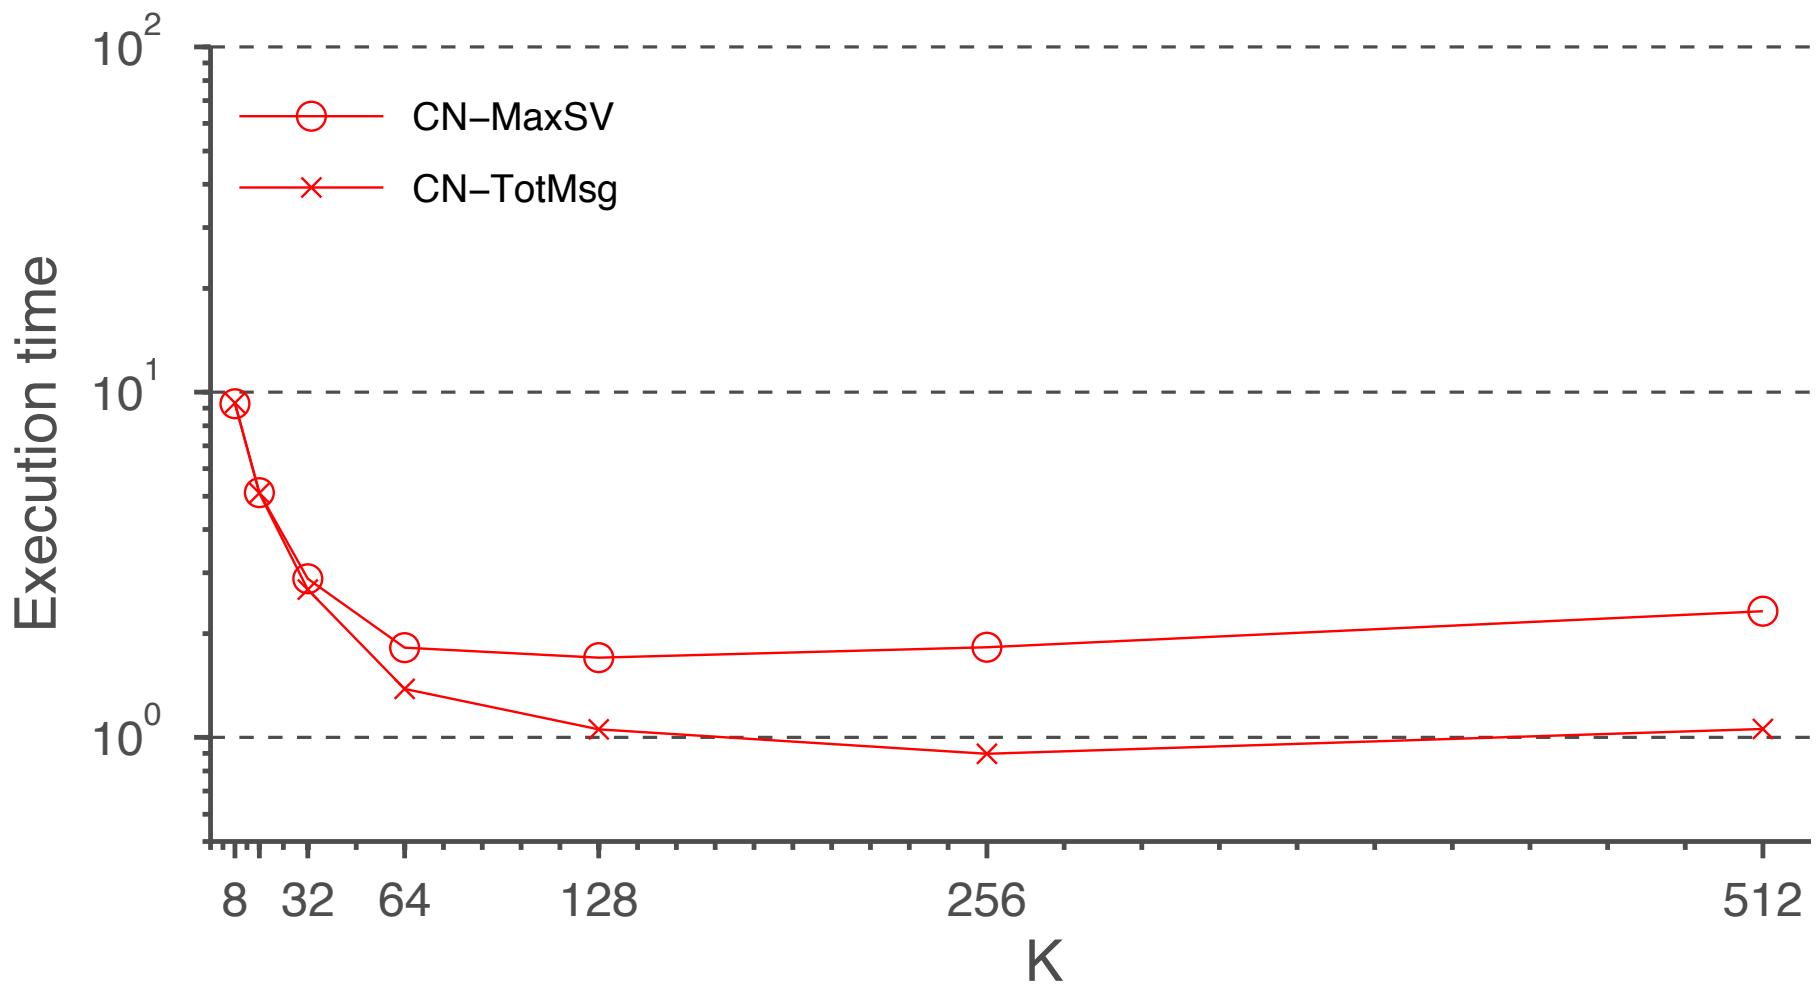
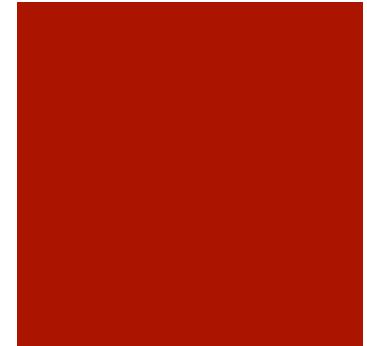
Örnek problem: SpMV



Kapasite metrikleri
Gönderilen
Alınan
Toplam **Byte**

Hız metrikleri
Gönderilen
Alınan
Toplam **Mesaj**

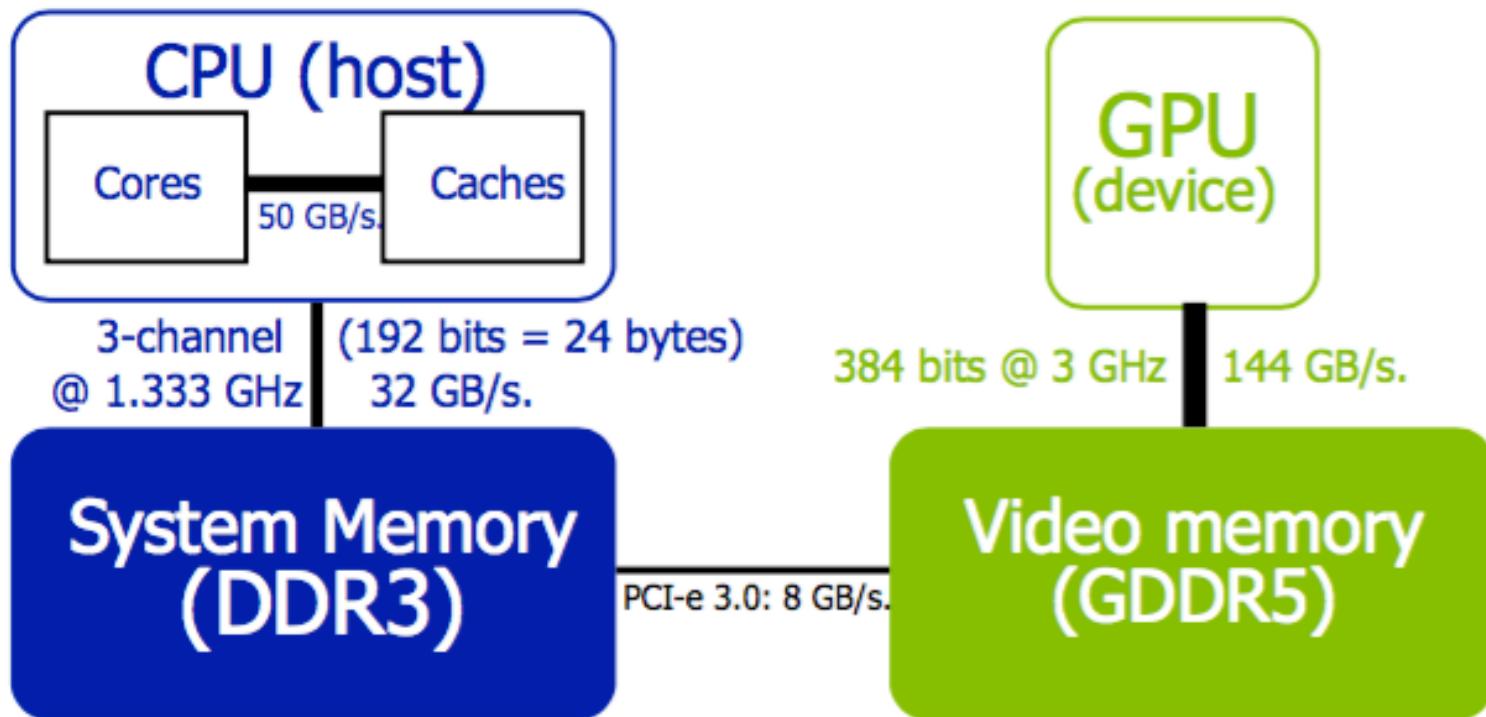
Örnek problem: SpMV



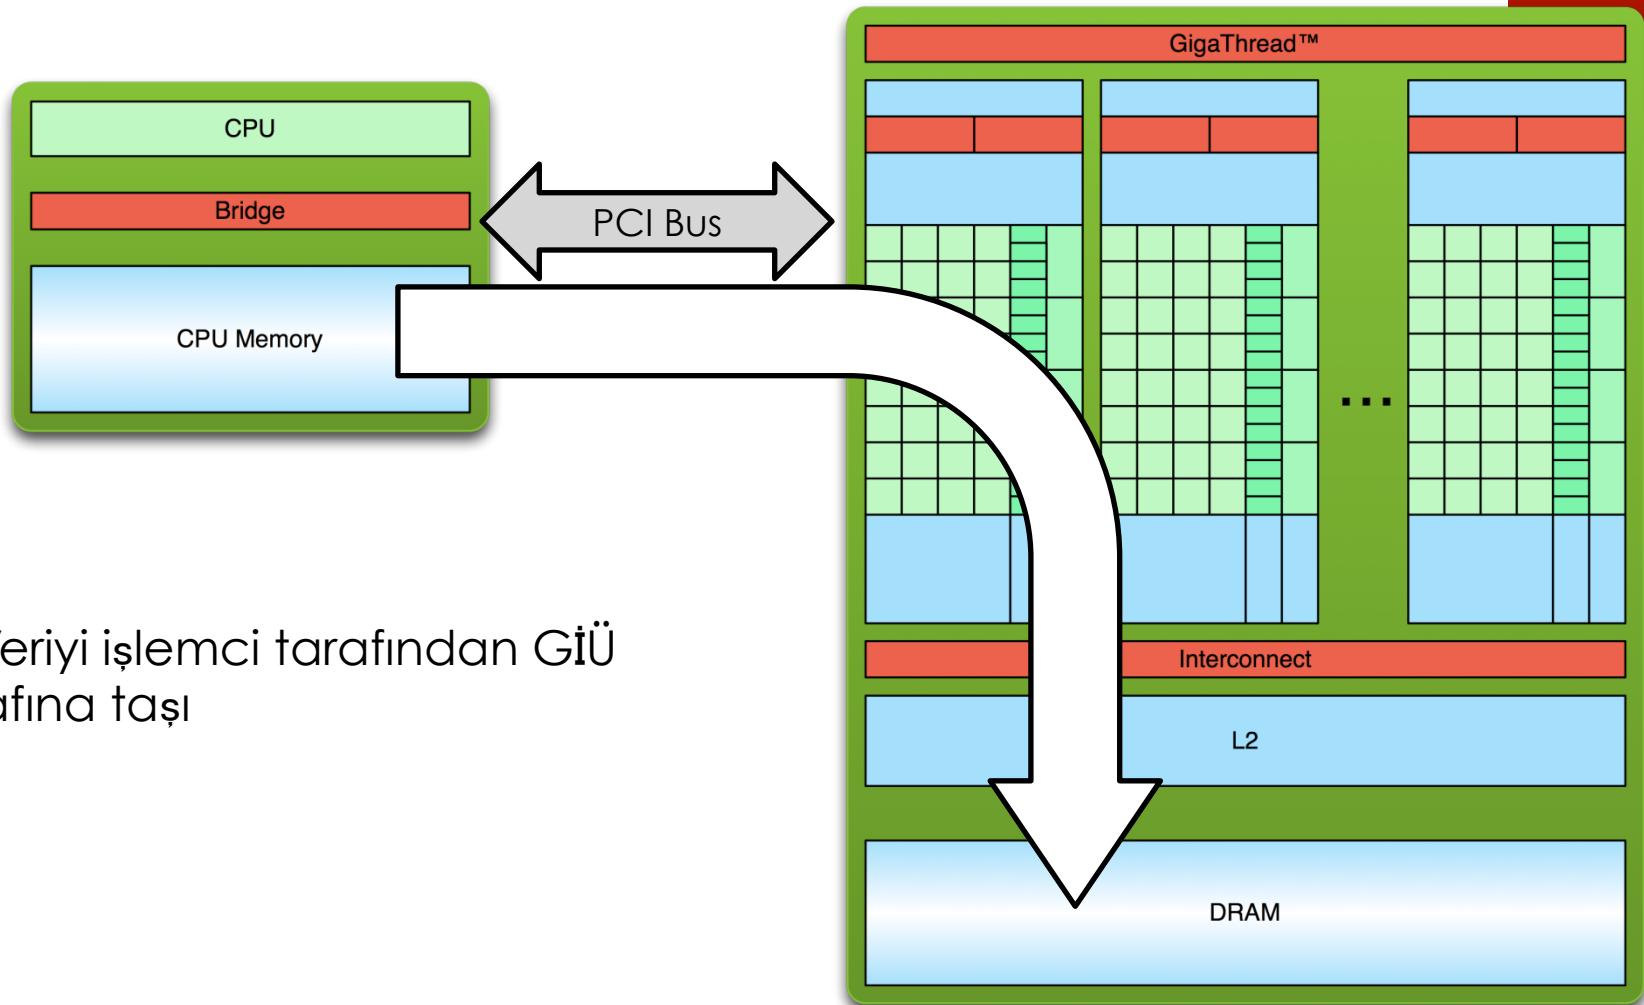
Grafik İşleme Üniteleri



Grafik İşleme Üniteleri

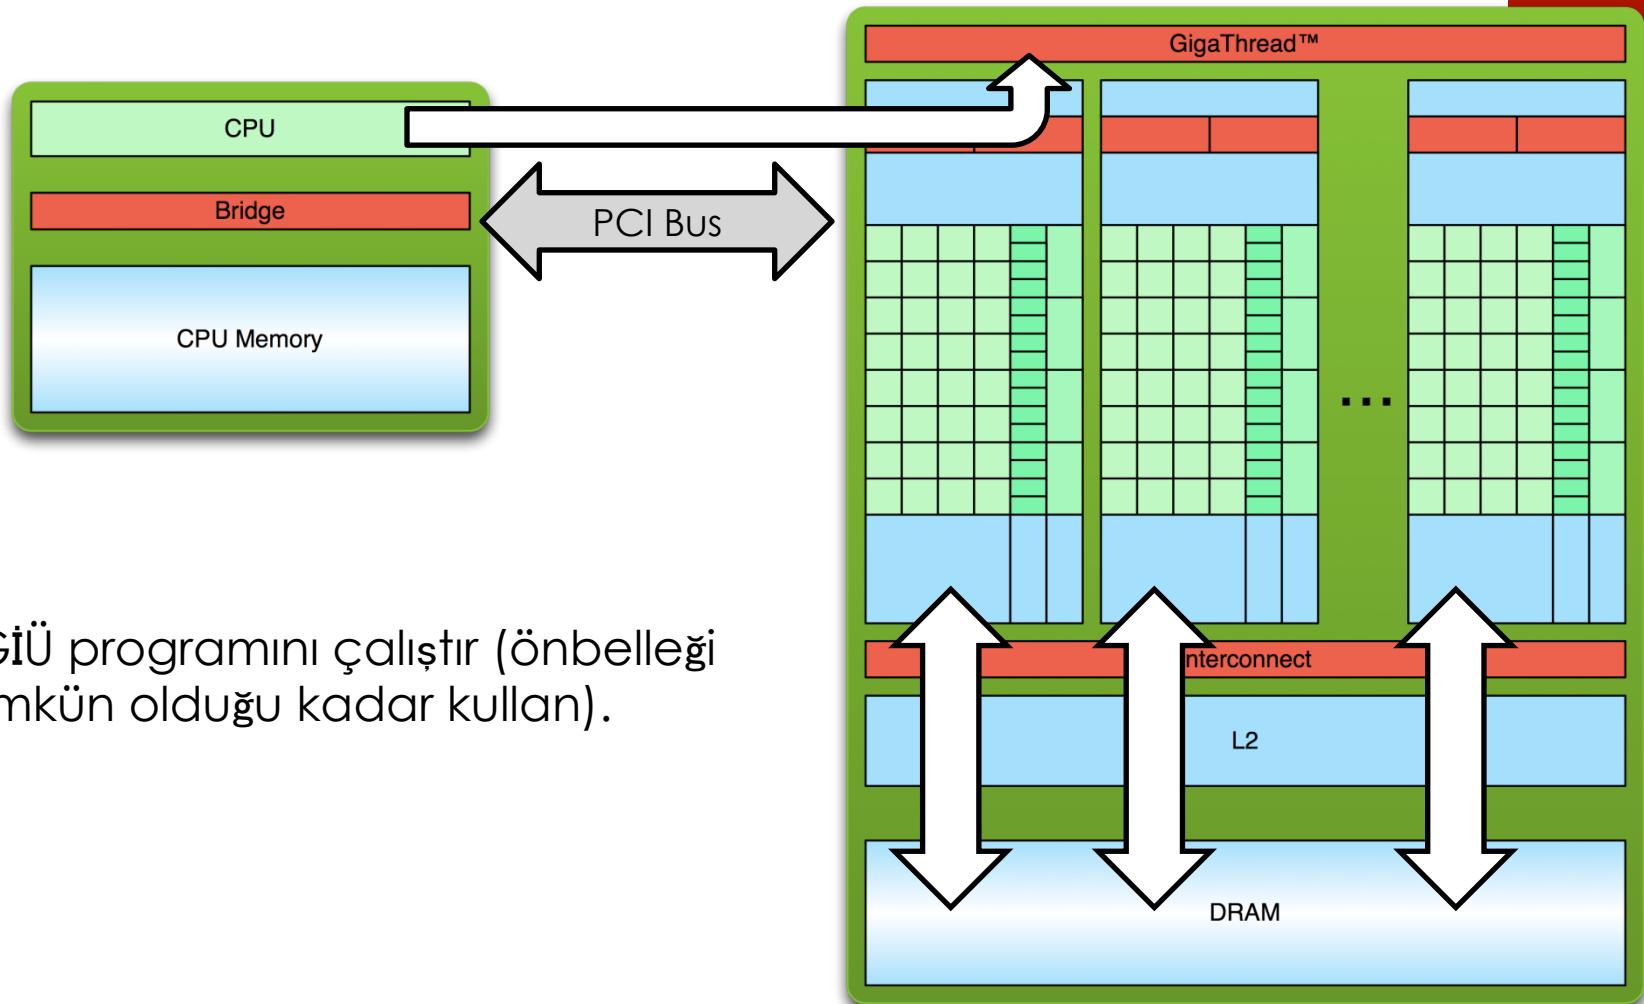


(Basitçe) GIÜ kullanımı



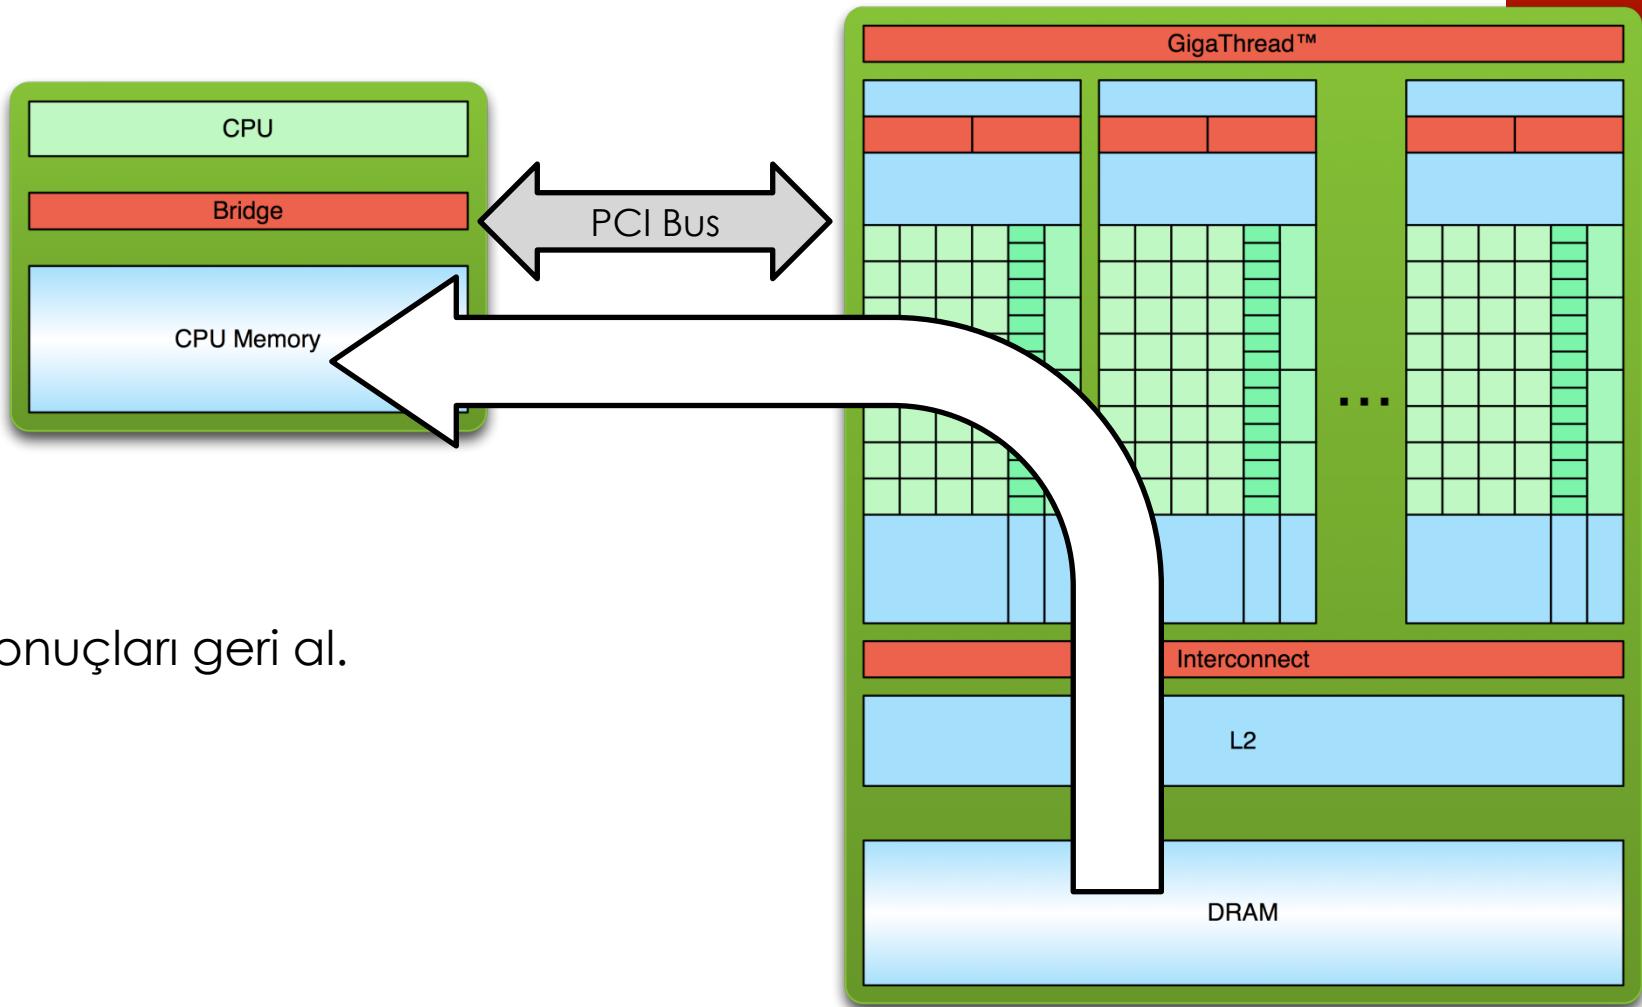
1- Veriyi işlemci tarafından GIÜ tarafına taşı

(Basitçe) GIÜ kullanımı



2- GIÜ programını çalıştır (önbeliği mümkün olduğu kadar kullan).

(Basitçe) GIÜ kullanımı



3- Sonuçları geri al.

GIÜ: Özellikler

Öznitelikler	Tesla K80	Tesla K40
GPU	2x Kepler GK210	1 Kepler GK110B
Peak double precision float. point performance	2.91 Tflops (GPU Boost Clocks) 1.87 Tflops (Base Clocks)	1.66 Tflops (GPU Boost Clocks) 1.43 Tflops (Base Clocks)
Peak single precision float. point performance	8.74 Tflops (GPU Boost Clocks) 5.6 Tflops (Base Clocks)	5 Tflops (GPU Boost Clocks) 4.29 Tflops (Base Clocks)
Memory bwidth (ECC off) ²	480 GB/sec (240 GB/sec per GPU)	288 GB/sec
Memory size (GDDR5)	24 GB (12GB per GPU)	12 GB
<u>CUDA cores</u>	4992 (2496 per GPU)	2880

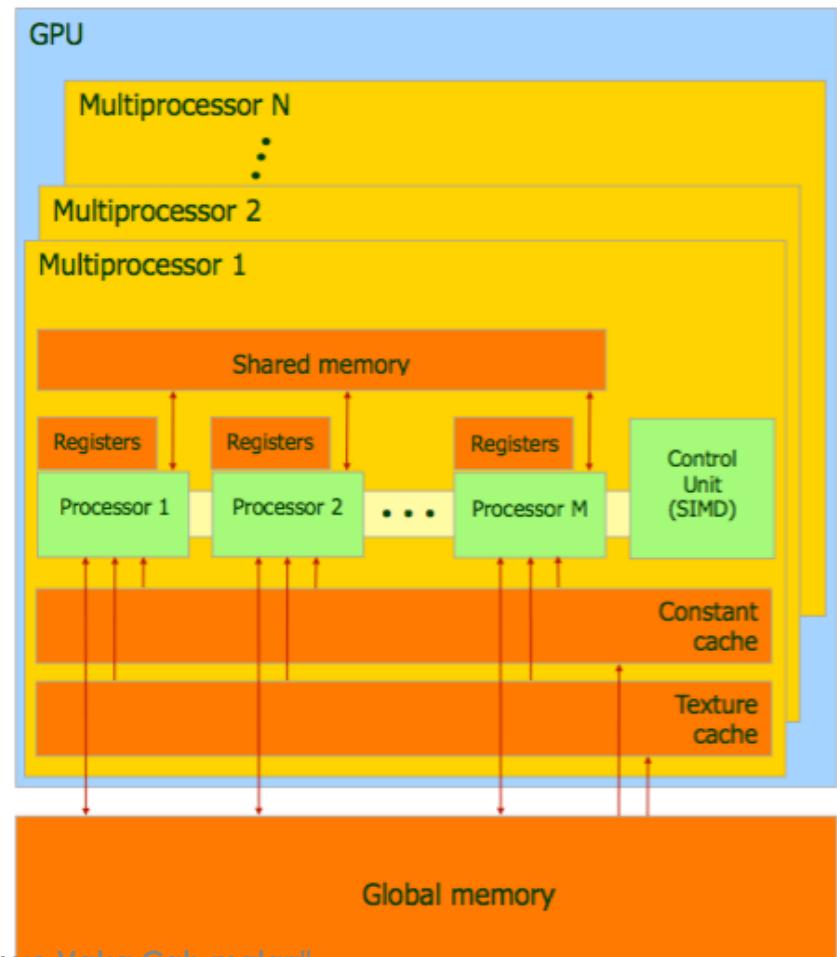
Grafik İşleme Üniteleri

SPECIFICATIONS

GPU Architecture	NVIDIA Pascal
NVIDIA CUDA® Cores	3584
Double-Precision Performance	5.3 TeraFLOPS
Single-Precision Performance	10.6 TeraFLOPS
Half-Precision Performance	21.2 TeraFLOPS
GPU Memory	16 GB CoWoS HBM2
Memory Bandwidth	732 GB/s
Interconnect	NVIDIA NVLink
Max Power Consumption	300 W

NVIDIA TESLA P100

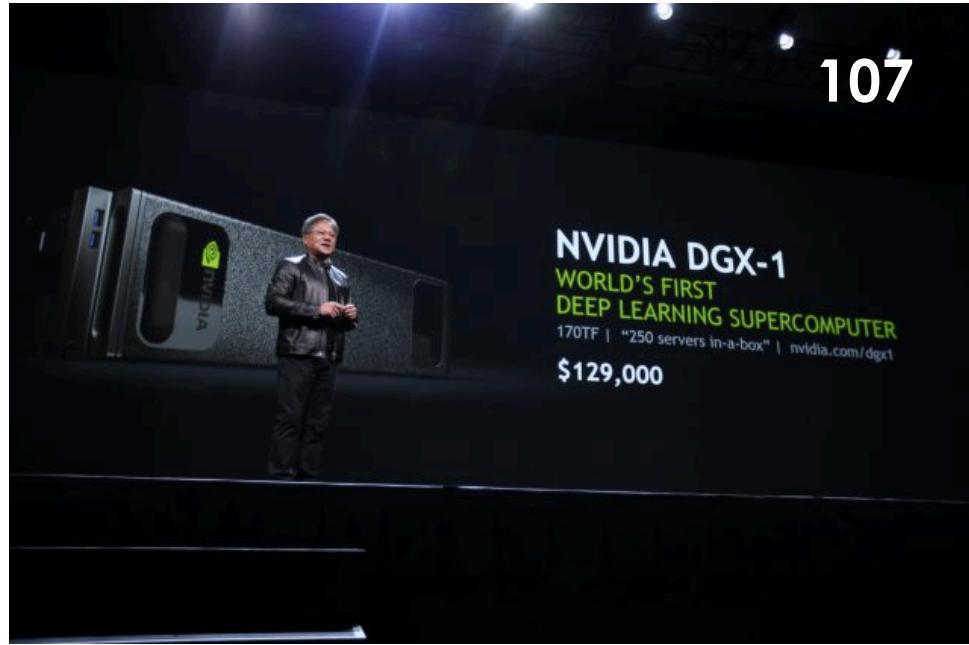
7000\$



Grafik İşleme Üniteleri

SYSTEM SPECIFICATIONS

GPUs	8x Tesla P100
TFLOPS [GPU FP16 / CPU FP32]	170/3
GPU Memory	16 GB per GPU
CPU	Dual 20-core Intel® Xeon® E5-2698 v4 2.2 GHz
NVIDIA CUDA® Cores	28672
System Memory	512 GB 2133 MHz DDR4
Storage	4x 1.92 TB SSD RAID 0
Network	Dual 10 GbE, 4 IB EDR
Software	Ubuntu Server Linux OS DGX-1 Recommended GPU Driver
System Weight	134 lbs
System Dimensions	866 D x 444 W x 131 H (mm)
Packing Dimensions	1180 D x 730 W x 284 H (mm)
Maximum Power Requirements	3200W
Operating Temperature Range	10 - 30°C



GIÜ: Programlama - CUDA

```
__global__
void hello(char *a, int *b) {
    a[threadIdx.x] += b[threadIdx.x];
}
```

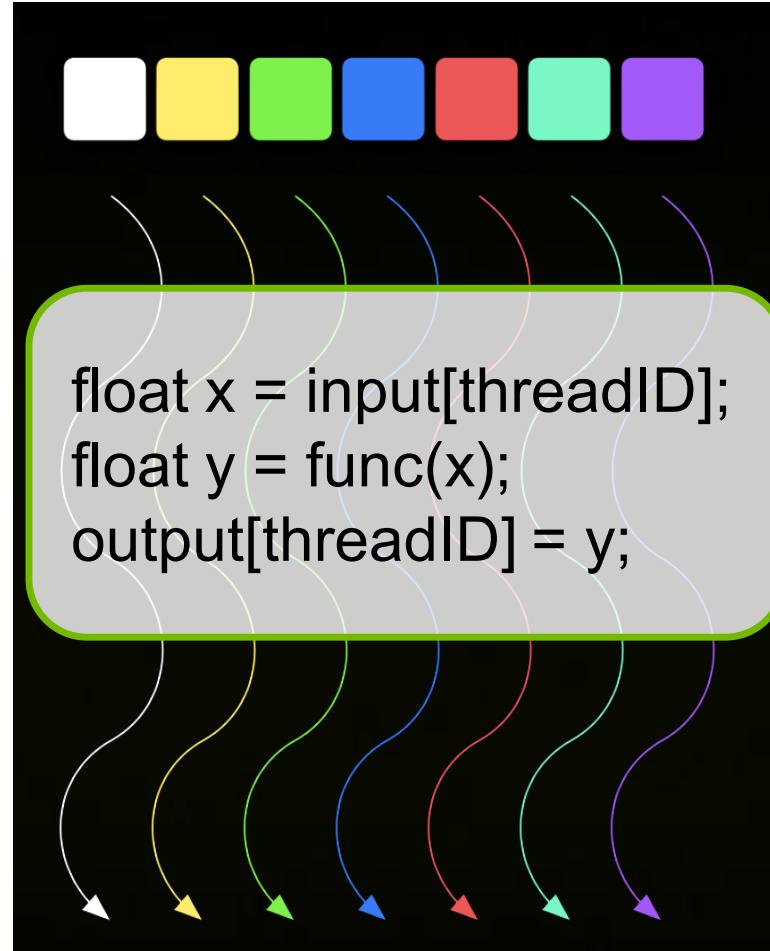
GIÜ: Programlama - CUDA

```
cudaMalloc( (void**)&bd, isize );
cudaMemcpy( ad, a, csize, cudaMemcpyHostToDevice );
cudaMemcpy( bd, b, isize, cudaMemcpyHostToDevice );

dim3 dimBlock( blocksize, 1 );
dim3 dimGrid( 1, 1 );
hello<<<dimGrid, dimBlock>>>(ad, bd);
cudaMemcpy(a, ad, csize, cudaMemcpyDeviceToHost );
cudaFree( ad );
cudaFree( bd );
```

GiÜ: Programlama – CUDA Kernel

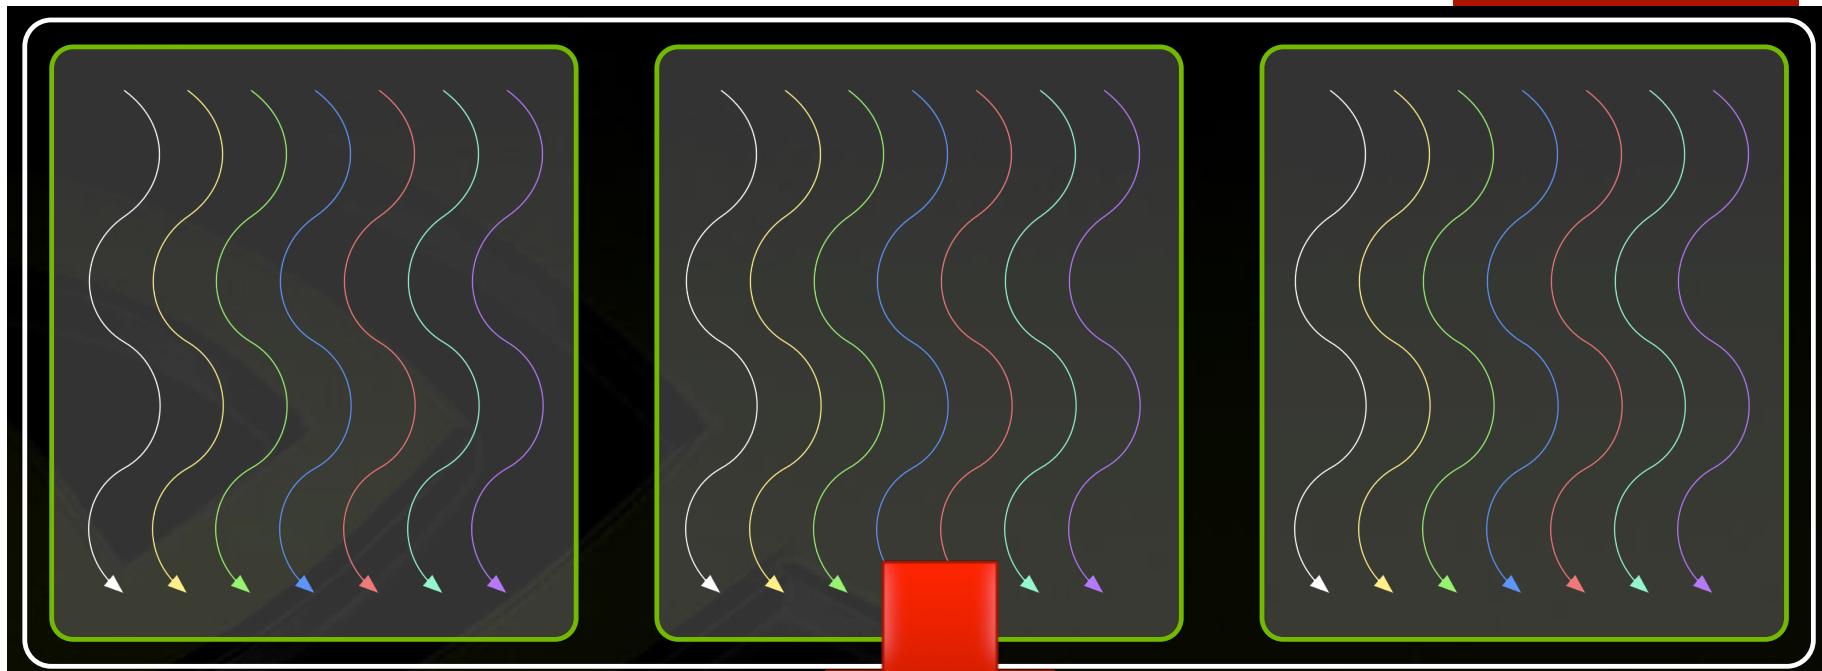
110



GiÜ: Programlama – CUDA

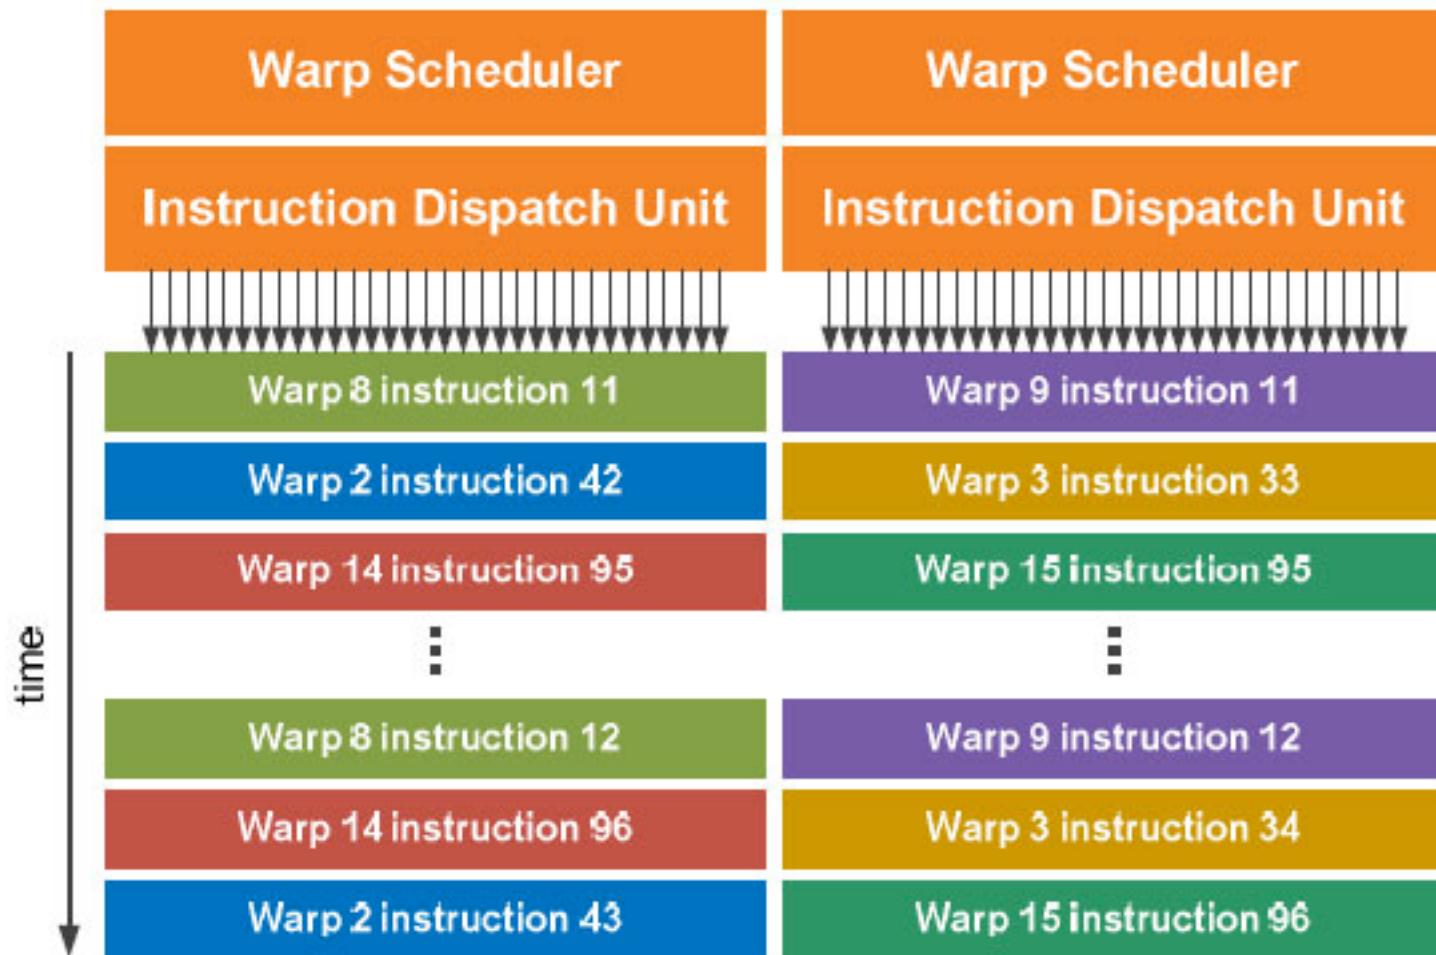
Kernelden Bloklara

111



GiÜ: Programlama – CUDA Bloklardan Warplara

112



GiÜ: Programlama – CUDA

Bloklardan Warplara

113

En önemli üç nokta

Warp içinde dallanan kodlar olmamalı

Warp içindeki ipliklere eşit yük paylaşımı

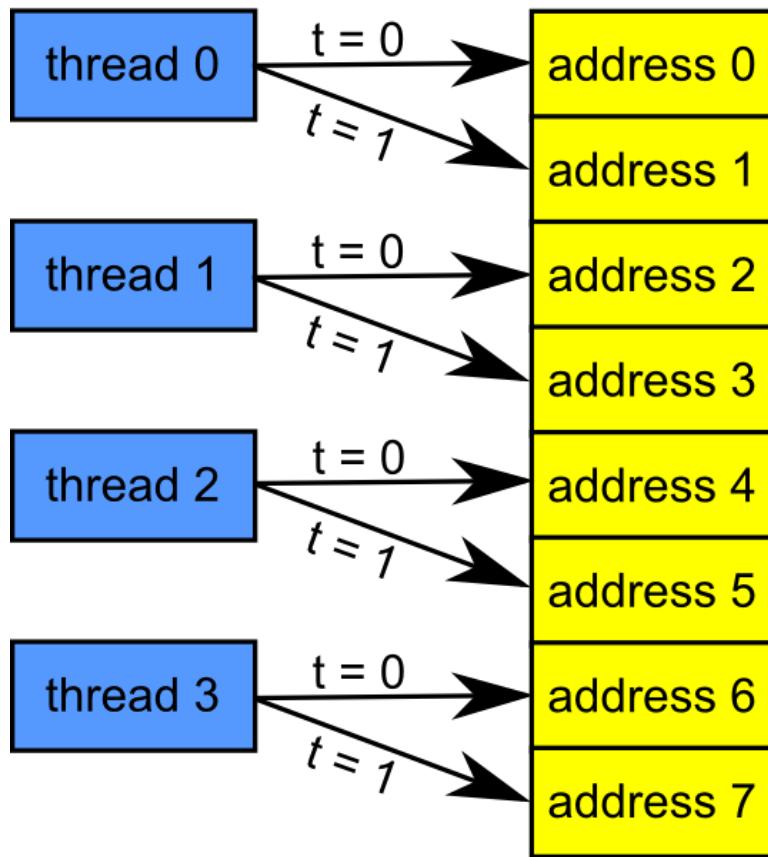
Warp içindeki ipliklerin kompakt hafıza erişimi

GiÜ: Programlama – CUDA

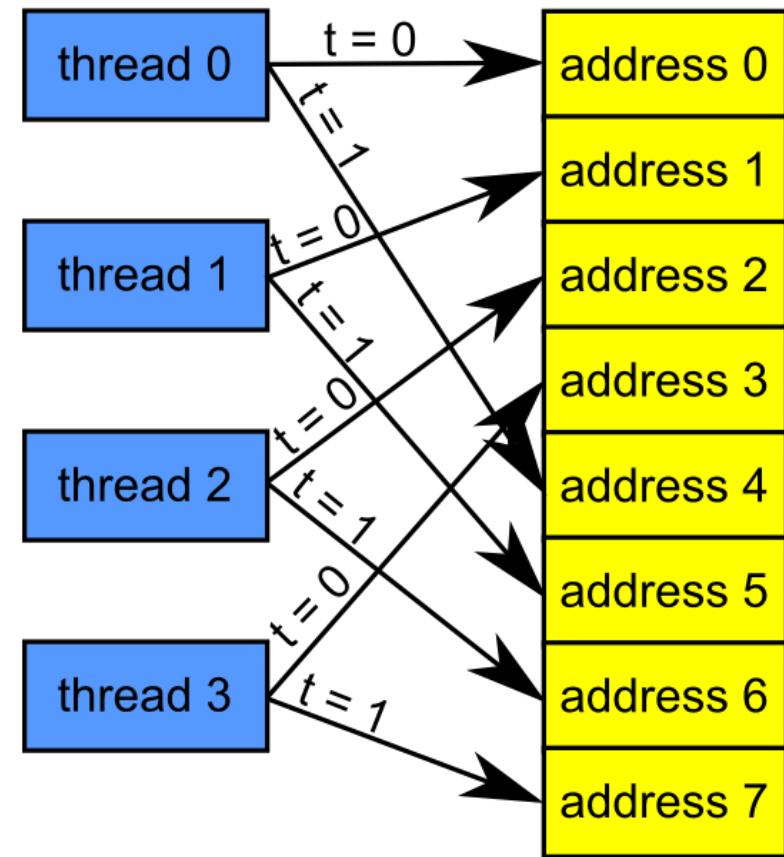
Bloklardan Warplara

114

traditional multi-core
optimal memory access pattern



many-core GPU
optimal memory access pattern



GiÜ: Programlama – CUDA Bloklardan Warplara

115

Hangisi bantgenişliğini daha verimli kullanır?

```
int a = input[i];
```

```
int b = input[2*i];
```

Örnek program: switch

```
int a = threadID.x;
```

```
switch(a) {
```

```
    case(1): doA();
```



```
    case(2): doB();
```



```
    case(3): doC();
```



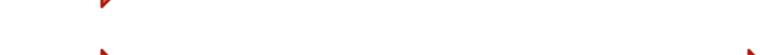
```
    case(4): doD();
```



```
    case(5): doE();
```



```
    case(6): doF();
```



```
    case(7): doG();
```

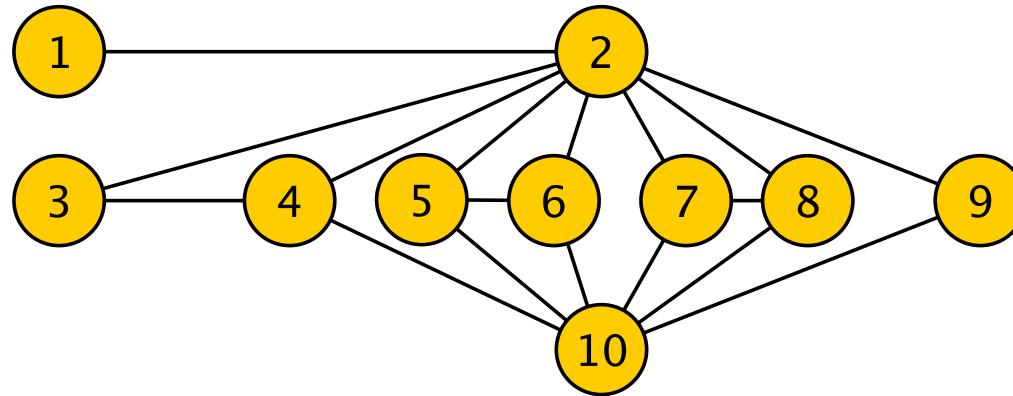


```
    case(8): doH();
```

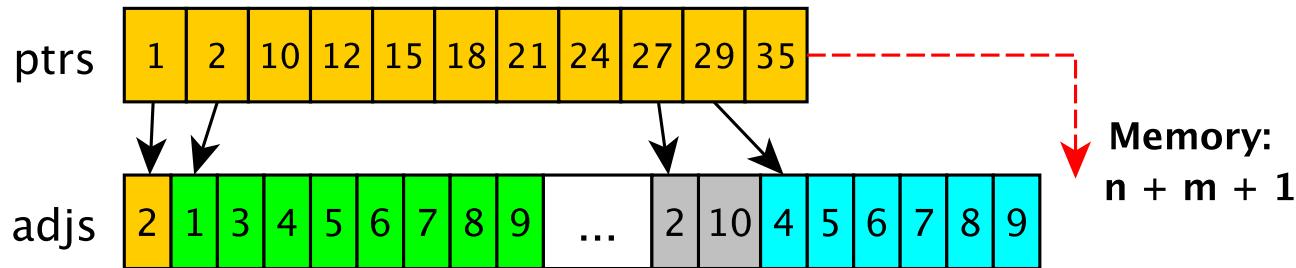


```
}
```

Örnek program: PageRank, BFS vs.

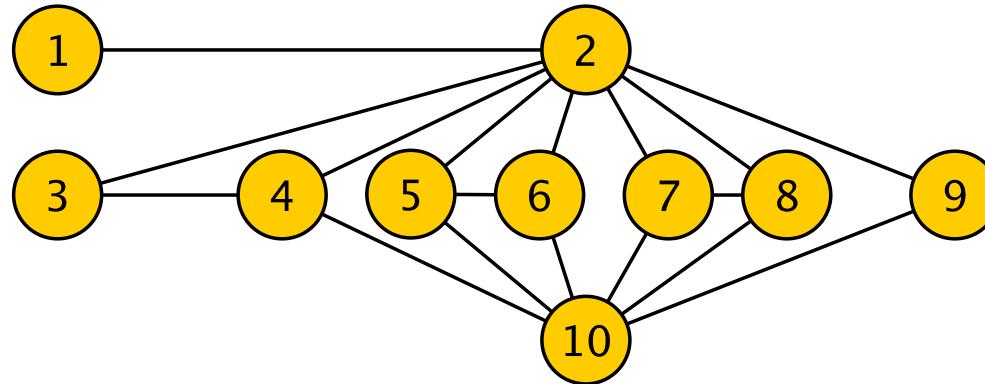


(a) A toy graph G



(b) CSR representation of G

Örnek program: PageRank, BFS vs.



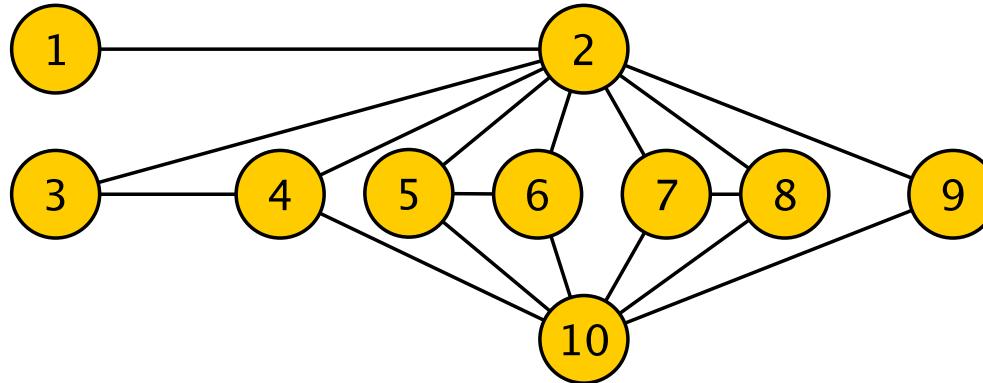
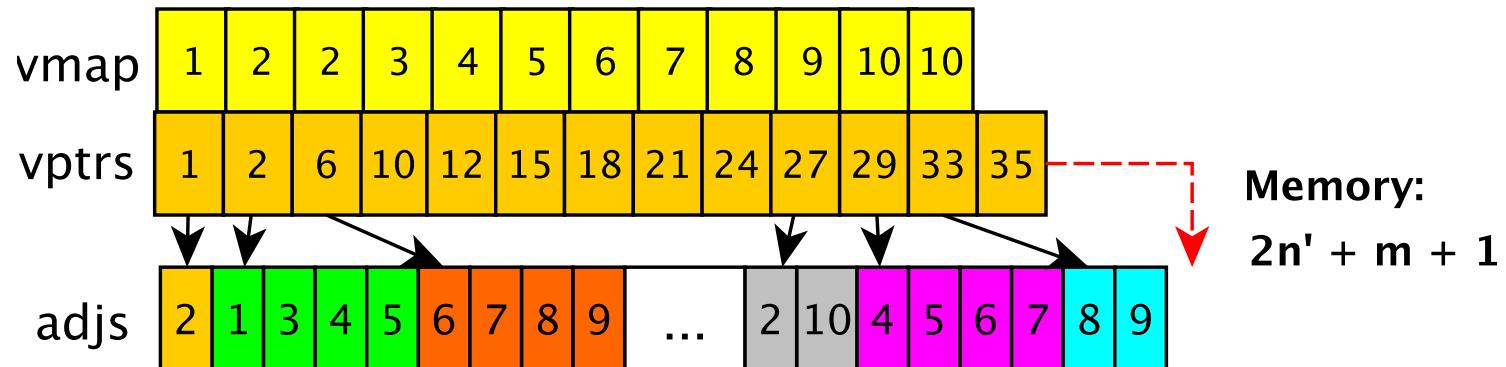
(a) A toy graph G

is	1	2	2	2	2	2	2	2	2	2	...	9	9	10	10	10	10	10	10
adjs	2	1	3	4	5	6	7	8	9	...	2	10	4	5	6	7	8	9	...

Memory:
2m

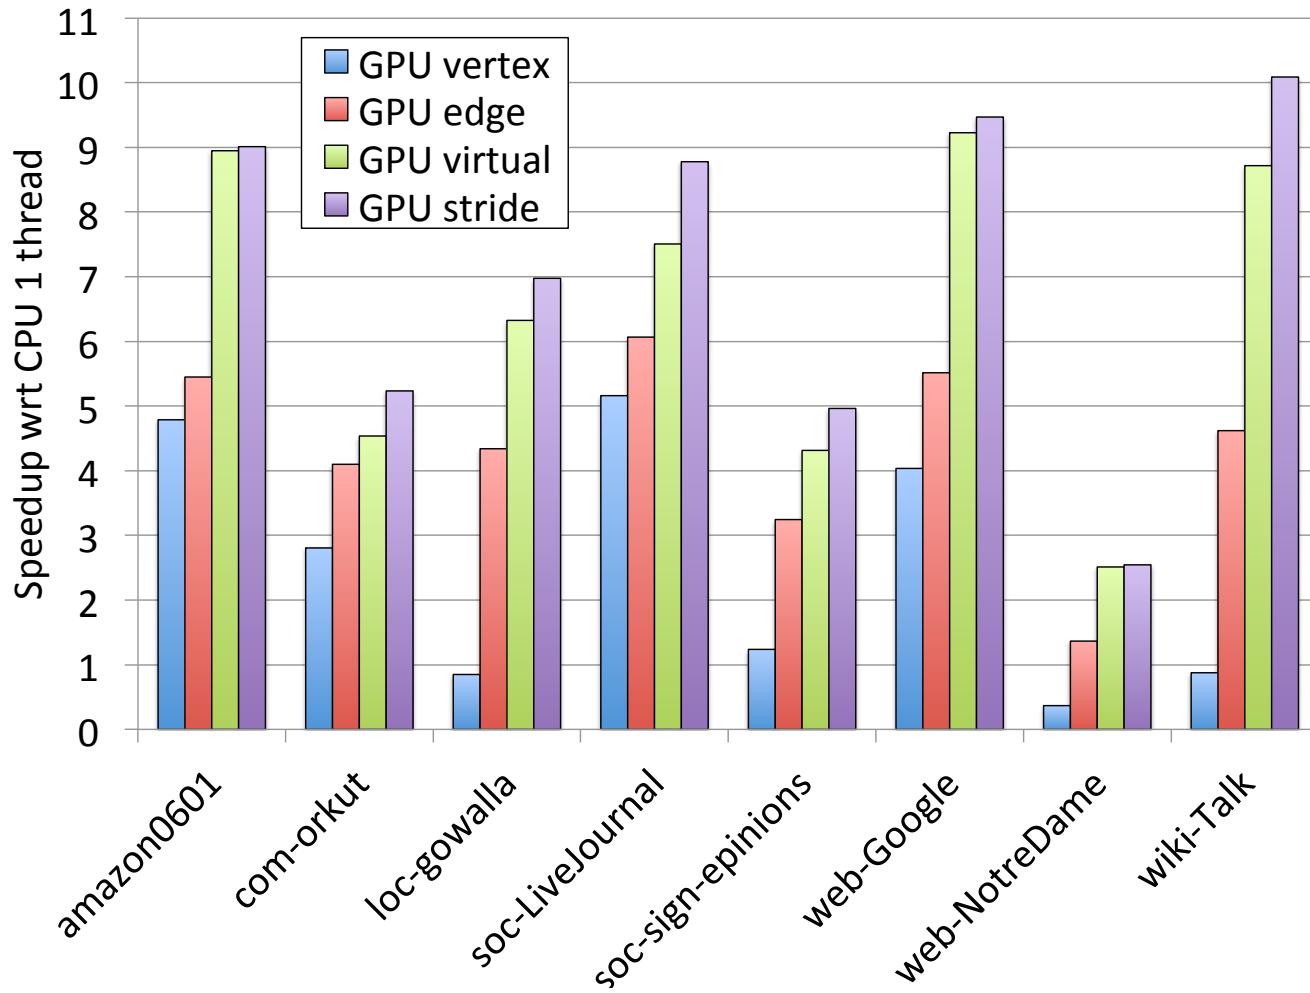
(c) COO representation of G

Örnek program: PageRank, BFS vs.

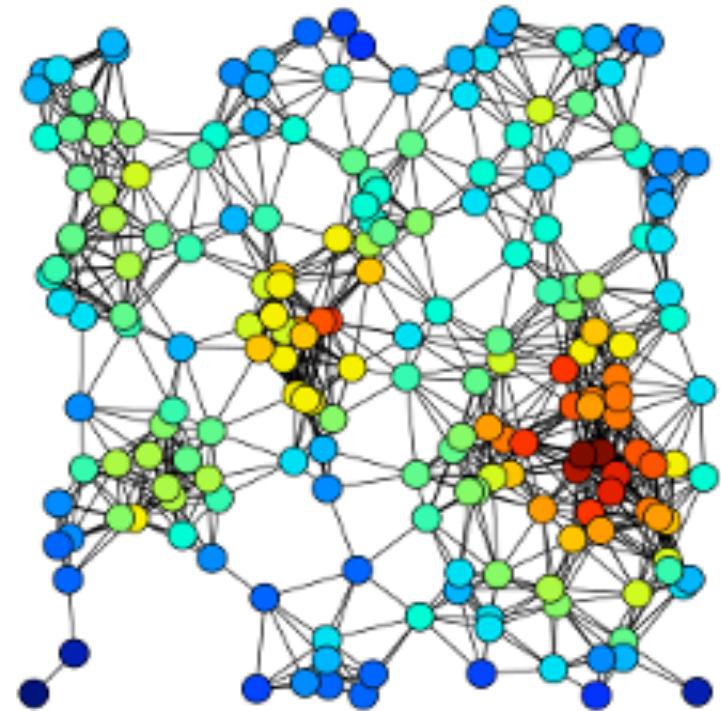
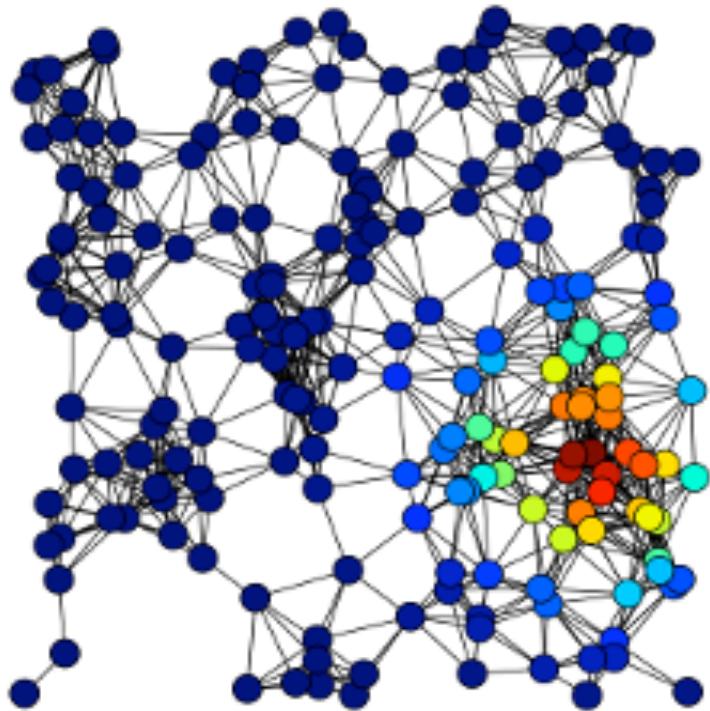
(a) A toy graph G (d) Virtual-CSR representation of G

Örnek program: PageRank, BFS vs.

120

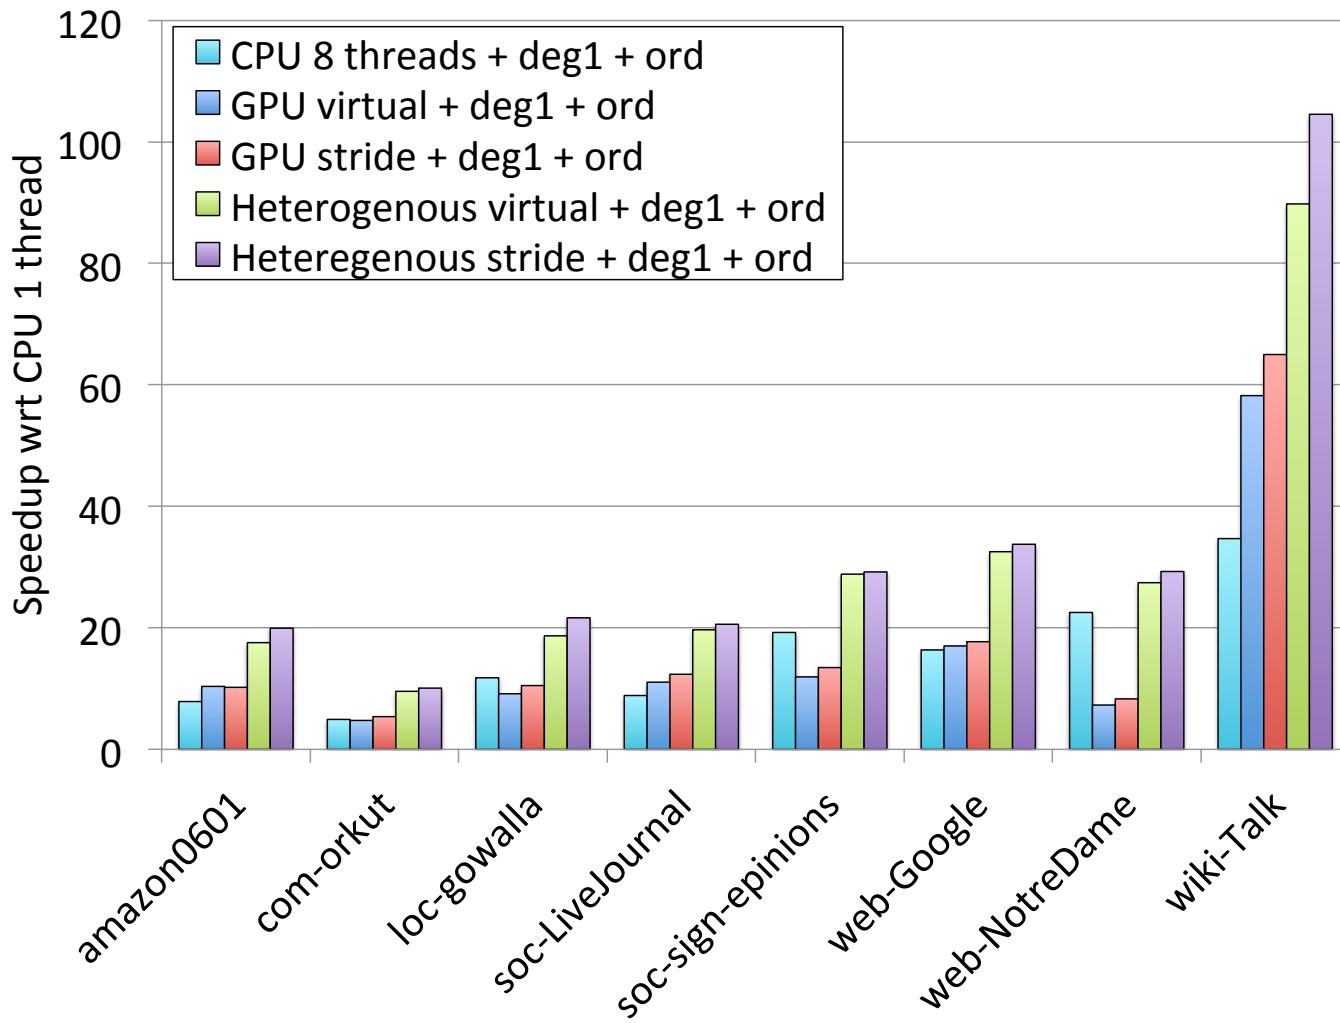


Örnek problem: Merkezi noktaları bulmak



Örnek problem: Merkezi noktaları bulmak

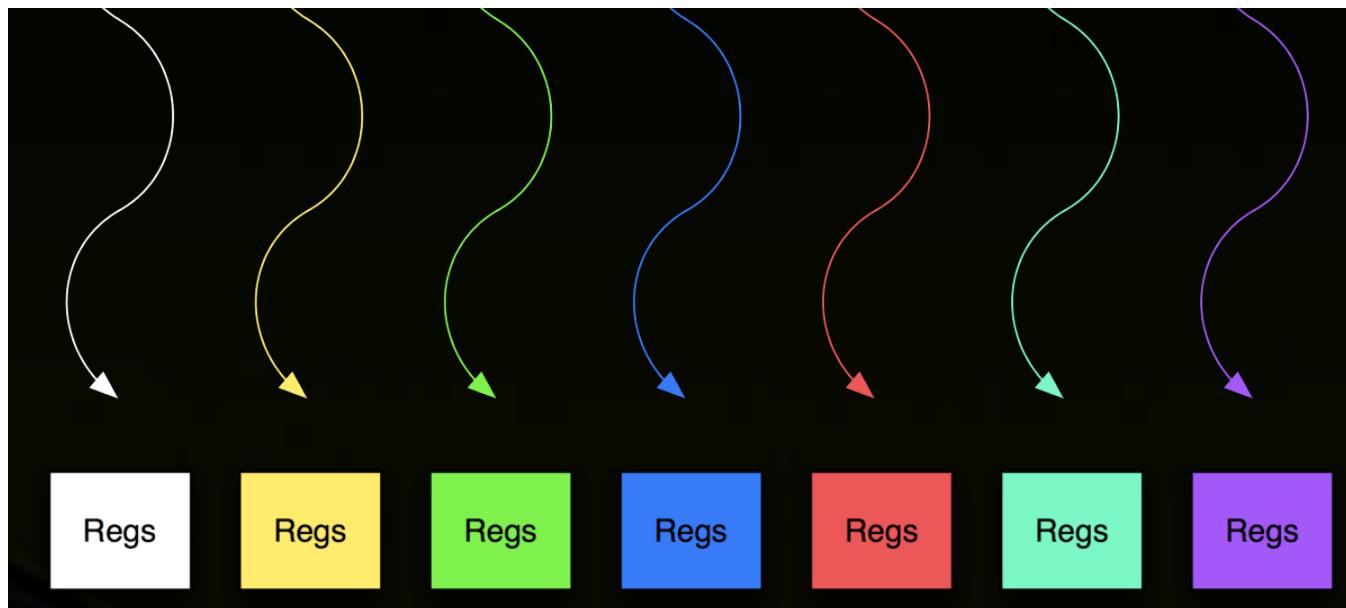
122



GiÜ: Programlama – CUDA

Hafıza yapısı

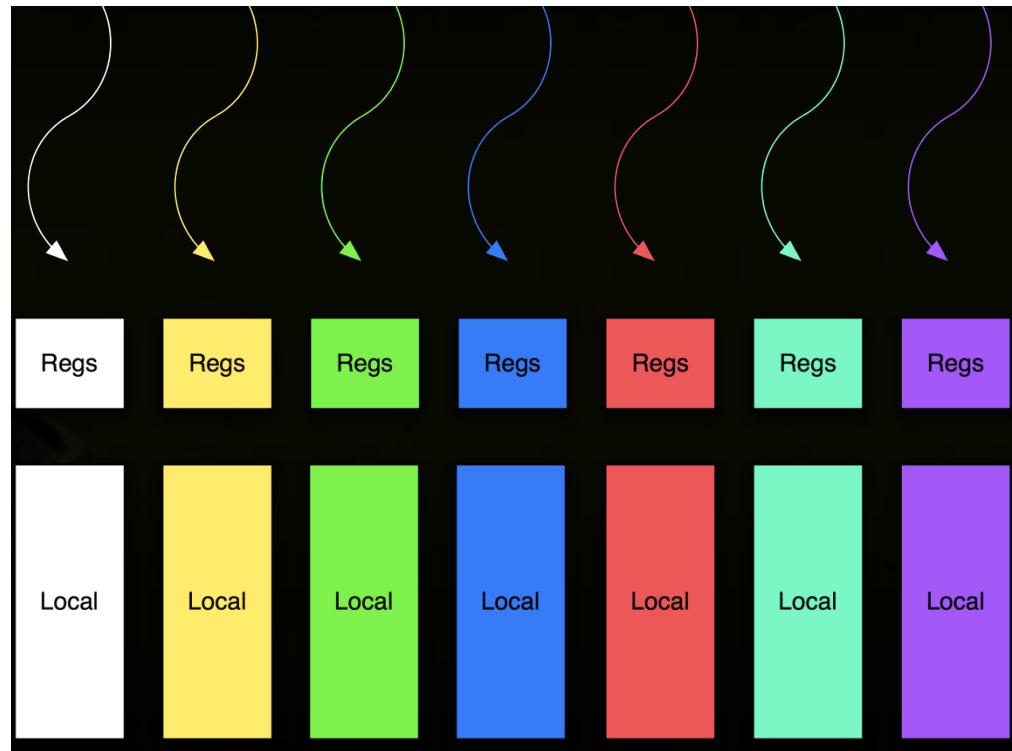
123



GiÜ: Programlama – CUDA

Hafıza yapısı

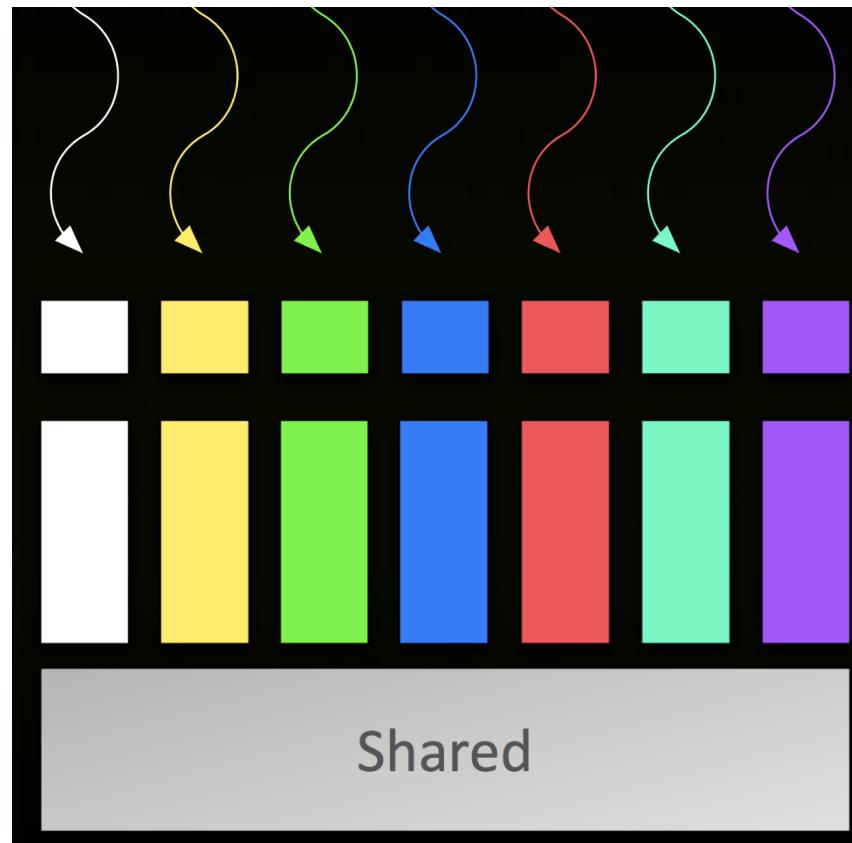
124



GiÜ: Programlama – CUDA

Hafıza yapısı

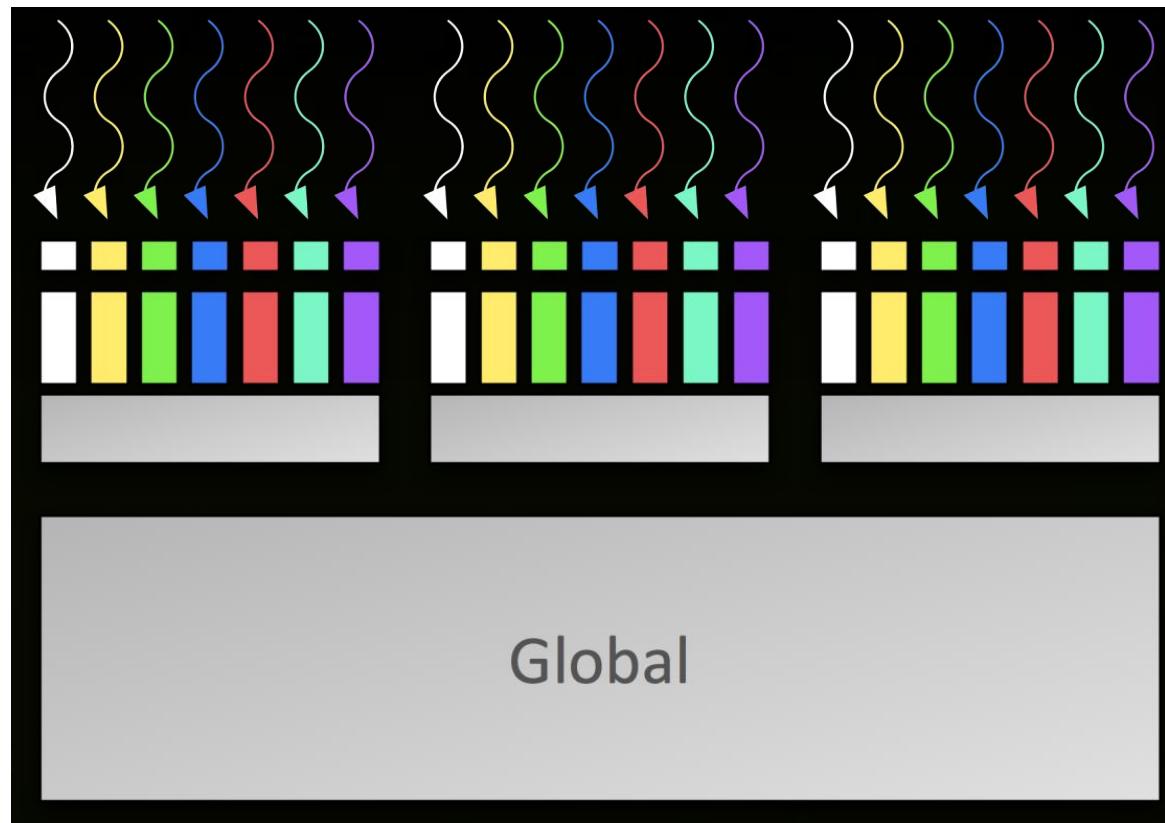
125



GiÜ: Programlama – CUDA

Hafıza yapısı

126



Atomics

2

- Guarantee that only a single thread has access to a piece of memory during an operation
- No dropped data, but ordering is still arbitrary
- Different types of atomic instructions
- Atomic Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor
- Can be done on device memory and shared memory
- Much more expensive than load + operation + store

Example: Histogram

13

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    int c = colors[i];  
    buckets[c] += 1;  
}
```

Example: Histogram

14

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    int c = colors[i];  
    buckets[c] += 1;  
}
```

Why is this incorrect?

Example: Histogram

15

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter atomically  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    int c = colors[i];  
    atomicAdd(&buckets[c], 1);  
}
```

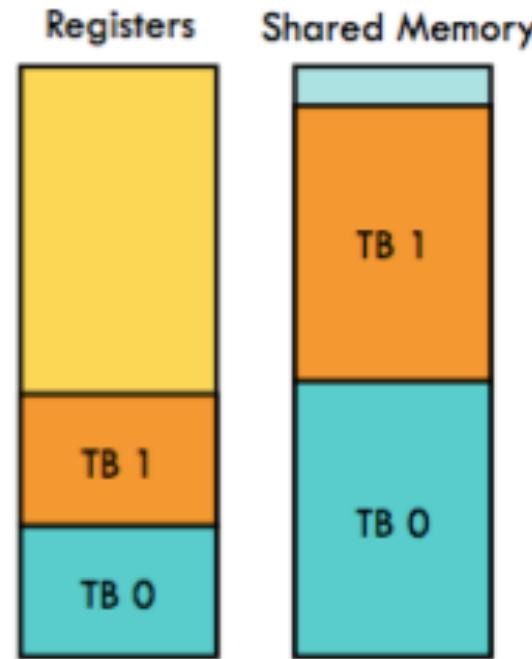
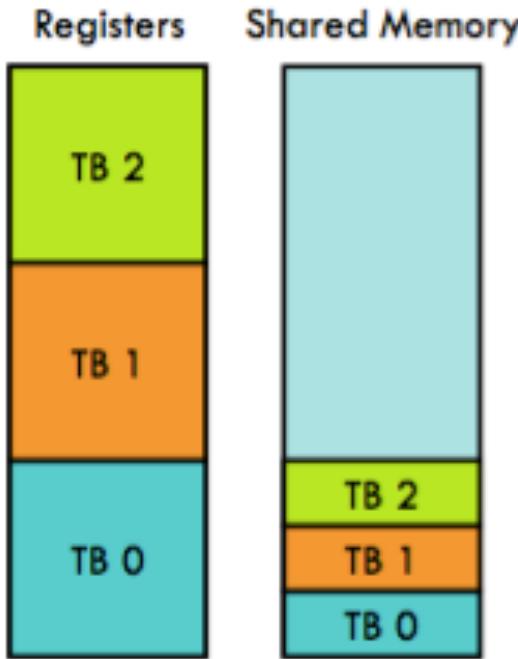
Occupancy

38

- What determines occupancy?
- Limited resources!
 - Register usage per thread
 - Shared memory per thread block

Resource Limits (1)

132



- Pool of registers and shared memory per SM
 - Each thread block grabs registers & shared memory
 - If one or the other is fully utilized → no more thread blocks

Resource Limits (2)

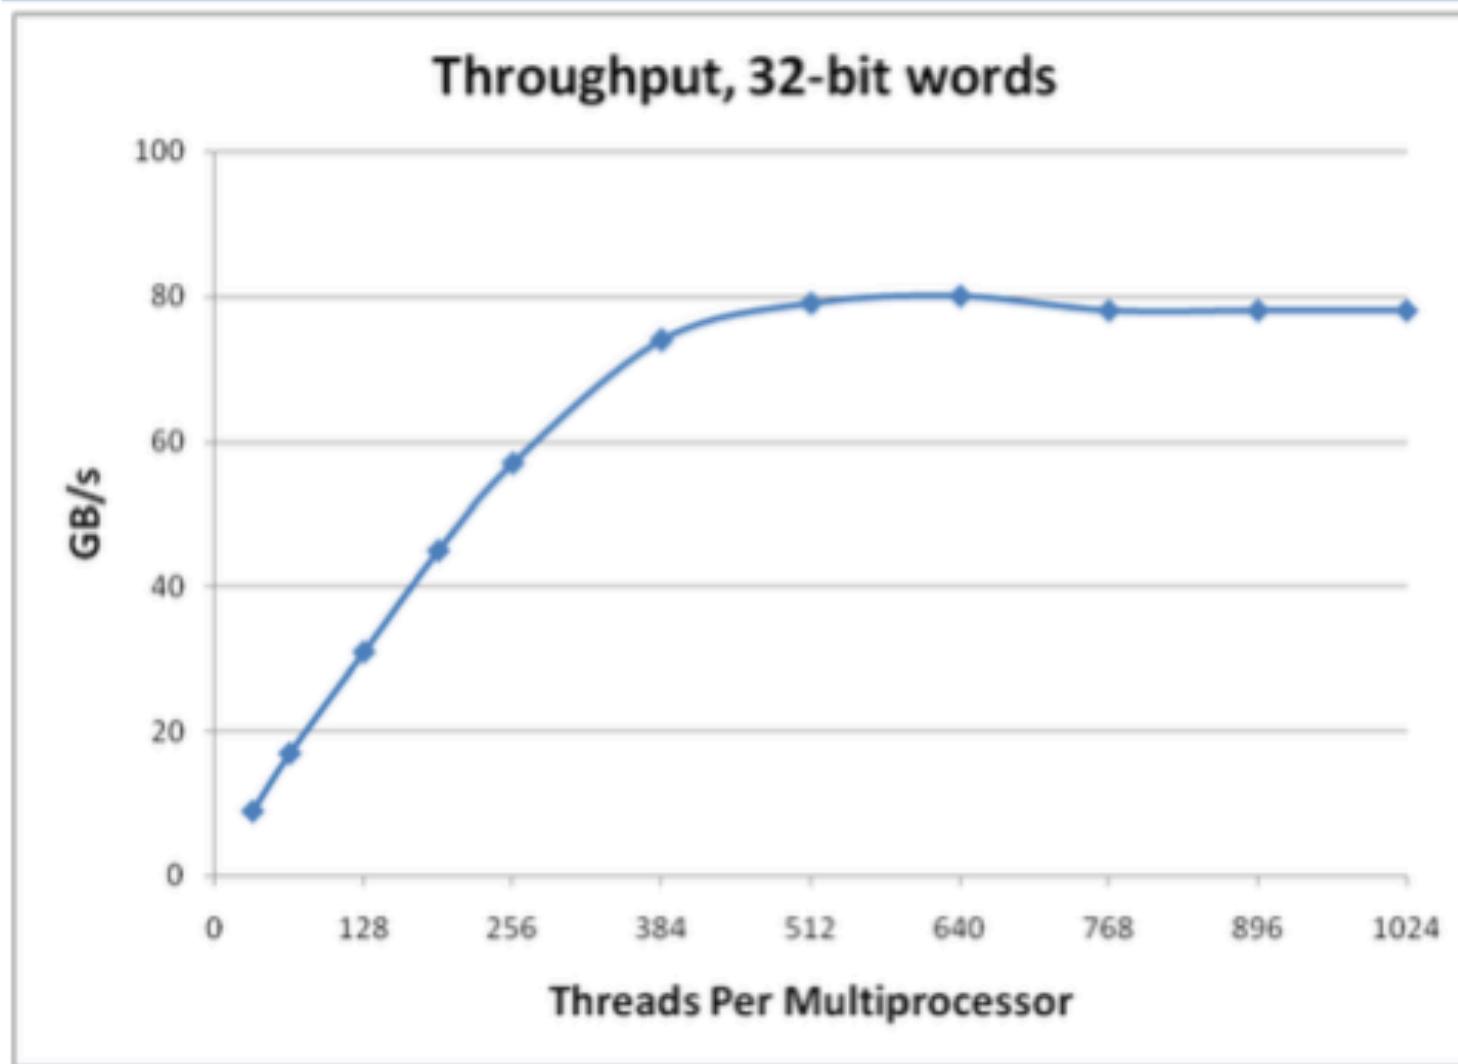
10

- Can only have 8 thread blocks per SM (16 for Kepler)
 - If they're too small, can't fill up the SM
 - Need 128 threads / block on gt200 (4 cycles/instruction)
 - Need 192 threads / block on Fermi (6 cycles/instruction)
- Higher occupancy has diminishing returns for hiding latency

Hiding Latency with more threads

134

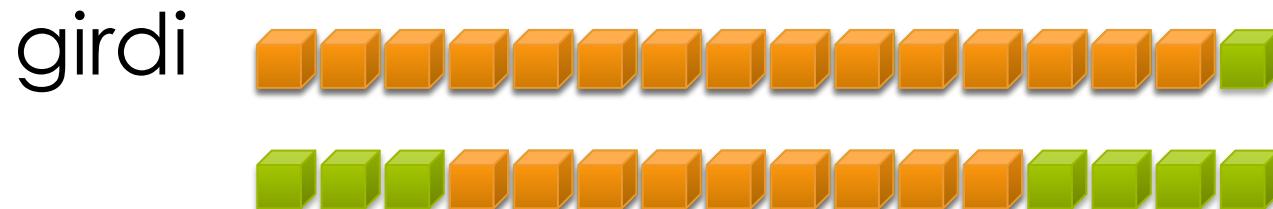
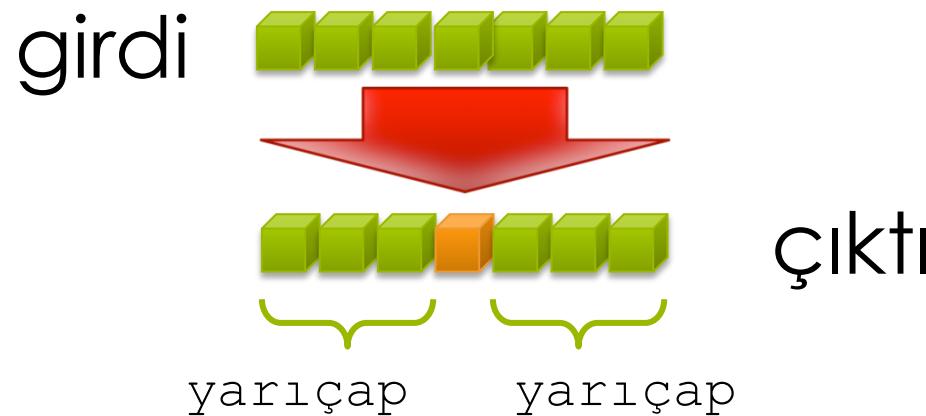
41



Örnek problem: 1D-stencil

Tek boyutlu dizinde

Her çıktı elemanı belli bir yarıçaptaki girdi elemanlarının toplamına eşit



Örnek problem: 1D-stencil

```
__global__ void stencil_1d(int *in, int *out) {  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    for (int offset = -R; offset <= R; offset++)  
        out[i] += in[i + offset];  
}
```

Örnek problem: 1D-stencil

```
__global__ void stencil_1d(int *in, int *out) {  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    int result = 0;  
    for (int offset = -R; offset <= R; offset++)  
        result += in[i + offset];  
    out[i] = result;  
}
```

Örnek problem: 1D-stencil

138

```
__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int lindex = threadIdx.x + RADIUS;

    // Read input elements into shared memory
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        temp[lindex + BLOCK_SIZE] =
            in[gindex + BLOCK_SIZE];
    }
}
```

Örnek problem: 1D-stencil

139

```
// Apply the stencil  
  
int result = 0;  
  
for (int o = -RADIUS ; o <= RADIUS ; o++)  
    result += temp[lindex + o];  
  
  
// Store the result  
  
out[gindex] = result;  
}
```

Tek kod. Bütün işlemciler ve hızlandırıcılar

OpenMP 5.0, OpenACC, OpenCL

OpenMP to Intel Xeon Phi

```
#pragma omp target device(0) map(tofrom:B)
#pragma omp parallel for
for (i=0; i<N; i++)
    B[i] += sin(B[i]);
```

OpenMP to GPGPU

```
#pragma omp target device(0) map(tofrom:B)
#pragma omp teams num_teams(num_blocks)
    num_threads(bsize)
#pragma omp distribute
for (i=0; i<N; i += num_blocks)
    #pragma omp parallel for
    for (b = i; b < i+num_blocks; b++)
        B[b] += sin(B[b]);
```

OpenACC to GPGPU

```
#pragma acc parallel copy(B[0:N]) \
num_gangs(numblocks) \
vector_length(bsize)
#pragma acc loop gang vector
for (i=0; i<N; ++i) {
    B[i] += sin(B[i]);
}
```

```
__global__ void CUDA
vectorAdd(float* a, float* b, float* c) {
    int index = blockIdx.x * blockDim.x + threadIdx.x;
    c[index] = a[index] + b[index];
}
```

```
__kernel void
vectorAdd(__global float* a, __global float* b,
          __global float* c) {
    int index = get_global_id(0);
    c[index] = a[index] + b[index];
}
```

OpenCL

Benchmarking OpenCL, OpenACC, OpenMP, and CUDA: programming productivity, performance, and energy consumption

Suejb Memeti¹ Lu Li² Sabri Pllana³ Joanna Kołodziej⁴ Christoph Kessler⁵

Our major observations include: (1) programming with OpenCL requires significantly more effort than programming with OpenACC for SPEC Accel benchmark suite; (2) programming with OpenCL on average requires about two times more effort than programming with CUDA for Rodinia benchmark suite; (3) the human factor can impact the fraction of code lines to parallelize the code; (4) OpenMP, OpenCL, and CUDA have comparable performance and energy consumption on Ida; and (5) OpenCL performs better than OpenACC for SPEC Accel benchmark suite on Ida; (6) programming with OpenMP requires less effort than programming with OpenCL and CUDA.

Python ile basitçe paralel programlama

- Çok çekirdekli işlemciler...
- Birden fazla iplik kullanımı (OpenMP, CUDA vb.)
- Fakat python'da bu çok verimli olmuyor
- Global Interpreter Lock
 - Güvenli çok çekirdekli işleme

Python ile basitçe paralel programlama

- Python'un multiprocessing kütüphanesinin bir altkümesini inceleyeceğiz.
 - İlk örneklerimizde Process sınıfını kullanacağız.
 - Daha sonra Pool sınıfı ile basit bir paraleleştirme elde edeceğiz.
 - apply, map, apply_async, map_async

Bibliyografi ve ilgili bağlantılar

- **Where HPC & Big Data Intersect**, Bruce Hendrickson, Sandia National Labs
- **A Tale of Two Data-Intensive Paradigms: Applications, Abstractions, and Architectures**, Jha, S., Qiu, J., Luckow, A., Mantha, P., 2014 IEEE International Congress on Big Data (BigData Congress)
- **How Big Data Is Redefining High Performance Computing**, Earl Joseph, Steve Conway, Chirag Dekate, Bob Sorensen
- <http://www.extremetech.com/extreme/185057-supercomputer-stagnation-new-list-of-the-worlds-fastest-computers-makes-exascale-by-2020-seem-unlikely>
- **Accelerating Iterative Big Data Computing Through MPI**, Fan Liang and Xiaoyi Lu
- **A Tale of Two Data-Intensive Paradigms: Applications, Abstractions, and Architectures** Shantenu Jha, Judy Qiu, Andre Luckow, Pradeep Mantha, Geoffrey C.Fox
- <http://www.theplatform.net/2015/02/22/flink-sparks-next-wave-of-distributed-data-processing/>
- http://mechanitis.blogspot.com.tr/2011_07_01_archive.html
- [top500.org/featured/top-systems/](http://sebastianraschka.com/Articles/2014_multiprocessing.html)
- http://sebastianraschka.com/Articles/2014_multiprocessing.html
- <https://docs.python.org/3/library/multiprocessing.html>